

Programming Languages: Semantics and Types

Stephanie Weirich

December 9, 2025

Contents

Preface	vii
What is this course all about?	ix
1 Type Safety for STLC	1
1.1 Syntax	1
1.2 Type system	4
1.3 Operational Semantics	5
1.4 Preservation and Progress	6
1.5 What is type safety?	10
1.6 Further reading	12
2 Natural number recursion	13
2.1 Examples	15
2.2 Exercises	15
2.3 Further Reading	16
3 Big-step semantics	17
3.1 Big-step semantics	18
3.2 Big-step semantics and type safety?	19
3.3 Further Reading	19
4 Termination	21
4.1 Big-step preservation	21
4.2 Big Step Safety	22
4.3 Rephrasing semantic soundness	24
4.4 Variations	25
4.5 Further reading	26
5 REC: Recursive definitions	27
5.1 Recursive definitions in CBV languages	27
5.1.1 Mutual recursion via recursive tuples	27
5.1.2 Another example with recursive tuples	28
5.1.3 Recursive values via recursive types	29
5.2 A fine-grained CBV language	31
5.2.1 Type system	33
5.2.2 Recovering expressiveness	34
5.3 Recursive values	34

5.4	Recursive types	35
5.4.1	Recursive type variations	35
5.5	Type safety	36
5.6	Further reading	36
6	Step-indexing: Semantic type safety for REC	37
6.1	A failing proof: semantic sets	37
6.1.1	Semantic lemmas with small-step evaluation	38
6.1.2	A problem	39
6.2	Fixing the problem: counting evaluation steps	40
6.3	Step-indexed semantic safety	41
6.3.1	Step-indexed propositions	42
6.3.2	Logical relation	42
6.3.3	Semantic typing and semantic typing lemmas	45
6.4	Further Reading	47
7	Type and Effect systems	49
7.1	Core system	49
7.2	Recursive values and types	51
7.3	Syntactic metatheory: Preservation and progress	52
7.4	Effect soundness	53
7.5	Generalizing effects	54
7.6	Expressivity	54
7.7	Variation: Syntax-directed effects	57
7.8	Historical notes and further reading	58
8	A Nontermination monad	59
8.1	Language definition	59
8.1.1	Examples	61
8.2	Restricted typing rule for let	61
8.3	Translation to Dependent Type Theory	62
8.4	The translation to EFF	63
8.5	What does a CBN Monadic language look like?	64
9	Explicit control stacks	67
9.1	Stack-based semantics	67
9.1.1	Primitive vs. control reductions	68
9.2	Stack-based semantics	69
9.2.1	Stack-based small-step semantics	70
9.2.2	Example: small-step semantics without and with stacks	70
9.2.3	Stack-based big-step semantics	72
9.2.4	Example: big-step semantics with stacks	72
9.3	Equivalence between these semantics	73
9.3.1	Stack completeness	73
9.3.2	Soundness	74
9.3.3	Small-step and big-step stack equivalence	74
9.4	References and Further reading	74
10	Control Effects	75
10.1	Exceptions	75
10.1.1	Type Safety in the presence of exceptions	76

10.2 Continuations	78
10.2.1 Continuations and the Curry-Howard Isomorphism	79
10.2.2 CPS conversion	82
10.3 Effect handlers	86
10.3.1 Grammar and stack-based semantics	86
10.3.2 Typing rules and type safety	88
10.4 Further reading	90
11 (Untyped) program equivalence	91
11.1 Preliminaries	91
11.1.1 Scopes and scoped relations	93
11.2 Scoped preorders and equivalences	94
11.2.1 CIU pre-order \leq_{CIU}	94
11.2.2 Contextual pre-order and equivalence (\leq_C)	94
11.2.3 Contextual Equivalence (Pitts) (\leq_{CTX})	95
11.2.4 A logical relation \leq_{log}	96
11.3 Examples	97
11.4 Relating the relations	98
11.4.1 \leq_C implies \leq_{CTX}	99
11.4.2 Show that \leq_{CTX} implies \leq_C	99
11.4.3 \leq_{log} implies \leq_{CIU}	99
11.4.4 \leq_{CIU} implies \leq_{log}	100
11.4.5 \leq_{CIU} implies \leq_{CTX}	100
11.4.6 \leq_{CTX} implies \leq_{CIU}	101
11.5 Further Reading	101
12 Typed Program Equivalence	103
12.1 Preliminaries	103
12.1.1 Nat-equivalence	104
12.1.2 Typed (equivalence) relations	104
12.2 Contextual, CIU, and logical equivalence	105
12.3 Examples	106
12.4 Equivalence between the relations	106
12.5 What next?	107

Preface

This document is a work in progress and summarizes lectures given in CIS 7000-1 during the Fall 2025 semester at the University of Pennsylvania. Expect significant changes to the material over the course of this semester.

The website for the course is:

<https://sweirich.github.io/pl-semantics-and-types/>

The source files for this document are available (<https://github.com/sweirich/pl-semantics-and-types/tree/main/notes>) and have been processed using the LaTeX and Ott [SNO⁺07] tools. Comments, typos and suggestions are appreciated.

All mathematical definitions and proofs have been mechanized using the Rocq proof assistant, with the assistance of the autosubst-ocaml tool. The proof scripts are available at <https://github.com/sweirich/pl-semantics-and-types/tree/main/rocq> and the text includes hyperlinks to each corresponding part of the artifact.

Acknowledgement: These lecture notes draw inspiration from a number of sources, most notably Harper's *Practical Foundations for Programming Languages* [Har16] and Pierce's *Types and Programming Languages* [Pie02]. Thanks to Roger Burtonpatrick, Chris Henson, Jason Liu, and Evelyn Yu for comments and bugfixes.

What is this course all about?

This semester, we are going to consider two big questions that one might ask about programs in a programming language.

1. When are two programs equivalent?

This question is the heart of programming language semantics. Our goal is to understand what programs mean, so understanding when a program is equivalent to some result of evaluation is a way of defining meaning.

However, program equivalence is more than that. It is also essential for software development, for compiler correctness, for program verification, and for security analysis.

For example, when you refactor a program, you are replacing a part of it with "equivalent" code: you want the code to still work, but be more general in some way. Or when a compiler optimizes your program, it replaces part of the code with an "equivalent" but faster sequence of instructions. Program verification often means showing that code is "equivalent" to some mathematical specification. And security analysis is showing that code is not equivalent to programs with certain security flaws.

What makes a good definition of equivalence? How can we reason about this definition mathematically? How can we use this definition in practice?

2. What does it mean to type check a program?

Many programming languages come with static type systems, such as Rust, Java, OCaml or Rocq. Where do these type systems come from? What does it mean when a program type checks? What does it mean when we say that a language is type safe? What can we model using static types? How do types interact with the definition of equivalence?

How will we study these two questions?

Studying these two questions is a study in definitions. And constructing definitions involves design work. A given definition may not be intrinsically wrong, it just may not be useful. So to evaluate our designs we need to identify what we want to do with these definitions and judge them using that metric.

In general, we will study these questions using the tools of programming language theory. This means that we will model idealized versions of programming languages using mathematical objects and then prove properties about these definitions, using the techniques such as induction and coinduction.

How will we stay grounded?

Programming language theory can seem both *trivial* (only applying to tiny "toy" languages) and at the same time *esoteric* (filled with unfamiliar jargon).

Programming languages found "in the wild" are complex and rapidly changing. It can be the work of several years to complete a mathematical specification of a language. A notable example is Andreas Rossberg's work on the semantics of WebAssembly, which you can find more about by watching his keynote from ICFP 23¹.

While such work is important, it does not suit our purposes. Instead we need programming language models that we can understand in hours, not years. We need these models to be representative (i.e. they should describe common features of many different programming languages) and illustrative (i.e. they should describe these features independently of other language constructs). Just as biologist gain understanding through the study of a model organisms such yeast, nematodes or fruit flies, that capture the essence of cell biology, neuroscience, and genetics, we will look at model programming languages that capture the essence of computation, using functions, data structures, control effects, and mutability.

To make sure that we can transfer what we learn from these essential models, we will talk about their connection to languages that you already know, their extension with new features, and their ability to do more than immediately apparent through macro encoding and compilation. Some properties that we prove for small languages will immediately transfer to larger contexts. Others provide a blueprint for how we might redo similar proofs at scale. And still others don't scale, so we also need to be aware of the limitations of our work and when we need to adopt new approaches.

The unfamiliar vocabulary of any discipline is a sign of maturity. It means that researchers have examined and analyzed programming languages for more than seven decades. In the process they have made discoveries, and communicating those discoveries requires precise naming and notation. One goal of this semester is to learn these concepts, along with their unfamiliar names, and understand how they may be put to use in practice. By the end of the semester, you should be better equipped to talk about languages and read research papers about programming language advances.

What is the role of proof assistants?

If you have read Software Foundations [PdAC⁺²⁵] or have taken a course like CIS 5000², you have already used the metalanguage of the Rocq proof assistant to mathematically model small programming languages, such as the simply-typed lambda calculus.

I am a *strong* believer in the value of these tools. They both solidify your understanding of the metalogic that we are working in as well as give immediate feedback on your progress. They are also fun to use.

However, this is *not* a course on the use of proof assistants. You do not need to have experience with Rocq or any similar tool in order to benefit from this course. The homework assignments will be in LaTeX and the exams will be on paper and

¹<https://www.youtube.com/watch?v=Lb45xIcqGjg>

²<https://www.seas.upenn.edu/~cis5000/current/index.html>

will cover topics related to programming language theory.

That said, I will ensure that the material that we study this semester is amenable to mechanical development. I will be developing and checking the topics that we cover throughout the semester using Rocq and will make my code available. (Caveat: I hope that I can keep up!) You can use this code as the basis for the homework assignments and I will gladly answer any questions and provide assistance during office hours. If you want to learn how to use Rocq, this is a good opportunity. I will also assist if you would like to translate this code to another framework (such as Agda or LEAN).

1

Type Safety for a Simply-Typed Lambda Calculus

This section gives a precise definition of the syntax of a simply-typed lambda calculus, its type system and small-step operational semantics. For conciseness, we often refer to this language as STLC.

If you are new to programming language theory, this section also introduces some of the mathematical concepts that we will be using throughout the semester, such as inductively defined grammars, recursive definitions, and proofs by structural induction.

STLC is actually a family of simple languages, with some freedom in the sorts of features that are included. There must always be some sort of “primitive” type such as booleans, numbers, or even a unit type. And STLC always includes first-class functions, i.e. λ -terms, making it a simplified version of typed functional languages such as ML or Haskell. However, in other contexts you may see it extended with various other features, such as records, products, disjoint unions, variant types, etc.

1.1 Syntax

The syntax of the simply-typed lambda calculus is defined by a set of terms and their associated set of types. By convention, we will use the metavariable e to refer to some arbitrary term and τ to refer to some arbitrary type. If you are familiar with algebraic datatypes, or inductive datatypes, you can think of the following definitions along those lines.

Definition 1.1.1 (Types). The set of **types** is inductively defined by the following rules:

1. A base type, Nat , is a type.
2. If τ_1 and τ_2 are types, then $\tau_1 \rightarrow \tau_2$ is a type.

The type $\tau_1 \rightarrow \tau_2$ represents the type of functions that take an argument of type τ_1 and return a value of type τ_2 .

Definition 1.1.2 (Terms). The set of **terms** is inductively defined by the following rules:

1. A natural number k is a term.
2. A variable x is a term.
3. If e is a term and x is a variable, then $\lambda x.e$ is a term (called a lambda abstraction). The variable x is the parameter and e is the body of the abstraction.
4. If e_1 and e_2 are terms, then $e_1 e_2$ is a term (called a function application).

The definition of terms refers to two other sets: natural numbers and variables. The set of natural numbers, \mathbb{N} , are an infinite set of numbers $0, 1, \dots$; we will use i, j and k to refer to arbitrary natural numbers. We treat variables more abstractly. We assume that there is some infinite set of variable *names*, called \mathcal{V} , and that given any finite set of variables, we can always find some variable that is not in contained in that set. (We call such a variable *fresh* because we haven't used it yet.) If you like, you can think of names more concretely as strings or numbers, but we won't allow all of the usual operations on strings and numbers to be applied to names.

Now, the above definitions are a wordy way of describing an inductively-defined grammar of abstract syntax trees. In the future, we will use a more concise notation, called Bakus-Naur form. For example, in BNF, we can provide a concise definition of the grammars for types and terms as follows.

Definition 1.1.3 (STLC Syntax (concise form)).

$$\begin{array}{lll} \text{numbers} & i, j, k & \in \mathbb{N} \\ \text{variables} & x & \in \mathcal{V} \\ \text{types} & \tau & ::= \text{Nat} \mid \tau_1 \rightarrow \tau_2 \\ \text{terms} & e & ::= k \mid x \mid \lambda x.e \mid e_1 e_2 \end{array}$$

Free variables Because types and terms are inductively defined sets, we can reason about them using recursion and induction principles. The recursion principle means that we define recursive functions that takes terms or types as arguments and know that the functions are total, as long as we call the functions over smaller subterms.

For example, one function that we might define calculates the set of *free* variables in a term.

Definition 1.1.4 (Free variables). We define the operation $\text{fv}(e)$, which calculates the set of variables that occur *free* in some term e , by structural recursion.

$$\begin{array}{lll} \text{fv}(k) & = \emptyset & \text{empty set} \\ \text{fv}(x) & = \{x\} & \text{a singleton set} \\ \text{fv}(e_1 e_2) & = \text{fv}(e_1) \cup \text{fv}(e_2) & \text{union of sets} \\ \text{fv}(\lambda x.e) & = \text{fv}(e) - \{x\} & \text{remove variable } x \end{array}$$

Each of the lines above describes the behavior of this function on the different sorts of terms. If the argument is a natural number constant k , then it contains no

free variables, so the result of the function is the \emptyset . Otherwise, if the argument is a single variable, then the function returns a singleton set. If the argument is an application, then we use recursion to find the free variables of each subterm and then combine these sets using an “union” operation. Finally, in the last line of this function, we find the free variables of the body of an abstraction, but then remove the argument x from that set because it does not appear free in entire abstraction.

Variables that appear in terms that are not free are called *bound*. For example, in the term $\lambda x.x\ y$, we have x bound and y free. Furthermore, some variables may occur in both bound and free positions in terms; such as x in the term $(\lambda x.x\ y)\ x$.

Renaming Here is another example of a recursively defined function. Sometimes we would like to change the names of free variables in terms.

A *renaming*, ξ , is a mapping from variables to variables. A renaming has a *domain*, $\text{dom } \xi$ and a *range* $\text{rng } \xi$. We use the notation y/x for a single renaming that maps x to y , and the notation $y/x, \xi$ to extend an existing renaming with a new replacement for x .

Definition 1.1.5 (Renaming application). We define the application of a renaming to a term, written with postfix notation $e\langle\xi\rangle$, as follows:

$$\begin{aligned} k\langle\xi\rangle &= k \\ x\langle\xi\rangle &= \xi x \\ (e_1\ e_2)\langle\xi\rangle &= (e_1\langle\xi\rangle)\ (e_2\langle\xi\rangle) \\ (\lambda x.e)\langle\xi\rangle &= \lambda y.(e\langle y/x, \xi \rangle) \text{ for } y \text{ not in } \text{rng } \xi \end{aligned}$$

We can only apply a renaming to a term when its domain includes the free variables defined in the term. In that case, our renaming function is total: it produces an answer for any such term.

We have to be a bit careful in the last line of this definition. What if ξ already maps the variable x to some other variable? What if ξ already maps some other variable to x ? Our goal is to only rename free variables: the function should leave the bound variables alone. Inside the body of $\lambda x.e$, the variable x occurs bound, not free. On the other hand, if we introduce a new x through renaming an existing free variable, we do not want it to be *captured* by the function. For example, if we rename x to y , in the function $\lambda x.y\langle x/y \rangle$, we do not want to produce $\lambda x.x$.

Therefore, we pick some fresh variable y , and updating the renaming to $(y/x, \xi)$ in the recursive call. (If x is already fresh, we can keep using it.) That way, we force the renaming that we use for the body of the abstraction to not change the binding structure of the term.

Substitution There is one final definition of a function defined by structural recursion over terms: the application of a *substitution* that applies to all free variables in the term.

A *substitution*, σ is a mapping from variables to terms. As above, it has a *domain* (a set of variables) and a *range* (this time a set of terms). We use the notation $e/x, \sigma$ to refer to the substitution that maps variable x to term e , but otherwise acts like σ .

As before, this definition only applies when the free variables of the term are contained within the domain of the substitution. Furthermore, when substituting in the body of an abstraction, we must be careful to avoid variable capture.

Definition 1.1.6 (Substitution application). We define the application of a substitution function to a term, written with postfix notation $e[\sigma]$, as follows:

$$\begin{aligned} k[\sigma] &= k \\ x[\sigma] &= \sigma x \\ (e_1 e_2)[\sigma] &= (e_1[\sigma]) (e_2[\sigma]) \\ (\lambda x.e)[\sigma] &= \lambda y.(e[y/x, \sigma]) \text{ when } y \notin \text{fv}(\text{rng } \sigma) \end{aligned}$$

Variable binding, alpha-equivalence and all that At this point, we will start to be somewhat informal when it comes to bound variables in terms. As you see above, we need to be careful about variable capture when doing renaming and substitution. But we don't want to pollute our reasoning later with these details.

Fortunately, we also don't want to distinguish between terms that differ only in their use of bound variables, such as $\lambda x.x$ and $\lambda y.y$. There is a relation called α -equivalence that relates such terms, and from this point forward we will say that our definitions are “up-to- α -equivalence”. What this means practically is that on one hand, we must be sure that our definitions don't really depend on the names of bound variables. In return, we can always assume that any bound variable is distinct from any other variable, if we need it to be. This practice is called the “Barendregt Variable Convention” [Bar84].

But, note that this is an informal convention, allowing us to follow the common practice of describing lambda calculus terms as we have done above (sometimes called using a named or nominal representation of variables). But getting the details right is difficult (it requires maintaining careful invariants about all definitions) and subtle. If you are working with a proof assistant, you really do need to get the details right. In that context, it also makes sense to use an approach (such as de Bruijn indices [de 72]) where the details are easier to get right. This is what we will do in the accompanying mechanized proofs.

However, because using a named representation is standard practice, we will continue to use that approach in these notes, glossing over details. This will allow us to stay roughly equivalent to the proof scripts (which have other details). Because of the informal nature of our discussion, there will be minor omissions related to variable naming; we won't stress about them.

1.2 Type system

Next we will define a typing relation for STLC. This relation has the form $\Gamma \vdash e \in \tau$, which is read as “in the typing context Γ , the term e has type τ .” If a term is in this relation we say that it “type checks”. The **typing context** Γ , tells us what the types of free variables should be. Therefore, we model it as a finite map from variables to types, and write it by listing all of the associations, for example $x:\tau_1, y:\tau_2, z:\tau_3$. If the context has no associations, we call it the *empty context* and leave it blank, writing for example $\vdash 2 \in \text{Nat}$. Terms that type check with an empty context are called *closed* and have no free variables.

We define the typing relation inductively, using the following rules. A term type checks if we can find some tree that puts these rules together in a *derivation*. In each rule, the part below the line is the conclusion of the rule, and the rule may have multiple premises. In a derivation tree, each premise must be satisfied by subderivations, bottoming out with rules such as rule **T-VAR** or rule **T-LIT** that do not have any premises for the same relation.

Definition 1.2.1 (STLC type system).

$$\begin{array}{c}
 \boxed{\Gamma \vdash e \in \tau} \qquad \qquad \text{(in context } \Gamma, \text{ term } e \text{ has type } \tau\text{)} \\
 \begin{array}{cccc}
 \text{T-LIT} & \text{T-VAR} & \text{T-ABS} & \text{T-APP} \\
 \frac{}{\Gamma \vdash k \in \mathbf{Nat}} & \frac{x : \tau \in \Gamma}{\Gamma \vdash x \in \tau} & \frac{\Gamma, x : \tau_1 \vdash e \in \tau_2}{\Gamma \vdash \lambda x. e \in \tau_1 \rightarrow \tau_2} & \frac{\begin{array}{c} \text{T-APP} \\ \Gamma \vdash e_1 \in \tau_1 \rightarrow \tau_2 \\ \Gamma \vdash e_2 \in \tau_1 \end{array}}{\Gamma \vdash e_1 e_2 \in \tau_2}
 \end{array}
 \end{array}$$

In the variable rule, we look up the type of the variable in the typing context. This variable must have a definition in Γ for this rule to be used. If there is no type associated with x , then we say that the variable is unbound and that the term fails to *scope check*.

In rule T-ABS, the rule for abstractions, we type check the body of the function with a context that has been extended with a type for the bound variable. The type of an abstraction is a function type $\tau_1 \rightarrow \tau_2$, that states the required type of the parameter τ_1 and the result type of the body τ_2 .

Rule T-APP, which checks the application of functions, requires the argument to have the type required by the function.

1.3 Operational Semantics

Is this type system meaningful? Our type system makes a distinction between terms that type check (such as $(\lambda x. x) 3$) and terms that do not, such as $(2 5)$. But how do we know that this distinction is useful? Do we have the right rules?

The key property that we want is called *type safety*. If a term type checks, we should be able to evaluate it without triggering a certain class of errors.

One way to describe the evaluation of programs is through a *small-step* operational semantics. This is a mathematical definition of a relation between a program e and its value. We build up a small-step semantics in two parts. First, we define a single step relation, written $e \rightsquigarrow e'$, to mean that a term reduces to e in one step. Then we iterate this relation, called the multistep relation and written $e \rightsquigarrow^* e'$, to talk about all of the different programs that e could reduce to after any number of steps, including 0.

The multistep evaluations that we are interested in are the ones where we do some number of small steps and get to an e' that has a very specific form, a *value*. If we have $e \rightsquigarrow^* v$ then we say that e evaluates to v .

Definition 1.3.1 (Value). A *value* is an expression that is either a natural number constant or an abstraction.

$$v ::= k \mid \lambda x. e$$

We define the single step relation inductively, using the inference rules below that state when one term steps to another.

Definition 1.3.2 (Small-step relation).

$$\begin{array}{c}
 \boxed{e \rightsquigarrow e'} \qquad \qquad \text{(term } e \text{ steps to } e'\text{)} \\
 \begin{array}{ccc}
 \text{S-BETA} & \text{S-APP-CONG1} & \text{S-APP-CONG2} \\
 \frac{(\lambda x. e) v \rightsquigarrow e[v/x]}{} & \frac{\begin{array}{c} e_1 \rightsquigarrow e'_1 \\ e_1 e_2 \rightsquigarrow e'_1 e_2 \end{array}}{} & \frac{\begin{array}{c} e_2 \rightsquigarrow e'_2 \\ v e_2 \rightsquigarrow v e_2 \end{array}}{}
 \end{array}
 \end{array}$$

In each of these three rules, the part below the line says when the left term steps to the right term. Rule S-BETA describes what happens when an abstraction is applied to an argument. In this case, we substitute the argument for the parameter in the body of the function. Note in this rule that the argument must be a value before substitution. If it is not a value, then we cannot use this rule to take a step. This rule is the key of a *call-by-value* semantics.

The second two rules each have premises that must be satisfied before they can be used. Rule S-APP-CONG1 applies when the function part of an application is not (yet) an abstraction. Similarly, the last rule applies when the argument part of an application is not (yet) a value.

This small-step relation is intended to be deterministic. Any term steps to at most one new term.

Lemma 1.3.1 (Determinism). If $e \rightsquigarrow e_1$ and $e \rightsquigarrow e_2$ then $e_1 = e_2$.

The small-step relation is *not* a function. For some terms e , there is no term e' such that $e \rightsquigarrow e'$. For example, if we have a number in the function position, e.g. $(3\ e)$, then the term does not step and these terms do not evaluate to any value.

This is important. These terms are called *stuck* and correspond to crashing programs. For example, if we tried to use a number as function pointer in the C language, then we might get a segmentation fault.

1.4 Preservation and Progress

Type safety is a crucial property of a typed programming language. It ensures that a well-typed program will never “go wrong” during execution. For the simply-typed lambda calculus, this means a program will not get stuck in a state where it cannot take a reduction step but is not a final value.

The type safety proof is usually defined through two lemmas: Preservation and Progress.

Preservation The *preservation* lemma property states that if a term e has type τ , and it takes a single reduction step to e' , then the new term must also have the exact same type τ . In other words, the type is “preserved” through evaluation.

Lemma 1.4.1 (Preservation). If $\vdash e \in \tau$ and $e \rightsquigarrow e'$ then $\vdash e' \in \tau$.

We can prove this lemma in three separate ways: by structural induction on the syntax of e , or by induction on the derivations of $\vdash e \in \tau$ or $e \rightsquigarrow e'$. This flexibility is enabled by the simplicity of this type system. For example, we have exactly one typing rule for each syntactic form, and each typing rule has a corresponding premise for each subterm.

Because this is our first inductive proof, we will first prove it by induction on the syntax and then by induction on the step relation.

Proof. Proof is by induction on the syntax of e . We want to prove that the statement of the lemma holds for any arbitrary e . Our induction principle requires that we prove that it holds for the four different syntactic forms, natural numbers, variables, applications and abstractions. In the latter two cases, we will be able to assume that the lemma holds for each of the subterms of the form.

- If e is a natural number k , then we want to show that if $\vdash k \in \tau$ and $k \rightsquigarrow e'$ then $\vdash e' \in \tau$. However, by looking at our operational semantics, we see that constants don't step, so there cannot be any such e' . So this case is immediate.
- If e is a variable x , then we want to show that if $\vdash x \in \tau$ and $x \rightsquigarrow e'$ then $\vdash e' \in \tau$. Again, this case is impossible, and for two reasons. We cannot type check variables in an empty context and there is no rule of the operational semantics that applies.
- If e is $\lambda x.e$, then this case also impossible as abstractions don't step.
- If e is of the form $e_1 e_2$, then we want to show that if $\vdash e_1 e_2 \in \tau$ and $e_1 e_2 \rightsquigarrow e'$ then $\vdash e' \in \tau$. In this case, we can look at the typing rules to observe that there is only one way to type check an application. For this application to type check, the derivation must also show that $\vdash e_1 \in \tau_2 \rightarrow \tau$ and $\vdash e_2 \in \tau_2$. (This reasoning principle is called *inversion*.) We can also use inversion on the step relation. But this time there are three ways that an application could step.
 - The application could step using rule **S-BETA**. In this case, we know that e_1 is $\lambda x.e'_1$ and e_2 is some value v . Furthermore, we have that the result of the application step e' is $e'_1[v/x]$. Therefore, we need to show that the result of this substitution has type τ . By inverting the typing judgement for $\vdash \lambda x.e'_1 \in \tau_2 \rightarrow \tau$, we know that $x : \tau_2 \vdash e'_1 \in \tau$. At this point we, we will appeal to a *substitution lemma* (see 1.4.1 below) to finish this case of the proof.
 - The application could step using rule **S-APP-CONG1**. In this case, we know that $e_1 \rightsquigarrow e'_1$ and that e' is $e'_1 e_2$. Because e_1 is a subterm of $e_1 e_2$, we can assume that the preservation holds for that term. Therefore, we know that $\vdash e'_1 \in \tau_2 \rightarrow \tau$. We then use this fact with rule **T-APP** to conclude that $\vdash e'_1 e_2 \in \tau$.
 - The application could step using rule **S-APP-CONG2**. In this case, we know that e_1 is a value and that $e_2 \rightsquigarrow e'_2$. As above, we can assume that the preservation holds for e_2 . Therefore, we know that $\vdash e'_2 \in \tau_2$. We then use this fact with rule **T-APP** to conclude that $\vdash e_1 e'_2 \in \tau$.

□

For comparison, we also prove this theorem by induction on one of the derivations. The proof is not significantly different: the same pieces of the argument are necessary. However, we put them together in a different way, replacing some uses of inversion with the induction principle.

Proof. The proof is by induction on the derivation of $e \rightsquigarrow e'$. There are three cases, one for each of the rules that could have been used to conclude $e \rightsquigarrow e'$.

- In the case of rule **S-BETA**, we have that e is of the form $(\lambda x.e) v$ and e' is $e[v/x]$. We also know that the first term type checks, i.e. that $\vdash (\lambda x.e) v \in \tau$. For this term to type check, we must have used rule **T-APP**, so by inversion, we also know that $\vdash (\lambda x.e) \in \tau_1 \rightarrow \tau$ and $\vdash v \in \tau_1$. We can do this again, because the only way to make an abstraction to type check is rule rule **T-ABS**, so we must have also shown $x : \tau_1 \vdash e \in \tau$. At this point we, we will again appeal to the substitution lemma (see 1.4.1 below) to finish this case of the proof.

- In the case of rule **S-APP-CONG1**, we have the conclusion $e_1 e_2 \rightsquigarrow e'_1 e_2$, and premise $e_1 \rightsquigarrow e'_1$. For the first term to type check, we again must have also used rule **T-APP**, so we know that $\vdash e_1 \in \tau_1 \rightarrow \tau$ and $\vdash e_2 \in \tau_1$. In this case we can use induction, because we know that e_1 , a term in the subderivation both steps and type checks. So we know that $\vdash e'_1 \in \tau_1$. Now we can use rule **T-APP** to conclude that $\vdash e'_1 e_2 \in \tau$.
- This case is similar to the one above.

□

Substitution In the rule **S-BETA** case, our proof above relies on this lemma, that we can write more formally:

Corollary 1.4.1 (Single Substitution). If $x : \tau_1 \vdash e \in \tau_2$ and $\vdash v \in \tau_2$ then $\vdash e[v/x] \in \tau_1$

However, to prove this lemma, we must first generalize it. We cannot prove the lemma directly as stated, because we need a version that gives us a stronger induction hypothesis. To see where we run into trouble, let's walk through a potential proof and see where we get stuck.

We can try to prove this lemma by structural induction on e . That means that we have four cases. The cases for variables, constants and applications go through without difficulty. Now consider the case for abstractions. Say we have $x : \tau_1 \vdash \lambda y. e_1 \in \tau_2$. We want to show that $\vdash (\lambda y. e_1)[v/x] \in \tau_2$. From the definition of substitution, we know that this is equivalent to showing that $\vdash \lambda y. e_1[v/x] \in \tau_2$, implicitly using the variable convention to assume that y is not the same as x . By inversion on the typing judgement, we also know that τ_2 is some function type $\tau_3 \rightarrow \tau_4$ and that $x : \tau_1, y : \tau_3 \vdash e_1 \in \tau_4$. From rule **T-ABS**, it suffices to show $y : \tau_1 \vdash e_1[v/x] \in \tau_4$. However, this result is not available to us through induction: the lemma only applies to terms that type check in a context with exactly one variable. Although e_1 is a subterm, it type checks in a context with two variables. So we cannot make any more progress on this proof.

Therefore, we generalize the substitution lemma in two ways. First, we allow the term to type check in any context Γ . Then, the lemma works for any substitution σ that replaces every variable in $\text{dom } \Gamma$, i.e. in scope, to a term of the appropriate type. However, the range of σ need not be closed: we use Δ to describe the types of variables that can appear in the range of σ .

Lemma 1.4.2 (Simultaneous substitution). If $\Gamma \vdash e \in \tau$ and for all $x \in \text{dom } \Gamma$, we have $\Delta \vdash \sigma x \in \Gamma x$, then $\Delta \vdash e[\sigma] \in \tau$.

Proof. Proof is by structural induction on e . That means that we have four cases. Again, the cases for variables, constants and applications go through without difficulty. Now consider the case for abstractions. Say we have $\Gamma \vdash \lambda y. e_1 \in \tau_2$. We want to show that $\vdash (\lambda y. e_1)[\sigma] \in \tau_2$.

From the definition of substitution, we know that this is equivalent to showing that $\vdash \lambda y. (e_1[y/y, \sigma]) \in \tau_2$, implicitly using the variable convention to assume that y is not in the domain of σ or free in the range of σ . (The substitution $(y/y, \sigma)$ is a map that is just like σ , but maps the variable y to itself.) By inversion on the typing judgement, we also know that τ_2 is some function type $\tau_3 \rightarrow \tau_4$ and that $\Gamma, y : \tau_3 \vdash e_1 \in \tau_4$. From rule **T-ABS**, it suffices to show $\Delta, y : \tau_3 \vdash e_1[y/y, \sigma] \in \tau_4$. This time we can use our inductive hypothesis on the typing derivation for e_1 .

However, to do so, we need to show that for all x in $\text{dom } \Gamma$, $y : \tau_3$, we have $\Delta, y : \tau_3 \vdash (y/y, \sigma) x \in (\Gamma, y : \tau_3) x$. But we know that $\Delta, y : \tau_3 \vdash y \in \tau_3$ and we already know that the rest of the substitution is well-typed. \square

This is not the only way to strengthen our substitution lemma. If our typing contexts are not ordered we can stick with single substitutions. Not ordered means that we consider $\Gamma, x : \tau_1$ to be the same context as $x : \tau_1, \Gamma$.

Lemma 1.4.3 (Substitution (Unordered context)). If $\Gamma, x : \tau_1 \vdash e \in \tau$ and $\Gamma \vdash v \in \tau_1$, then $\Gamma \vdash e[v/x] \in \tau$.

If contexts are ordered, then another way to strengthen this lemma is to let the variable being substituted for appear anywhere in the middle of the context. However, the proof of this version of the lemma requires an additional property called weakening (shown below).

Lemma 1.4.4 (Substitution (Ordered context)). If $\Gamma, x : \tau_1, \Gamma' \vdash e \in \tau$ and $\Gamma \vdash v \in \tau_1$, then $\Gamma, \Gamma' \vdash e[v/x] \in \tau$.

Weakening is a corollary of our strongest substitution lemma where we set $\Delta = \Gamma, x : \tau$ and σ to be the identity function.

Lemma 1.4.5 (Single Weakening). If $\Gamma \vdash e \in \tau$ then $\Gamma, x : \tau_1 \vdash e \in \tau$.

Progress The second lemma, called *progress*, states that any well-typed term that has not been completely reduced can always take at least one more reduction step. It ensures that a well-typed term is not “stuck.” (i.e. is not a value but cannot step).

Lemma 1.4.6 (Progress). If $\vdash e \in \tau$ then either e is a value or there exists an e' such that $e \rightsquigarrow e'$.

Proof. We prove this lemma by induction on the typing derivation. In the rules where e is already a value, then the proof is trivial. Therefore we only need to consider when e is an application of the form $e_1 e_2$, where $\vdash e_1 \in \tau_1 \rightarrow \tau$ and $\vdash e_2 \in \tau_1$. By induction on the first premise, we know that either e_1 is a value or that it takes a step to some e'_1 . If it takes a step, the entire application takes a step by rule **S-APP-CONG1** and we are done. Otherwise, if it is a value, then we know that it must be of the form $\lambda x. e'$, because it must have a function type. By induction on the second premise, we know that either e_2 is a value or that it takes a step to some e'_2 . In the former case, the application steps to $e'[e_2/x]$ by rule **S-BETA**, in the latter case, the application steps to $(\lambda x. e') e'_2$ by rule **S-APP-CONG2**. \square

Part of this proof involves inferring the structure of a closed value from its type. Although this is a straightforward bit of reasoning by inversion, this is a key step of any progress proof. Therefore, we explicitly state the lemma (and its analogue for natural numbers below).

Lemma 1.4.7 (Canonical forms (arrow types)). If $\vdash v \in \tau_1 \rightarrow \tau_2$ then v is some $\lambda x. e$.

Lemma 1.4.8 (Canonical forms (nat)). If $\vdash v \in \text{Nat}$ then v is some natural number k .

1.5 What is type safety?

We above claimed that type safety means that well-typed programs do not get stuck. But what does this mean? Is that what we have really proven?

There are languages and type systems that do not satisfy both of these lemmas, yet we still might like to say that they are type safe. Can we come up with a more general definition? Something that is implied by preservation/progress but doesn't itself require them to be true.

Perhaps we would like to prove something like below, where the multistep relation \rightsquigarrow^* is iteration of the single-step relation any number of times. If a closed term type checks then it must evaluate to a value with the same type.

Conjecture 1.5.1 (Terminating Type Safety). If $\vdash e \in \tau$ then there exists some value v such that $e \rightsquigarrow^* v$ and $\vdash v \in \tau$.

This conjecture seems straightforward to prove from progress and preservation. By progress we know that either a term is a value or that it steps. By preservation, we know that if it steps, it has the same type. But what we are missing from a straightforward proof is the fact that this conjecture says that evaluation *terminates*. How do we know that we will eventually reach a value in some finite number of steps?

It turns out that this conjecture is true, but we are not yet ready to prove it directly. But even though the conjecture is true, it is not a good definition of type safety: even though all well-typed STLC programs halt, that is not true of most programming languages. We would like to have a definition of type safety that also applies to nonterminating languages—one that shows that well-typed programs do not get stuck, while not requiring them to produce values.

There are several solutions to this issue.

Well-typed programs don't get stuck The most straightforward approach is to define what it means for a program to get stuck, and then show that this cannot happen.

Definition 1.5.1 (Stuck). A term e is *stuck* if it is not a value and there does not exist any e' such that $e \rightsquigarrow e'$.

Theorem 1.5.1 (Type safety (no stuck terms)). If $\vdash e \in \tau$ then for all e' , such that $e \rightsquigarrow^* e'$, e' is not stuck.

Proof. We prove this by induction on the derivation of $e \rightsquigarrow^* e'$. If there are no steps in this reduction sequence, then e is equal to e' . By the progress lemma, we know that e is not stuck. Otherwise, say that there is at least one step, i.e. there is some e_1 such that $e \rightsquigarrow e_1$ and $e_1 \rightsquigarrow^* e'$. By preservation, we know that $\vdash e_1 \in \tau$. Then we can use induction to say that e' is not stuck. \square

A coinductive definition What if we want to state type safety a little more positively? In other words, we want to say that a well-typed term either produces a value or runs forever, without having to talk about stuckness.

We can do that using the following *coinductive* definition.

Definition 1.5.2 (Runs safely). A program e *runs safely*, if it is a value or if $e \rightsquigarrow e'$, and e' *runs safely*.

This is exactly the definition we want to use in a type safety theorem.

Theorem 1.5.2 (Type Safety (runs safely)). If $\vdash e \in \tau$ then e runs safely.

Just as in an inductive definitions, the definition of “runs safely” refers to itself. But we are interpreting this definition coinductively, so it includes both finite and infinite runs. In other words, if a program steps to another program, which steps to another program, and so on, infinitely, then it is included in this relation.

Coinductive definitions come with *coinduction* principles. We usually use induction principles to show that some property holds about an element of an inductive definition that we already have. As we “consume” this definition, we can assume, by induction, that the property is true for the subterms of the definition. For example, when proving the preservation lemma, we assumed that the lemma held for the subterms of the evaluation derivation.

The principle of coinduction applies when we want to “generate” an element of a coinductive definition. Watch!

We will prove type safety through coinduction. Given a well typed term $\vdash e \in \tau$, the progress lemma tells us that it is either a value or that it steps. If it is a value, then we know directly that it runs safely. If it steps, i.e. if we have $e \rightsquigarrow e'$, then by preservation, we know that $\vdash e' \in \tau$. By the principle of coinduction, we know that e' runs safely. So we can conclude that e runs safely.

When are we allowed to use a coinductive hypothesis? With induction, we were limited to “consuming” subterms or smaller derivations. But when we use a coinductive hypothesis, it cannot be the last step of the proof. We need to do something with the result of this hypothesis to generate our coinductive definition.

This can be a bit confusing at first, and I encourage you to look at proofs completed with coinduction in the first place to get the hang of using this principle.

An inductive definition Alternatively, if you are still uncomfortable with coinduction, we can define what it means to run safely another way.

We say that an expression e steps to e' in k steps using the following inductive definition.

Definition 1.5.3 (Counted steps). $\boxed{e \rightsquigarrow^k e'}$ (k steps)

$$\frac{\text{MS-K-REFL}}{e \rightsquigarrow^0 e} \qquad \frac{\text{MS-K-STEP}}{e_0 \rightsquigarrow e_1 \quad e_1 \rightsquigarrow^k e_2}{\frac{}{e_0 \rightsquigarrow^{\mathbf{S} k} e_2}}$$

Definition 1.5.4 (Safe for k). An expression evaluates safely for k steps if it either there is some e' , such that $e \rightsquigarrow^k e'$, or there is some number of steps j strictly less than k where the term terminates with a value (i.e. there is some v and $j < k$ such that $e \rightsquigarrow^j v$).

We can now state type safety using this step-counting definition. We can't really talk about an infinite computation, but we can know that for an arbitrarily long time, e will run safely during that time.

Theorem 1.5.3 (Type Safety (step-counting)). If $\vdash e \in \tau$ then for all natural numbers k , e is safe for k .

We show this result by induction on k . If k is 0, then the result is trivial. All

expressions run safely for zero steps. If k is nonzero, then progress states that e is either a value or steps. If it is a value, we are also done, as values are safe for any k . If it steps to some e' , then preservation tells us that $\vdash e' \in \tau$. By induction, we know that e' is safe for $k - 1$. So either $e' \rightsquigarrow^j v$, i.e. e' steps to some value v within j steps, for some $j < k - 1$, or $e' \rightsquigarrow^{k-1} e''$. In the first case, we have $e \rightsquigarrow^{j+1} v$ which is a safe evaluation for e . In the second case, we have $e \rightsquigarrow^k e''$, which is also a safe evaluation for e .

1.6 Further reading

The type safety proof for the simply-typed lambda calculus is explained in a number of textbooks including TAPL [Pie02], PFPL [Har16] and Software Foundations [PdAC⁺25]. Each of these sources defines type safety as the conjunction of preservation and progress.

Milner [Mil78] proved a *type soundness* theorem, which states that well-typed ML programs cannot “go wrong”. To do so, he constructed a denotational semantics of the ML language that maps every ML program to either some mathematical value (like a number or continuous function), to a special element indicating divergence (\perp), or to a special element called “wrong” that indicates a run-time error. He then proved that if a program type checks, then its denotation does not include the “wrong” element.

Wright and Felleisen [WF94] observed that run-time errors could be ruled out by using a small-step operational semantics. They defined syntactic type soundness as showing preservation (inspired by subject reduction from combinatory logic), characterizing “stuck” or “faulty” expressions, and then showing that faulty expressions are not typeable (i.e. progress). They put these together with a strong soundness theorem that says that well-typed programs either diverge or reduce to values of the appropriate type.

2

Natural number recursion

STLC is rather *simple*. It lacks the computational power of most typed programming languages. All STLC expressions terminate! In due time, we will extend this language with arbitrary recursive definitions, which make the language Turing complete.

However, before we do that let's extend this system with a limited form of recursion. Our definition of STLC includes the *natural numbers* as constants, i.e. numbers starting from zero. Natural numbers can be defined using an *inductive datatype*. Any natural number is either zero or the successor of some natural number.

Let's redefine the syntax of natural numbers to make this structure explicit.

$$k ::= 0 \mid S k$$

Now, instead of saying 1, or 2, or 3, we could say $S 0$, or $S (S 0)$, or $S (S (S 0))$. Isn't that better? Ok, perhaps maybe not. We will keep the syntax 1, 2, 3 around for clarity, but remember that these Arabic numerals stand for this unary structure.

The advantage of working with an inductive structure of natural numbers is that they now come with an induction principle (for reasoning mathematically) and a recursion principle (for creating new definitions). This induction principle is the justification that we used in the previous section for the step-counting definition of type safety. Natural number induction is a common proof technique so we will see more of in the future!

Now that we have observed the inductive structure of natural numbers, let's incorporate this structure into our programming language so that STLC programs can work with natural numbers.

In this chapter, we add *two* new expression forms, as shown in the grammar below.

$$e ::= \dots \mid \mathbf{succ} \ e \mid \mathbf{nrec} \ e \ \mathbf{of} \ \{0 \Rightarrow e_0; S x \Rightarrow e_1\}$$

The first form, written $\mathbf{succ} \ e$, lifts the natural number successor $S k$ to be a primitive operation. Instead of only being able to take the successor of a literal

$$\begin{array}{c}
 \boxed{\Gamma \vdash e \in \tau} \qquad \qquad \qquad \text{(in context } \Gamma, \text{ term } e \text{ has type } \tau\text{)} \\
 \\
 \text{T-SUCC} \qquad \qquad \qquad \text{T-NREC} \\
 \frac{\Gamma \vdash e \in \mathbf{Nat}}{\Gamma \vdash \mathbf{succ} \ e \in \mathbf{Nat}} \qquad \qquad \qquad \frac{\Gamma \vdash e \in \mathbf{Nat} \qquad \Gamma \vdash e_0 \in \tau}{\Gamma, x:\mathbf{Nat} \vdash e_1 \in \tau \rightarrow \tau} \\
 \frac{}{\Gamma \vdash \mathbf{nrec} \ e \ \mathbf{of} \ \{0 \Rightarrow e_0; \mathbf{S} \ x \Rightarrow e_1\} \in \tau} \\
 \\
 \boxed{e \rightsquigarrow e'} \qquad \qquad \qquad \text{(term } e \text{ steps to } e'\text{)} \\
 \\
 \text{S-SUCC-LIT} \qquad \qquad \qquad \text{S-SUCC-CONG} \qquad \qquad \qquad \text{S-NREC-ZERO} \\
 \frac{}{\mathbf{succ} \ k \rightsquigarrow \mathbf{S} \ k} \qquad \qquad \qquad \frac{e \rightsquigarrow e'}{\mathbf{succ} \ e \rightsquigarrow \mathbf{succ} \ e'} \qquad \qquad \qquad \frac{}{\mathbf{nrec} \ 0 \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x \Rightarrow e_2\} \rightsquigarrow e_1} \\
 \\
 \text{S-NREC-SUCC} \\
 \frac{}{\mathbf{nrec} \ (\mathbf{S} \ k) \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x \Rightarrow e_2\} \rightsquigarrow (e_2[k/x]) \ (\mathbf{nrec} \ k \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x \Rightarrow e_2\})} \\
 \\
 \text{S-NREC-CONG} \\
 \frac{e \rightsquigarrow e'}{\mathbf{nrec} \ e \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x \Rightarrow e_2\} \rightsquigarrow \mathbf{nrec} \ e' \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x \Rightarrow e_2\}}
 \end{array}$$

Figure 2.1: Natural number operations: successor and recursion

number, with this operation, we can compute the successor of any expression that evaluates to a number.

Second, we add a *primitive recursion* operator to STLC. This extension allows us to define STLC programs by recursion. This form of recursion will always be bounded; we will be able to write interesting computations but will be able to know that all expressions terminate.

Along with these new expression forms, we also add the typing and small-step semantics rules shown in Figure 2.1.

The successor operation $\mathbf{succ} \ e$, adds one to its argument. This operation is specified by the two rules of the operational syntax that trigger when the argument is a literal value (rule **S-SUCC-LIT**) and when the argument itself steps (rule **S-SUCC-CONG**). The typing rule (rule **T-SUCC**) requires the argument to have type **Nat** and asserts that its successor is also a natural number. (Note: don't confuse the successor operation $\mathbf{succ} \ e$ of the expression language, with the syntax $\mathbf{S} \ k$ of natural numbers. The former is never a value (it steps by one of the two rules) while the latter is a way of writing a natural number).

The (primitive) recursion operation $\mathbf{nrec} \ e \ \mathbf{of} \ \{0 \Rightarrow e_0; \mathbf{S} \ x \Rightarrow e_1\}$ recurses over e . This operation compares e to see if it is 0 or some larger number. In the first case, the expression steps to e_0 . If the argument is equal to $\mathbf{S} \ k$ for some k , then the expression steps to e_1 where k replaces x . But that is not all! The rule also applies the result of this substitution to the recursive execution of the loop on k .

2.1 Examples

Let's use these operation!

double First, we define a doubling function on natural numbers with the following definition.

$$\text{double } x = \mathbf{nrec} \ x \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ y \Rightarrow \lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)\}$$

Here's how this doubling function might evaluate when given the number 2:

$$\begin{aligned} \text{double } 2 &= \mathbf{nrec} \ 2 \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ y \Rightarrow \lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)\} \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \\ &\quad (\mathbf{nrec} \ 1 \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ y \Rightarrow \lambda z. (\mathbf{succ} \ (\mathbf{succ} \ z))\}) \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \\ &\quad ((\lambda z. (\mathbf{succ} \ (\mathbf{succ} \ z))) \ (\mathbf{nrec} \ 0 \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ y \Rightarrow \lambda z. (\mathbf{succ} \ (\mathbf{succ} \ z))\})) \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \ ((\lambda z. (\mathbf{succ} \ (\mathbf{succ} \ z))) \ 0) \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \ (\mathbf{succ} \ (\mathbf{succ} \ 0)) \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \ (\mathbf{succ} \ 1) \\ &\rightsquigarrow (\lambda z. \mathbf{succ} \ (\mathbf{succ} \ z)) \ 2 \\ &\rightsquigarrow \mathbf{succ} \ (\mathbf{succ} \ 2) \\ &\rightsquigarrow \mathbf{succ} \ 3 \\ &\rightsquigarrow 4 \end{aligned}$$

Because the successor case is applied to the recursive execution of the loop, the typing rule requires that it have a function type.

Why do we specify the operation in this way? Sometimes you may see a semantics for primitive recursion that directly substitutes the result of the recursive execution in the successor case instead of indirectly doing so via application. The reason is that we want a *call-by-value* semantics for iteration. The rules should fully evaluate the recursive call on the predecessor before evaluating e_1 . Because our operational semantics for application is already call-by-value, we get this behavior automatically.

pred Note that the way that we have defined this natural number recursor through pattern matching makes it particularly simple to define a *predecessor* function:

$$\text{pred } x = \mathbf{nrec} \ x \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ y \Rightarrow \lambda z. y\}$$

This is not the case for all recursion principles. A more restricted form, sometimes called *iteration* does not bind y in the successor case.

2.2 Exercises

Proof assistants, such as Rocq, have built-in natural number types with induction and recursion principles. To become more familiar with using `nrec`, you can try defining operations in Rocq using its primitive recursion principles.

These new extensions satisfy the properties of *substitution*, *progress*, and *preservation* that we saw in the previous chapter. As an exercise, extend those proofs with appropriate new cases.

2.3 Further Reading

This chapter is adapted from Chapter 9 of Harper [Har16], with the recursor modified to for our call-by-value semantics. Harper calls this language Gödel’s System T [Gö58], which was designed to study the consistency of arithmetic. The terminology that we use for “primitive recursion” is not quite the same as the related concept from computability theory. In that context, the natural number recursor is restricted to produce functions with types of the form $\text{Nat} \rightarrow \text{Nat} \rightarrow \dots \rightarrow \text{Nat}$, i.e. functions that take any number of naturals as arguments and return a natural number. This operator does not have that restriction, and can define functions that are not usually considered “primitive recursive”.

In general, primitive recursion is not just for natural numbers. Any inductive type, such as lists or trees, can be equipped with its own primitive recursion operation (see Mendler’s dissertation [Men87]).

3

Big-Step Operational Semantics

So far, we have only considered small-step operational semantics for STLC. This semantics is useful because it provides a substitution-based, step-by-step explanation of how each expression evaluates.

But, as we have seen, our operational semantics is deterministic! Why didn't we express the semantics using an interpreter instead.

Consider the following definition in the Roq programming language. We represent the syntax of the language using the `Tm` datatype and the interpreter using the `eval` function.

```
Inductive Tm : Type :=
  var : Var -> Tm          (* variable, with names of type Var *)
  | abs : Var -> Tm -> Tm    (* abstraction *)
  | app : Tm -> Tm -> Tm    (* application *)
  | lit : nat -> Tm.          (* literal natural number constant *)

Fixpoint eval (e : Tm) : option Tm :=
  match e with
  | var x      => None
  | lit k      => Some (lit k)
  | abs x e1   => Some (abs x e1)
  | app e1 e2 =>
    match eval e1 , eval e2 with
    | Some (abs x e1') , Some v2 => eval (e1' [v2 / x])
    | _ , _ => None
  end
end.
```

The interpreter is partial because we may try to evaluate an expression with a free variable, or because we may have a type error. In either case, the interpreter returns `None`. Otherwise, when evaluating an application, the interpreter evalu-

ates the function to some abstraction value, the argument to some other value and then calls itself recursively after substituting the argument for the parameter.

But this definition is **not** accepted by Rocq. It rejects the definition of `eval` with the following error message:

```
Recursive call to eval has principal argument equal to
"e1' [v2/x]" instead of one of the following variables:
"e1" "e2".
```

The reason for this error is that Rocq assumes that we are defining this evaluation function via *structural recursion* on expressions. That means that we are allowed to call `eval` on any subterm of the argument (such as `e1` and `e2`, the subterms in the application case). However, the third recursive call is not to a subterm—instead it is to the body of the closure. Rocq cannot determine that this function terminates, so it must reject this definition.

3.1 Big-step semantics

To work around this issue in our metalogic, we can work with a *relational* version of `eval` instead. We define the inductive relation $e \Rightarrow v$ that holds when e evaluates v .

Definition 3.1.1 (Big-step semantics).

$$\boxed{e \Rightarrow v} \quad (\text{term } e \text{ big-steps to } v)$$

$$\frac{\begin{array}{c} \text{BS-VAL} \\ \hline v \Rightarrow v \end{array}}{} \quad \frac{\begin{array}{c} \text{BS-APP} \\ e_1 \Rightarrow \lambda x. e'_1 \quad e_2 \Rightarrow v_1 \\ e'_1[v_1/x] \Rightarrow v_2 \\ \hline e_1 e_2 \Rightarrow v_2 \end{array}}{}$$

Notably, this definition requires only two rules! The first rule states that values evaluate to themselves. The second evaluates an application and holds when the function evaluates to an abstraction, the argument evaluates to a value and the substitution of the argument the parameter also evaluates to a value.

Theorem 3.1.1 (Equivalence of semantics). For closed expressions e , we have $e \rightsquigarrow^* v$ if and only if $e \Rightarrow v$.

We prove each direction of this lemma separately. For the forward direction (**small-step implies big-step**), we need to show the following lemma:

Lemma 3.1.1 (Step expansion). If $e \rightsquigarrow e'$ then for all v , if $e' \Rightarrow v$ then $e \Rightarrow v$.

For the backwards direction (**big-step implies small-step**), we need to define multi-step analogues of the two evaluation rules of the big-step semantics.

Lemma 3.1.2 (`s_val`). $v \rightsquigarrow^* v$

Lemma 3.1.3 (`s_app`). If $e_1 \rightsquigarrow^* \lambda x. e'_1$ and $e_2 \rightsquigarrow^* v_1$ and $e'_1[v_1/x] \rightsquigarrow^* v_2$ then $e_1 e_2 \rightsquigarrow^* v_2$.

This lemma itself relies on showing multi-step analogues of the single-step congruence rules for the small step semantics.

Lemma 3.1.4 (`ms_app_cong1`). If $e_1 \rightsquigarrow^* e'_1$ then $e_1 e_2 \rightsquigarrow^* e'_1 e_2$.

Lemma 3.1.5 (`ms_app_cong2`). If $e_2 \rightsquigarrow^* e'_2$ then $v_1 e_2 \rightsquigarrow^* v_1 e'_2$.

3.2 Big-step semantics and type safety?

The big-step semantics has fewer rules, and in some situations, may be easier to understand as it is more directly connected to an interpreter. But, note that while the relation $e \Rightarrow v$ is deterministic, it represents a *partial function*. There are many expressions e that are not related to values.

This partiality leads to a significant drawback of a big-step semantics: it handles both the partiality of a runtime type error and the partiality of divergence in exactly the same way. This is in contrast to the small-step semantics: runtime errors made the single-step relation partial, while divergence can be modeled using a coinductive definition of multi-step reduction.

Because the big step semantics does not distinguish runtime errors from diverging programs, we run into difficulty when stating and proving type safety. While preservation holds for this semantics:

Lemma 3.2.1 (Preservation). If $e \Rightarrow v$ and $\vdash e \in \tau$ then $\vdash v \in \tau$.

There is no way to define an analogue for the progress lemma for the big step semantics.

Furthermore, it is tempting to define type safety as follows:

Conjecture 3.2.1 (Big Step Safety). If $\vdash e \in \tau$ then $e \Rightarrow v$ and $\vdash v \in \tau$.

But this is a strong lemma—it rules out both forms of partiality. We know that the program doesn't crash, but we also know that the program doesn't diverge either. While this lemma *is* true for STLC, proving this lemma is not a straightforward induction. And, if we were to extend the language to include nontermination, it would no longer be true.

3.3 Further Reading

Kahn [Kah87] initially proposed the use of *natural semantics*, which has since been referred to as *big-step* semantics to contrast with *small-step* semantics.

Leroy [Ler06] explores the ramifications of using coinductive definitions for big-step semantics.

Charguéraud develops a compromise between big-step semantics and small-step semantics that he calls *pretty-big-step semantics* [Cha13].

4

Big-step termination and Semantic soundness

In this chapter we will show how to prove the strong version of type safety for a language with a big-step semantics. After completing the proof, not only will we know that programs do not crash, we will also know that they do not diverge either. Every well-typed program can produce a value.

4.1 Big-step preservation

Before we start with type safety, let's consider the preservation and progress lemmas from the small step relation.

Lemma 4.1.1 (Preservation). If $\vdash e \in \tau$ and $e \Rightarrow v$ then $\vdash v \in \tau$.

Proof. Proof by induction on $e \Rightarrow v$.

- For the value case, we have $v \Rightarrow v$, so the result type checks by assumption.
- For the application case, e is $e_1 e_2$ where $e_1 \Rightarrow \lambda x.e'_1$ and $e_2 \Rightarrow v_2$ and $e'_1[v_2/x] \Rightarrow v$. We also know (by inversion of the evaluation relation) that $\vdash e_1 \in \tau_1 \rightarrow \tau$ and $\vdash e_2 \in \tau_1$. By induction, we know that $\vdash \lambda x.e'_1 \in \tau_1 \rightarrow \tau$ and $\vdash v_2 \in \tau_1$. By inverting the former, we also know that $x:\tau_1 \vdash e'_1 \in \tau$. This means that we can use our substitution lemma to show $\vdash e'_1[v_2/x] \in \tau$, and induction on the third evaluation to conclude $\vdash v \in \tau$.

□

The preservation lemma goes through readily. (However, note that unlike the small step relation, this proof must be done by induction on the evaluation relation, and not the typing judgement.) By itself, this lemma isn't a type safety theorem because it does not give use any confidence that the type system rules out runtime errors. The lemma is stated so that we can only make a conclusion when we

already have a successful run to a value. But, even if a term type checks, we can't tell whether it will produce anything.

The small step semantics had a progress lemma that reassured us that good things would happen if a program type checks — that program would make some incremental progress. But with this semantics we do not have a notion of incremental progress. So there isn't a ready analogue of this lemma available for us.

4.2 Big Step Safety

Instead, consider this statement of type safety for the big-step semantics. It states that any closed, well-typed term e must evaluate to some value v . This means that the term cannot crash or diverge.¹

Conjecture 4.2.1 (Big Step Safety). If $\vdash e \in \tau$ then $e \Rightarrow v$.

We won't be able to prove this theorem directly by induction on e or on the typing derivation. But, looking at how this proof fails will tell us how to strengthen it so that it is provable. We will do this strengthening incrementally, generalizing the induction hypothesis as necessary.

Our first attempt to prove the big step safety theorem is by induction on e .

- For the literal case, we know that e is a constant k (which is a value) because it is well typed, then τ is the type `Nat`. Immediately we have $k \Rightarrow k$.
- Similarly for the abstraction case, we know that e is $\lambda x.e'$ and that τ is $\tau_1 \rightarrow \tau_2$. Again, immediately, we have $\lambda x.e' \Rightarrow \lambda x.e'$.
- We don't need to consider the variable case as the term is closed.
- Finally, consider the application case, where e is $e_1 e_2$ and by inversion on the typing derivation we know that $\vdash e_1 \in \tau_1 \rightarrow \tau$ and $\vdash e_2 \in \tau_1$. We can use induction on e_1 and e_2 to get $e_1 \Rightarrow v_1$ and $e_2 \Rightarrow v_2$. Furthermore, preservation tells us that $\vdash v_1 \in \tau_1 \rightarrow \tau$, and we also know (via canonical forms) that the only closed values with function types are functions, so v_1 must be some $\lambda x.e'_1$. However, here we are stuck. We want to show that there is some v_3 such that $e'_1[v_2/x] \Rightarrow v_3$. But we don't have a way to conclude this. We can't use induction because $e'_1[v_2/x]$ isn't a subterm of e .

The solution at this point is to observe that our induction hypothesis was not strong enough. We needed to know more about e_1 than just that it terminates. We also need to know that when given any argument, *the application will also terminate*.

Therefore, let's strengthen the property that we prove. Above, in the application case we want to know something about $e'_1[v_2/x]$, so we cannot use the induction hypothesis on this term. However, we get this term from evaluating e_1 and e_2 , so we can strengthen the conclusion of our lemma. Above, all we know about v is that it is a value. But, depending on this type, we can assert stronger properties about this value.

Definition 4.2.1 (**Value set**). Define the family of sets $\mathcal{V}[\tau]$ by structural recursion on τ .

$$\begin{aligned}\mathcal{V}[\text{Nat}] &= \mathbb{N} \\ \mathcal{V}[\tau_1 \rightarrow \tau_2] &= \{\lambda x.e \mid \forall v, v \in \mathcal{V}[\tau_1] \text{ implies } \exists v', e[v/x] \Rightarrow v' \text{ and } v' \in \mathcal{V}[\tau_2]\}\end{aligned}$$

¹For simplicity, we will work with a version of type safety that is a little weaker than the version in the previous chapter; it doesn't require the result to be well typed.

This definition interprets each type as a set of values. The values in each set make sense for the associated type. The set for `Nat` contains all natural numbers, and the sets for each function type only contain abstractions. These abstractions must themselves act like they have the right type: they must take any value from the interpretation of the argument type and, after substitution, evaluate to a value in that is in the result type. It is this requirement for function types that will help us complete the proof above—we know that the substitution we need $e'[v_2/x]$ will terminate (and satisfy the condition of our stronger theorem).

This idea of evaluating a term to a value in a particular set will come up again, so let's name it. The computational interpretation of a type includes all expressions that evaluate to a value in the value set for that type.

Definition 4.2.2 (Computation set).

$$\mathcal{C}[\tau] = \{ e \mid e \Rightarrow v \text{ and } v \in \mathcal{V}[\tau] \}$$

Now we can use these definitions to strengthen our safety theorem. Not only do we require that a term evaluate to a value, but that value must be in the appropriate set for that type.

Lemma 4.2.1 (Semantic soundness). If $\vdash e \in \tau$ then $e \in \mathcal{C}[\tau]$.

However, as you might guess, our induction hypothesis is *still* not strong enough to prove the result. Let's try a case where it works, and one where it doesn't.

Let's prove the theorem by induction on $\vdash e \in \tau$. The variable case is impossible and the literal case is immediate. Now consider the application case, where we have $\vdash e_1 e_2 \in \tau$, with $\vdash e_1 \in \tau_1 \rightarrow \tau$ and $\vdash e_2 \in \tau_1$. By induction, we know that $e_1 \in \mathcal{C}[\tau_1 \rightarrow \tau]$ and $e_2 \in \mathcal{C}[\tau_1]$. So that means that both of these terms must evaluate to values in their appropriate sets. i.e. $e_1 \Rightarrow v_1$ where $v_1 \in \mathcal{V}[\tau_1 \rightarrow \tau]$ and $e_2 \Rightarrow v_2$ where $v_2 \in \mathcal{V}[\tau_1]$. To construct an evaluation in this case, we need to know that v_1 is some abstraction, and that if we substitute v_2 for the parameter the program terminates. But this is exactly what the set $\mathcal{V}[\tau_1 \rightarrow \tau]$ gives us! We have fixed this case.

However, now consider the application case. In this case we have $\vdash \lambda x. e \in \tau_1 \rightarrow \tau_2$. We need to show that this abstraction is in $\mathcal{C}[\tau_1 \rightarrow \tau_2]$. We know already that this term has the right type. It also terminates as it is already a value; i.e. $\lambda x. e \Rightarrow \lambda x. e$. Our remaining goal is to show $\lambda x. e \in \mathcal{V}[\tau_1 \rightarrow \tau_2]$. But, we cannot make any more progress. We don't have an induction hypothesis that we can use in this case, because our lemma only applies to closed terms.

Therefore, we will strengthen our lemma *again*, so that it applies not just to closed terms, but also gives us an induction hypothesis for open terms. But, our sets are only sets of closed terms! How can we do this?

The answer is that our revised lemma should quantify over *closing substitutions*, i.e. substitutions that replace all free variables with closed values. Furthermore, to make sure that our theorem is strong enough, we also require that these closed values be part of our semantic sets.

Definition 4.2.3 (Semantic substitution). Define $\sigma \in \mathcal{G}[\Gamma]$ when $\forall x \in \text{dom } \Gamma, \sigma x \in \mathcal{V}[(\Gamma x)]$.

With this definition, we can now restate our lemma in its final form.

Lemma 4.2.2 (Semantic soundness). If $\Gamma \vdash e \in \tau$ then for all $\sigma \in \mathcal{G}[\Gamma]$, $e[\sigma] \in \mathcal{C}[\tau]$.

Proof. Proof is by induction on the derivation of $\Gamma \vdash e \in \tau$.

- If $\Gamma \vdash x \in \Gamma x$, then given an arbitrary $\sigma \in \mathcal{G}[\Gamma]$, we need to show that $x[\sigma] \in \mathcal{C}[(\Gamma x)]$. However, in our definition of substitution, we have $x[\sigma] = \sigma x$, so we know that $\sigma x \in \mathcal{V}[(\Gamma x)]$ by assumption. This also implies that $\sigma x \in \mathcal{C}[(\Gamma x)]$ because values evaluate to themselves.
- If $\Gamma \vdash k \in \text{Nat}$, then given an arbitrary $\sigma \in \mathcal{G}[\Gamma]$, we need to show that $k[\sigma] \in \mathcal{C}[\text{Nat}]$. This is the same as showing that $k \Rightarrow k$ and $k \in \mathbb{N}$.
- Say $\Gamma \vdash e_1 e_2 \in \tau$ where $\Gamma \vdash e_1 \in \tau_1 \rightarrow \tau$ and $\Gamma \vdash e_2 \in \tau_1$ are subderivations. Given an arbitrary $\sigma \in \mathcal{G}[\Gamma]$, we need to show that $(e_1 e_2)[\sigma] \in \mathcal{C}[\tau]$. By induction we know that $e_1[\sigma] \in \mathcal{C}[\tau_1 \rightarrow \tau]$ and $e_2[\sigma] \in \mathcal{C}[\tau_1]$. This means that the former steps to some $v_1 \in \mathcal{V}[\tau_1 \rightarrow \tau]$ and the latter steps to some $v_2 \in \mathcal{V}[\tau_1]$. Furthermore, we know that v_1 must be some λ -term $\lambda x.e$, and that $e[v_2/x] \Rightarrow v$ with $v \in \mathcal{V}[\tau]$. But, by rule **BS-APP**, we also have $e_1[\sigma] e_2[\sigma] \Rightarrow v$, which gives us our goal.
- Say $\Gamma \vdash \lambda x.e \in \tau_1 \rightarrow \tau$ where $\Gamma, x : \tau_1 \vdash e \in \tau$. Given an arbitrary $\sigma \in \mathcal{G}[\Gamma]$, we need to show that $(\lambda x.e)[\sigma] \in \mathcal{C}[\tau]$. We know that $(\lambda x.e)[\sigma] = \lambda x.e[x/x, \sigma]$ (assuming that x is not in the domain or range of σ). It is a value, so it steps to itself. To show that this term is in the appropriate value set, we need to assume some $v_1 \in \mathcal{V}[\tau_1]$ and show that $e[x/x, \sigma][v_1/x] \in \mathcal{V}[\tau]$. Substitutions compose, so this is the same as saying $e[v_1/x, \sigma] \in \mathcal{V}[\tau]$. This goal follows by our induction hypothesis as long as $v_1/x, \sigma \in \mathcal{G}[\Gamma, x:\tau]$. But this follows by definition because $\sigma \in \mathcal{G}[\Gamma]$ and $v_1 \in \mathcal{V}[\tau]$.

□

After strengthening our lemma twice, we have put it in a form that we can prove. Furthermore, this form implies the original version, so we have proven type safety.

4.3 Rephrasing semantic soundness

Where does the name *semantic soundness* come from in this lemma? What is semantic about this argument?

The type safety proofs we have seen so far are syntactic: they involve working with the syntax of programs, typing derivations (which are syntactic objects), substitution and rewriting relations (which is a syntactic manipulation). Semantic interpretations of the λ -calculus give us a more mathematical meaning for our programs. And one idea in semantics is to give a meaning to types based on sets of values.

Furthermore, some authors like to define $\mathcal{V}[\tau]$ by first defining operations on sets.

Definition 4.3.1 (Function set). Define $T_1 \Rightarrow T_2$, which constructs a set of terms from two given sets of terms T_1 and T_2 as follows:

$$T_1 \Rightarrow T_2 = \{ \lambda x.e_2 \mid \text{for all } e_1, \quad e_1 \in T_1 \text{ implies } e_2[e_1/x] \in T_2 \}$$

With this definition, we can define our value set using this operation:

$$\begin{aligned}\mathcal{V}[\text{Nat}] &= \mathbb{N} \\ \mathcal{V}[\tau_1 \rightarrow \tau_2] &= \mathcal{V}[\tau_1] \Rightarrow \mathcal{C}[\tau_2]\end{aligned}$$

Furthermore, we can also define a notion for “semantic typing”: the idea that a term is well typed semantically when given any closing substitution, it is an element of the appropriate substitution set.

Definition 4.3.2 (Semantic typing). Define $\Gamma \models e : \tau$ when forall $\sigma \in \mathcal{G}[\Gamma]$, $e[\sigma] \in \mathcal{C}[\tau]$.

With this definition, we can prove semantic typing rules that are analogous to each syntactic typing rule. The proofs of these lemmas are the subcases of the semantic soundness proof.

Lemma 4.3.1 (Semantic var rule). $\Gamma \models x : \Gamma x$.

Lemma 4.3.2 (Semantic lit rule). $\Gamma \models k : \text{Nat}$.

Lemma 4.3.3 (Semantic abs rule). If $\Gamma, x:\tau \models e : \tau_1 \rightarrow \tau$, then $\Gamma \models \lambda x.e : \tau_1 \rightarrow \tau$.

Lemma 4.3.4 (Semantic app rule). If $\Gamma \models e_1 : \tau_1 \rightarrow \tau_2$ and $\Gamma \models e_2 : \tau_1$, then $\Gamma \models e_1 e_2 : \tau_2$.

These lemmas allow us to rephrase the semantic soundness theorem nicely.

Theorem 4.3.1 (Semantic typing). If $\Gamma \vdash e \in \tau$ then $\Gamma \models e : \tau$

Proof. Induction on the typing derivation, applying the appropriate semantic typing rule in each case. \square

4.4 Variations

This proof is usually stated with a small-step semantics and tutorials are available from a variety of references. However, with a small-step semantics you need an additional lemma: closure under reverse evaluation. Terms that step to terms in the computation sets are themselves in the computation set.

Lemma 4.4.1 (Closure under expansion). If $e \rightsquigarrow e'$ and $e' \in \mathcal{C}[\tau]$ then $e \in \mathcal{C}[\tau]$.

Not all authors make a distinction between value sets $\mathcal{V}[\tau]$ and computation sets $\mathcal{C}[\tau]$, combining them together into a uniform definition.

You might also ask whether semantic *completeness* holds.

Lemma 4.4.2 (Semantic completeness). If $\Gamma \models e : \tau$ then $\Gamma \vdash e \in \tau$.

In other words, do our sets of terms contain only well-typed terms? The answer is no! A program can include any sort of stuck subterm, as long as that subterm is never executed. For example, consider $(\lambda x.3) (\lambda y.1\ 0)$. This term doesn’t type check because it has the stuck application $1\ 0$ as a subterm. But it evaluates to 3 which is in $\mathcal{V}[\text{Nat}]$, so we can say that $\models (\lambda x.3) (\lambda y.1\ 0) : \text{Nat}$.

The reason that semantic completeness fails, is that we include ill-typed terms in our sets as long as they are harmless. But, if you believe that only typed terms are worth discussing, you might prefer to restrict the sets to well-typed terms. In that case, completeness holds trivially.

4.5 Further reading

Robert Harper's note: *How to (Re)Invent Tait's Method*<https://www.cs.cmu.edu/~rwh/courses/chtt/pdfs/tait.pdf> includes this proof for a small-step semantics of STLC with booleans, unit, products and functions. Harper calls the computation sets "hereditarily terminating."

This proof is the simplest example of a general proof technique called *proof by logical relations*. The set $\mathcal{V}[\tau]$ is a *unary* logical relation, also called a logical predicate. Binary logical relations can be used to show program equivalence or noninterference. Or, extending the sets to include open terms means that we can show that full reduction for STLC (even inside functions) always terminates. However, these proofs take a few more steps than the one that we

5

REC: Recursive definitions

In this chapter, we introduce our first *effectful* programming language: a fine-grained call-by-value language with recursive definitions, called REC. That effect is *nontermination*, where evaluating a program may not actually result in a value.

We will add nontermination through two sorts of recursive definitions: first through recursive values and then through recursive types.

5.1 Recursive definitions in CBV languages

To build intuition about the language structures we are adding in this section, we will first start out with some examples written in the OCaml programming language.

For example, a straightforward definition of the doubling function, which we previously implemented via primitive recursion, looks like this:

```
let rec double : nat -> nat = fun x ->
  match x with
  | O -> O
  | S y -> succ (succ (double y))
```

However, we are not limited to primitive recursion in OCaml. We can do much more.

5.1.1 Mutual recursion via recursive tuples

First, OCaml allows the definition of functions that are *mutually* recursive using the keywords **rec** ... **and**. Any definition in a block can refer to any other in the same block.

A simple example is the mutual definition of the odd and even functions.

```
let rec even : nat -> bool = fun x ->
  match x with
  | O -> true
  | S y -> odd y
```

```

and
odd : nat -> bool = fun x ->
match x with
| O -> false
| S y -> even y

```

But what if we didn't have **and** available? It turns out that there are two ways that we could replace the code above.

One way is to inline one recursive definition into another. This requires us to repeat the definition of one of the operations later.

```

let rec even : nat -> bool = fun x ->
match x with
| O -> true
| S y -> match y with
| O -> false
| S z -> even z

let odd : nat -> bool = fun x ->
match x with
| O -> false
| S y -> even y

```

Bekić's theorem¹ states that this always works.

Another option is to use recursion on products of functions, instead of just on functions. In this version, to make the code a little easier to read, we first define a record type containing odd and even functions (i.e. a product). Then, by using recursion through the product value, we can mutually define the two functions.

```

type oe = { odd : nat -> bool ; even : nat -> bool }

let rec obj : oe = { odd = (fun x ->
match x with
| O -> false
| S y -> obj.even y) ;
even = (fun x ->
match x with
| O -> true
| S y -> obj.odd y
)
}

```

5.1.2 Another example with recursive tuples

However, it turns out that in OCaml there are definitions of recursive product values where the recursive reference is not hidden in a function body.

Consider the following stream type:

```

type stream = Cons of int * stream

let rec zeros : stream = Cons (0, zeros)

```

¹https://en.wikipedia.org/wiki/Beki%C4%87s_theorem

Perhaps surprisingly OCaml accepts the definition of `zeros` above. This line defines `zeros` as a recursive value, represented by a cycle in the heap, and displayed as `Cons (0, <cycle>)`.

We are limited with what we can do with the `stream` type. For example, although we define a mapping operation, thus:

```
let rec map_stream f = fun y ->
  match y with
  | Cons (x, s) -> Cons (f x, map_stream f s)
```

any use of the mapping function will go into an infinite loop:

```
let ones : stream = map_stream (fun x -> x + 1) zeros
```

We cannot use the stream for termination. Instead, we need to find some other finite value to iterate over.

For example, we can access any finite number of zeros from the stream:

```
let rec take : nat -> stream -> int list * stream = fun n s ->
  match n with
  | 0 -> ([], s)
  | S m -> match s with
    | Cons (x, s') ->
      let (tl, s'') = take m s' in (x :: tl, s'')
(* let (example,_) = take n5 zeros *)
```

But there is a natural number that we could give to `take` that would cause it to diverge.

```
let rec omega : nat = S omega
let example = take omega zeros
```

5.1.3 Recursive values via recursive types

So far, we have been showing examples that use recursive value definitions as well as recursive types. But, it turns out that we only need the latter to define the former. At least for recursive functions. The trick is that we can use a recursive type allow the Y-combinator to type check in OCaml.

Let's use a simple `length` function as an example.

```
let rec length : int list -> int = fun l ->
  match l with
  | [] -> 0
  | _ :: l' -> 1 + length l'
```

If we knew how big of list we wanted to call the `length` function on, then we would not need to use recursion for this definition. For example, say we only needed it to work for lists of size 0 or 1. Then we can use the definition of `length2` below:

```
let bottom : int list -> int =
  fun _ -> failwith "<loops>"

let length0 : int list -> int = bottom
```

```
let length1 : int list -> int = fun l ->
  match l with
  | [] -> 0
  | _ :: l' -> 1 + length0 l'

let length2 : int list -> int = fun l ->
  match l with
  | [] -> 0
  | _ :: l' -> 1 + length1 l'
```

When given a longer list, this function fails. But, this is a pretty long definition. Let's refactor it:

```
let mk_length : (int list -> int) -> int list -> int =
  fun length ->
    (fun l ->
      match l with
      | [] -> 0
      | _ :: l' -> 1 + length l')

let length2 : int list -> int =
  (fun f -> (f (f bottom)))
  mk_length
```

The refactored version is easy to see how to extend it to work with longer lists.

```
let length4 : int list -> int =
  (fun f -> (f (f (f (f bottom))))))
  mk_length
```

More generally, our recursive definition is

```
let rec length : int list -> int =
  fun l ->
    mk_length length
    l
```

which could also be written as:

```
let rec length : int list -> int =
  fun l ->
    mk_length (fun l -> mk_length length l)
    l
```

Here's the big leap! If we could do self application, then maybe we would write this example without using **rec**.

```
(* DOESN'T TYPE CHECK *)
let length : int list -> int =
  fun l ->
    (fun f -> f f)
    (fun g -> mk_length (fun l -> g g l))
    l
```

But this code does not type check in OCaml.

To get it to type check, we need to introduce a recursive type so that we can apply a function to itself. The key part of the type definition is that recursive reference occurs to the left of the arrow in its own definition. This is called a negative occurrence and would be rejected by Rocq, but it is not an inductive type.

```
type 'a dom = Abs of ('a dom -> 'a)

let app (f : 'a dom) (x : 'a dom) : 'a =
  match f with
  | Abs h -> h x
```

Indeed, with this type, we can write the simplest infinite loop in the untyped lambda-calculus. This term steps immediately to itself.

```
let rec f : int -> void =
  fun x -> f x

let loop : void = f 0
```

Putting this idea to work, we can use this type get the length function above to type check.

```
let length : int list -> int =
  fun l ->
    (fun f -> app f f)
    (Abs (fun g ->
      mk_length (fun l -> app g g l)))
  l
```

We can also factor out the call-by-value Y-combinator, which is the essence of the recursive definition.

```
let y_combinator : (('a -> 'b) -> ('a -> 'b)) -> ('a -> 'b)
= fun h ->
  (fun f -> app f f)
  (Abs (fun g -> h (fun x -> (app g g) x)))
```

let length = y_combinator mk_length

5.2 A fine-grained CBV language

Now, before we consider the semantics of recursive definitions, let's refactor the base language to make it easier to extend. At the same time, we will add a few more common language features, such as products, sums and an empty type.

Here is the [syntax](#) of the base language that we will work with.

$\tau ::= \mathbf{Void} \mid \mathbf{Nat} \mid \tau_1 \rightarrow \tau_2 \mid \tau_1 * \tau_2 \mid \tau_1 + \tau_2$	types
$v ::= x$	Variables
0	zero
$\mathbf{S} v$	successor
$\lambda x. e$	Functions
(v_1, v_2)	Pairs
$\mathbf{inj}_1 v \mid \mathbf{inj}_2 v$	Sums
$e ::= v_1 v_2$	application
$\mathbf{case} v \mathbf{of} \{0 \Rightarrow e_1; \mathbf{S} x \Rightarrow e_2\}$	test for zero
$\mathbf{prj}_1 v \mid \mathbf{prj}_2 v$	projection
$\mathbf{case} v \mathbf{of} \{\mathbf{inj}_1 x \Rightarrow e_1; \mathbf{inj}_2 x \Rightarrow e_2\}$	case
$\mathbf{ret} v$	value
$\mathbf{let} x = e_1 \mathbf{in} e_2$	sequencing

The first thing to notice about this language is that it makes a syntactic distinction between values and expressions. These two grammars are mutually defined. Terms appear inside values in function bodies. Values can appear in terms in several ways. First by being “returned” via $\mathbf{ret} v$. (Some versions of fine-grained CBV make the \mathbf{ret} implicit, as it often can be inferred from context. For clarity, we will make it explicit.)

Furthermore, notice that almost all terms restrict their “active” subterms to be values. The arguments of $\mathbf{succ} v$, the scrutinee for the various forms of pattern matching, and both the function and the argument in an application must be values.

There are several reasons for this modification. The one that we will discuss now is that it allows us to extend the small-step semantics with new constructs while skipping the associated congruence rules in the operational semantics. All we need to add are the main computation rules. As the active part of the expression must be a value, we do not need to add any rules to evaluate that subexpression.

Definition 5.2.1 (Active rules).	$e \rightsquigarrow e'$	<i>(term e steps to e')</i>
S-BETA		S-CASE-ZERO
$\frac{}{(\lambda x. e) v \rightsquigarrow e[v/x]}$		$\frac{}{\mathbf{case} 0 \mathbf{of} \{0 \Rightarrow e_1; \mathbf{S} x \Rightarrow e_2\} \rightsquigarrow e_1}$
S-CASE-SUCC		S-PRJ1
$\frac{}{\mathbf{case} (\mathbf{S} k) \mathbf{of} \{0 \Rightarrow e_1; \mathbf{S} x \Rightarrow e_2\} \rightsquigarrow e_2[k/x]}$		$\frac{}{\mathbf{prj}_1(v_1, v_2) \rightsquigarrow \mathbf{ret} v_1}$
S-PRJ2	S-CASE-INJ1	
$\frac{}{\mathbf{prj}_2(v_1, v_2) \rightsquigarrow \mathbf{ret} v_2}$	$\frac{}{\mathbf{case} (\mathbf{inj}_1 v) \mathbf{of} \{\mathbf{inj}_1 x \Rightarrow e_1; \mathbf{inj}_2 x \Rightarrow e_2\} \rightsquigarrow e_1[v/x]}$	
S-CASE-INJ2		
$\frac{}{\mathbf{case} (\mathbf{inj}_2 v) \mathbf{of} \{\mathbf{inj}_1 x \Rightarrow e_1; \mathbf{inj}_2 x \Rightarrow e_2\} \rightsquigarrow e_2[v/x]}$		

In this language, the let term controls the sequencing of evaluation. We only have a single congruence rule, which evaluates right-hand side of a let expression if it is not a returned value.

$$\boxed{e \rightsquigarrow e'} \quad (\text{term } e \text{ steps to } e')$$

$$\frac{\text{S-LET-RET}}{\mathbf{let} \ x = \mathbf{ret} \ v \mathbf{in} \ e \rightsquigarrow e[v/x]} \quad \frac{\text{S-LET-CONG} \quad e_1 \rightsquigarrow e'_1}{\mathbf{let} \ x = e_1 \mathbf{in} \ e_2 \rightsquigarrow \mathbf{let} \ x = e'_1 \mathbf{in} \ e_2}$$

In this language, there is no question about the ordering in which evaluation happens. In some CBV languages, we might evaluate the argument of an application before we evaluate the function. Or we might evaluate the function before the argument. Or, the language might say that this order is undefined, giving flexibility to the language implementation. For example, the OCaml byte code compiler chooses a different order than the native code compiler.

In STLC, the order doesn't matter, and either version will produce the same result. However, in some effectful languages, the order of this evaluation can matter.

Therefore, to rule out ambiguity, this language puts the onus on the programmer to explicitly say what they want, by writing either

$$\mathbf{let} \ x_1 = e_1 \mathbf{in} \ \mathbf{let} \ x_2 = e_2 \mathbf{in} \ x_1 \ x_2 \quad \text{or} \quad \mathbf{let} \ x_2 = e_2 \mathbf{in} \ \mathbf{let} \ x_1 = e_1 \mathbf{in} \ x_1 \ x_2$$

5.2.1 Type system

With a syntactic distinction between values and expressions, we have two separate typing judgments: one for values and one for terms.

Definition 5.2.2 (Type system).

$$\boxed{\Gamma \vdash v \in \tau} \quad (\text{in context } \Gamma, \text{ value } v \text{ has type } \tau)$$

$$\begin{array}{cccc}
 \text{TV-ZERO} & \text{TV-SUCC} & \text{TV-VAR} & \text{TV-ABS} \\
 \dfrac{}{\Gamma \vdash 0 \in \mathbf{Nat}} & \dfrac{\Gamma \vdash v \in \mathbf{Nat}}{\Gamma \vdash \mathbf{S} \ v \in \mathbf{Nat}} & \dfrac{x : \tau \in \Gamma}{\Gamma \vdash x \in \tau} & \dfrac{\Gamma, x : \tau_1 \vdash e \in \tau_2}{\Gamma \vdash \lambda x. e \in \tau_1 \rightarrow \tau_2} \\
 \\
 \text{TV-PAIR} & \text{TV-INJ1} & \text{TV-INJ2} & \\
 \dfrac{\Gamma \vdash v_1 \in \tau_1 \quad \Gamma \vdash v_2 \in \tau_2}{\Gamma \vdash (v_1, v_2) \in \tau_1 * \tau_2} & \dfrac{\Gamma \vdash v_1 \in \tau_1}{\Gamma \vdash \mathbf{inj}_1 v_1 \in \tau_1 + \tau_2} & \dfrac{\Gamma \vdash v_2 \in \tau_2}{\Gamma \vdash \mathbf{inj}_2 v_2 \in \tau_1 + \tau_2}
 \end{array}$$

$$\boxed{\Gamma \vdash e \in \tau} \quad (\text{in context } \Gamma, \text{ term } e \text{ has type } \tau)$$

$$\begin{array}{ccc}
 \text{TE-RET} & \text{TE-LET} & \text{TE-APP} \\
 \dfrac{}{\Gamma \vdash v \in \tau} & \dfrac{\Gamma \vdash e_1 \in \tau_1 \quad \Gamma, x : \tau_1 \vdash e_2 \in \tau}{\Gamma \vdash \mathbf{let} \ x = e_1 \mathbf{in} \ e_2 \in \tau} & \dfrac{\Gamma \vdash v_1 \in \tau_1 \rightarrow \tau_2 \quad \Gamma \vdash v_2 \in \tau_1}{\Gamma \vdash v_1 \ v_2 \in \tau_2} \\
 \\
 \text{TE-IFZ} & & \\
 \dfrac{\Gamma \vdash v \in \mathbf{Nat} \quad \Gamma \vdash e_1 \in \tau \quad \Gamma, x_2 : \mathbf{Nat} \vdash e_2 \in \tau}{\Gamma \vdash \mathbf{case} \ v \ \mathbf{of} \ \{0 \Rightarrow e_1; \mathbf{S} \ x_2 \Rightarrow e_2\} \in \tau} & \text{TE-PRJ1} & \text{TE-PRJ2} \\
 & \dfrac{\Gamma \vdash v \in \tau_1 * \tau_2}{\Gamma \vdash \mathbf{prj}_1 v \in \tau_1} & \dfrac{\Gamma \vdash v \in \tau_1 * \tau_2}{\Gamma \vdash \mathbf{prj}_2 v \in \tau_2}
 \end{array}$$

$$\text{TE-CASE} \quad \frac{\Gamma \vdash v \in \tau_1 + \tau_2 \quad \Gamma, x_1 : \tau_1 \vdash e_1 \in \tau \quad \Gamma, x_2 : \tau_2 \vdash e_2 \in \tau}{\Gamma \vdash \text{case } v \text{ of } \{\text{inj}_1 x_1 \Rightarrow e_1; \text{inj}_2 x_2 \Rightarrow e_2\} \in \tau}$$

These two judgments are mutually defined.

5.2.2 Recovering expressiveness

Note that this restriction does not limit the expressiveness of the language. We can *define* the standard expression forms.

For example, instead of application of the form $e_1 e_2$, the only application form in this language is $v_1 v_2$.

Definition 5.2.3 (Eager let). Define $\text{let } x \Leftarrow e_1 \text{ in } e_2$ as $e_2[v/x]$ when e_1 is $\text{ret } v$ and let $x = e_1$ in e_2 otherwise.

We use this eager let in the definition of some derived forms. For example, we can create a “macro” for the usual application form for expressions, with the following definition.

Definition 5.2.4 (Extended Application). Define $e_1 e_2$ as $\text{let } x_1 \Leftarrow e_1 \text{ in let } x_2 \Leftarrow e_2 \text{ in } x_1 x_2$.

The reason for the eager let is so that we can prove that our derived forms have the appropriate operational behavior.

For example, we can prove the following lemmas about the behavior of the extended application form.

Lemma 5.2.1 (Application steps).

1. $(\text{ret } (\lambda x. e)) (\text{ret } v) \rightsquigarrow e[v/x]$.
2. If $e_1 \rightsquigarrow e'_1$ then $e_1 e_2 \rightsquigarrow^* e'_1 e'_2$.
3. If $e_2 \rightsquigarrow e'_2$ then $v e_2 \rightsquigarrow^* v e'_2$.

5.3 Recursive values

Now, let’s add general recursion. We will do so by adding a new form of value, called a *recursive value*, written $\text{rec } x. v$. Here, the variable x is bound inside v and refers to the whole value.

$$v ::= \text{rec } x. v$$

Usually v will be a function, and this will give us a way to define recursive functions. But it doesn’t have to be, as we saw in OCaml.

The typing rule includes an auxiliary judgment written $\tau \text{ ok}$ that specifies which types may be used in recursive definition.

$$\text{TV-REC} \quad \frac{\tau \text{ ok} \quad \Gamma, x : \tau \vdash v \in \tau}{\Gamma \vdash \text{rec } x. v \in \tau}$$

For now, the rules specify that function and product types are ok.

$$\frac{\text{FUN-OK}}{(\tau_1 \rightarrow \tau_2) \text{ ok}} \quad \frac{\text{PROD-OK}}{(\tau_1 * \tau_2) \text{ ok}}$$

Analogously, we need small-step rules to unwind the value when that type is used.

$$\frac{\text{S-APP-REC}}{(\mathbf{rec} \ x. v_1) \ v_2 \rightsquigarrow v_1[\mathbf{rec} \ x. v_1/x] \ v_2}$$

$$\frac{\text{S-PRJ1-REC}}{\mathbf{prj}_1(\mathbf{rec} \ x. v) \rightsquigarrow \mathbf{prj}_1 v[\mathbf{rec} \ x. v/x]} \quad \frac{\text{S-PRJ2-REC}}{\mathbf{prj}_2(\mathbf{rec} \ x. v) \rightsquigarrow \mathbf{prj}_2 v[\mathbf{rec} \ x. v/x]}$$

Why do we do it this way? Because of our syntactic separation between values and computations! We can only substitute values for values.

5.4 Recursive types

Finally we add recursive types to the language. This means adding a recursive definition form μ and type variables to the syntax of types. And, in the syntax of terms, we introduce two coercions, for introducing and eliminating values with the recursive type.

$$\begin{aligned}\tau &::= \alpha \mid \mu\alpha.\tau && \text{variables and recursive types} \\ v &::= \mathbf{fold} \ v && \text{introduction form} \\ t &::= \mathbf{unfold} \ v && \text{elimination form}\end{aligned}$$

Although we have type variables in our types, they are only used for recursion in this language. It will be an invariant of our type system that we only work with closed types. Furthermore, a recursive type $\mu\alpha.\tau$ is closed, if and only if its unfolding $\tau[\mu\alpha.\tau/\alpha]$ is also closed.

The introduction form creates a value with a recursive type and the elimination rule exposes its structure.

$$\frac{\text{TV-FOLD}}{\Gamma \vdash v \in \tau[\mu\alpha.\tau/\alpha]} \quad \frac{\text{TE-UNFOLD}}{\Gamma \vdash \mathbf{unfold} \ v \in \mu\alpha.\tau} \quad \frac{}{\Gamma \vdash \mathbf{unfold} \ v \in \tau[\mu\alpha.\tau/\alpha]}$$

The single new rule of the operational semantics removes the coercions.

$$\frac{\text{S-UNFOLD}}{\mathbf{unfold} \ (\mathbf{fold} \ v) \rightsquigarrow \mathbf{ret} \ v}$$

5.4.1 Recursive type variations

The presence of $\mathbf{fold} \ v$ and $\mathbf{unfold} \ v$ introduction and elimination forms, means that language includes *iso*-recursive types. The types $\mu\alpha.\tau$ and $\tau[\mu\alpha.\tau/\alpha]$ are isomorphic, but not equal types. The introduction and elimination terms are coercions for the isomorphism. Iso-recursive types are the most straightforward way to add recursive types to a language, as it only requires α -equivalence for type equality.

Alternatively, some languages include a definitional type equivalence, i.e. an

equivalence relation that states when two types are equal and is coarser than α -equality. In this case, the language includes some sort of conversion rule, giving terms any equivalent types. This conversion rule can make the metatheory of the language more complex to work with as often the type system is not syntax-directed.

5.5 Type safety

Let's consider type safety for REC. Notably, we will do so by proving preservation and progress with the small-step semantics. While it is possible to define a big-step semantics for this language, the fact that computations may diverge means that we have to be careful.

5.6 Further reading

Fine-grained CBV is from [LPT03] as a way to control the sequence of effects. The definition of *eager let* is from Forster et al. [FSSS19].

6

Step-indexing: Semantic type safety for REC

We proved a semantic soundness theorem for STLC. Can prove the same result for REC? Of course not. We know that we are doomed to failure as the semantic soundness theorem for STLC also proves that all programs terminate, but that is not the case for this language.

6.1 A failing proof: semantic sets

However, let's take a look at the proof because it first gives us a chance to redo the argument from before, but using a small-step semantics this time instead of a big-step semantics. It is also instructive to look at where our proof breaks down, so we can think about how to modify the statement of the theorem so that we can prove something that is true.

First, we need to define our semantic sets for our fine-grained CBV language. Our definition is similar to last time, with a few changes.

Definition 6.1.1 (Semantic sets).

$$\begin{aligned}\mathcal{C}[\tau] &= \{ e \mid e \rightsquigarrow^* \mathbf{ret} v \text{ and } v \in \mathcal{V}[\tau] \} \\ \mathcal{V}[\mathbf{Nat}] &= \mathbb{N} \\ \mathcal{V}[\mathbf{Void}] &= \{\} \\ \mathcal{V}[\tau_1 \rightarrow \tau_2] &= \{ v \mid \forall v_2, v_2 \in \mathcal{V}[\tau_1] \text{ implies } v v_2 \in \mathcal{C}[\tau_2] \} \\ \mathcal{V}[\tau_1 * \tau_2] &= \{ v \mid \mathbf{prj}_1 v \in \mathcal{C}[\tau_1] \text{ and } \mathbf{prj}_2 v \in \mathcal{C}[\tau_2] \} \\ \mathcal{V}[\tau_1 + \tau_2] &= \{ \mathbf{inj}_1 v \mid v_1 \in \mathcal{V}[\tau_1] \} \cup \{ \mathbf{inj}_2 v \mid v_2 \in \mathcal{V}[\tau_2] \}\end{aligned}$$

First, we define the sets for computations and values mutually. That will allow us to use the definition of computation sets directly in the definition of the value sets, so we can localize our description of the evaluation.

Next, we have a few more types around, so we need more cases in our definition of the value sets. We now have cases for the empty type (there are no values of this type, so its semantics is an empty set), product types and sum types. These cases are in addition to the natural number and function type cases from last time.

One case that is missing from the definition is the case for recursive types. We might like to add a definition like the following: a “folded” value is in the set for a recursive type if the underlying value is in the set for the unfolded type.

$$\mathcal{V}[\mu\alpha.\tau] = \{\text{fold } v \mid v \in \mathcal{V}[\tau[\mu\alpha.\tau/\alpha]]\}$$

However, we **cannot** add this case to our definition. The reason is that with this case, the definition is no longer well-founded. Above, the semantic sets are defined by recursion over the type structure. This means that each recursive part of the definition must be to a set for a smaller type. However, the type $\tau[\mu\alpha.\tau/\alpha]$ is not necessarily smaller than $\mu\alpha.\tau$. So we cannot use $\mathcal{V}[\tau[\mu\alpha.\tau/\alpha]]$ to define $\mathcal{V}[\mu\alpha.\tau]$.

So to define these sets recursively, we need to recur on something else if we want to include recursive types.

The second difference in this definition is in the case for function types. We want to allow function values to be either explicit lambda expressions or recursive functions. By defining the set in terms of the elimination form for function types instead of the introduction form, we can implicitly include $\text{rec } x.v$ as a value in this set, as long as it behaves like a recursive function.

Consider what would happen if we tried to do it the other way around, and add an alternative for recursively defined values to our prior definition for lambda terms. We cannot do this definition either because our definition is not well-founded: in the first line we define $\mathcal{V}[\tau_1 \rightarrow \tau_2]$ in terms of $\mathcal{V}[\tau_1 \rightarrow \tau_2]$.

$$\begin{aligned} \mathcal{V}[\tau_1 \rightarrow \tau_2] &= \{ \text{rec } x.v \mid v[\text{rec } x.v/x] \in \mathcal{V}[\tau_1 \rightarrow \tau_2] \} \cup \\ &\quad \{ \lambda x.e \mid \forall v2, v_2 \in \mathcal{V}[\tau_1] \text{ implies } e[v_2/x] \in \mathcal{C}[\tau_2] \} \end{aligned}$$

6.1.1 Semantic lemmas with small-step evaluation

Continuing our (doomed) proof attempt, we can restate some of the definitions that were present in the previous proof. We need to know when a substitution only includes values in the semantic set, and can use this to define a semantic form of the value and computation typing relations.

Definition 6.1.2 (Semantic substitution and typing). 1. Define $\sigma \in \mathcal{G}[\Gamma]$ when $\forall x \in \text{dom } \Gamma, \sigma x \in \mathcal{V}[\Gamma x]$.

2. Define $\Gamma \models e : \tau$ when for all $\sigma \in \mathcal{G}[\Gamma], e[\sigma] \in \mathcal{C}[\tau]$.

3. Define $\Gamma \models v : \tau$ when for all $\sigma \in \mathcal{G}[\Gamma], v[\sigma] \in \mathcal{V}[\tau]$.

Because we are working with a small-step semantics, we will also need the following property.

Lemma 6.1.1 (Reverse evaluation). If $e \rightsquigarrow e'$ and $e' \in \mathcal{C}[\tau]$ then $e \in \mathcal{C}[\tau]$.

Proof. The proof is by definition—we just extend the multi-step evaluation one more step. \square

Now let’s think about which of our semantic typing rules are true. Some go

through very similarly to their analogs for the STLC (and the big-step semantics). And in some places the division between value typing rules and computation typing rules cleans things up nicely. No more need to reason about how values evaluate to themselves.

The base cases are straightforward.

Lemma 6.1.2 (Semantic var rule). $\Gamma \models x : \Gamma x$.

Lemma 6.1.3 (Semantic zero rule). $\Gamma \models 0 : \text{Nat}$.

Let's look that the case for (non-recursive) functions.

Lemma 6.1.4 (Semantic abs rule). If $\Gamma, x : \tau \models e : \tau_1 \rightarrow \tau_2$, then $\Gamma \models \lambda x. e : \tau_1 \rightarrow \tau_2$.

Proof. Let $\sigma \in \mathcal{G}[\Gamma]$ be arbitrary. We want to prove that for any $v_1 \in \mathcal{V}[\tau_1]$, we have $(\lambda x. e)[\sigma] v_1 \in \mathcal{C}[\tau_2]$. Using our assumption, we know that $e[v_1/x, \sigma] \in \mathcal{C}[\tau_2]$. Because $\mathcal{C}[\tau_2]$ is closed under reverse evaluation, and $(\lambda x. e)[\sigma] v_1 \rightsquigarrow e[v_1/x, \sigma]$ we are done. \square

Lemma 6.1.5 (Semantic app rule). If $\Gamma \models v_1 : \tau_1 \rightarrow \tau_2$ and $\Gamma \models v_2 : \tau_1$, then $\Gamma \models v_1 v_2 : \tau_2$.

Proof. Let $\sigma \in \mathcal{G}[\Gamma]$ be arbitrary. We want to prove that for any $v_1 v_2[\sigma] \in \mathcal{C}[\tau_2]$. By our assumptions, we know that $v_1[\sigma] \in \mathcal{V}[\tau_1 \rightarrow \tau_2]$ and $v_2[\sigma] \in \mathcal{V}[\tau_1]$. By definition, we know our desired result. \square

We have isolated all of our computation to let expressions.

Lemma 6.1.6 (Semantic let rule). If $\Gamma \models e_1 : \tau_1$ and $\Gamma, x : \tau_1 \models e_2 : \tau_2$, then $\Gamma \models \text{let } x = e_1 \text{ in } e_2 : \tau_2$.

Proof. Let $\sigma \in \mathcal{G}[\Gamma]$ be arbitrary. We want to prove that for any $(\text{let } x = e_1 \text{ in } e_2)[\sigma] \in \mathcal{C}[\tau_2]$. By our assumptions we know that $e_1[\sigma] \in \mathcal{C}[\tau_1]$. This means that there is some $e_1[\sigma] \rightsquigarrow^* v_1$ and $v_1 \in \mathcal{V}[\tau_1]$. Furthermore by assumption, we know that $e_2[v_1/x, \sigma] \in \mathcal{C}[\tau_2]$. We can put these evaluation sequences together to have $\text{let } x = e_1 \text{ in } e_2 \rightsquigarrow^* \text{let } x = \text{ret } v_1 \text{ in } e_2 \rightsquigarrow e_2[v_1/x, \sigma]$. And, because the computation set is closed under reverse evaluation, we are done. \square

6.1.2 A problem

Our typing rule for recursive values looks like this.

Lemma 6.1.7 (Semantic rec rule). If $\Gamma, x : \tau \models v : \tau$ then $\Gamma \models \text{rec } x. v : \tau$.

Let's consider the case only when τ is a function type, i.e. $\tau = \tau_1 \rightarrow \tau_2$.

Let $\sigma \in \mathcal{G}[\Gamma]$ be arbitrary. We want to show that $\text{rec } x. v[\sigma] \in \mathcal{V}[(\tau_1 \rightarrow \tau_2)]$. Let $v_1 \in \mathcal{V}[\tau_1]$ be arbitrary. We now want to show that

$$\text{rec } x. v[\sigma] v_1 \in \mathcal{C}[\tau_2]$$

By closure under reverse evaluation, this is the same as showing that

$$v[\text{rec } x. v[\sigma]/x, \sigma] v_1 \in \mathcal{C}[\tau_2]$$

By assumption, we can conclude this as long as we have

$$\text{rec } x. v[\sigma]/x, \sigma \in \mathcal{G}[\Gamma, x : \tau_1 \rightarrow \tau_2]$$

This requires showing that

$$\text{rec } x. v[\sigma] \in \mathcal{V}[\tau_1 \rightarrow \tau_2]$$

But now we are stuck! This is *exactly* the result that we were already trying to show.

6.2 Fixing the problem: counting evaluation steps

In the previous section, we saw that there were two problems with the approach. The first is that if we want to include recursive types, we cannot define the value relation by recursion over type structure. The semantics of recursive types are defined in terms of themselves, not in terms of smaller components.

The second issue is that when we are trying to prove that recursive values are in the relation, because they do a self-substitution, we need to know that they are already in the relation to show that they are in the relation.

Now, if our semantic sets were coinductively defined, we would solve both problems. We could define the set in terms of itself (allowing recursive types) and to show that an element was a member of a set, we could use coinduction. But, we cannot do that with a straightforward definition. The case for function types features a negative occurrence of the value set, so the naive definition cannot be used in a coinductive relation.

Instead, we are going to define these sets a little less directly.

Recall this definition from before, which says that safe programs do not get stuck. We use **irreducible** to refer to programs that cannot step.

Definition 6.2.1 (Safe program). A closed program is e is **safe** if, for any e' such that $e \rightsquigarrow^* e'$, if e' **irreducible** then e' is **ret** v for some v .

One way we showed type safety was to decompose this definition into steps.

Definition 6.2.2 (Safe for k). A program e is safe for k steps if the following conditions hold. If e **irreducible**, then e must be some **ret** v . Otherwise, if it steps $e \rightsquigarrow e'$ and if $0 < k$ then e' must run safely for $k - 1$ steps.

Definition 6.2.3 (Runs safely). A program **runs safely** if for all k , it is safe for k steps.

The definition of safe for k is inductively defined over k . It holds for all programs that do not get stuck within k steps. By requiring that this predicate hold for all k , we can define type safety without requiring programs to terminate. Diverging programs satisfy this relation because they do not get stuck after any number of steps of evaluation.

Lemma 6.2.1 (Programs that run safely are safe). If e runs safely, then it is safe.

Proof. This proof is by induction on the number of steps of evaluation of $e \rightsquigarrow^* e'$. If there are zero steps, because we know that the program runs safely for 0 steps, then we know it is safe. If the evaluation is $e \rightsquigarrow e_1$ and $e_1 \rightsquigarrow^* e'$, then, by induction we know that e_1 is safe, which gives us e is safe. But, we have to show that e' runs safely, which follows because e runs safely. \square

Another way to say this is to think about sets.

Let S_k be the set of all programs that are safe for k steps. That means that we have an infinite chain of decreasing sets.

$$S_0 \supseteq S_1 \supseteq S_2 \supseteq S_3 \dots$$

This sequence is *downward closed*: for each $i \leq j$ we have $S_j \subseteq S_i$. In other words, if any program is in some set S_j then that program is in all previous sets in the sequence.

Furthermore, the set of safe programs is the intersection of all of the sets in this sequence.

$$S = \bigcap_i S_i$$

Aside: Coinduction This is very similar to a definition of safety that uses coinduction. Another way to set up this definition is to consider this function from sets of terms to sets of terms.

$$FX = \{e \text{ irreducible implies } e = \text{ret } v \text{ and } e \rightsquigarrow e' \text{ implies } e' \in X\}$$

Say $F^i X$ is i applications of F to X . Then we have:

$$S = \bigcap_i F^i \mathcal{U}$$

The function F is monotonic: $X \subseteq Y$ implies that $FX \subseteq FY$, so we can prove a coinduction principle for this definition. This principle states that if we want to prove that some element is in S , then we need to show that it is in some set S_1 , such that $S_1 \subseteq F(S_1)$.

For example, say we have some term $(\text{rec } f.\lambda x.f\ x) 3$ and we want to show that this term runs safely. Intuitively we know that this is the case because we have the evaluation sequence:

$$(\text{rec } f.\lambda x.f\ x) 3 \rightsquigarrow (\lambda x.(\text{rec } f.\lambda x.f\ x)) 3 \rightsquigarrow (\text{rec } f.\lambda x.f\ x) 3 \rightsquigarrow \dots$$

Then we can pick S_1 to be the set containing just this term and its unfolding. Our proof obligation is to show that $S_1 \in FS_1$, i.e. that both terms step to an element of S_1 .

6.3 Step-indexed semantic safety

Let's use this idea to define our logical relation as a downward-closed step-indexed sequence of sets.

As before, we'll split our logical relation into two parts, mutually defining sets of values and sets of terms (computations). What is different this time is that these sets are also indexed by a step-count. So our relations are now three place relations, between types, values or terms, and step counts. This time, we write $\mathcal{V}[\![v \in \tau]\!]_k$ and $\mathcal{C}[\![e \in \tau]\!]_k$, when v and e are in their respective relations at step index k .

Furthermore, we will define these sets so that all terms in the computation set are safe.

Lemma 6.3.1 (Semantic safety). If for any k , we have $\mathcal{C}[\![e \in \tau]\!]_k$, then e is safe.

Our fundamental lemma will show that closed, well-typed terms $\vdash e \in \tau$ are in $\mathcal{C}[\![e \in \tau]\!]_k$ for any k . This will give us (the hard way) a type safety theorem for statically well-typed terms.

While this result is a bit underwhelming (the preservation/progress based theorem is a lot easier) it will give us a framework that we can extend to prove more interesting theorems.

6.3.1 Step-indexed propositions

Sometimes, we want to work with our relations as functions from steps to propositions. For example, if we write $\mathcal{V}[\![v \in \tau]\!]$, without the argument k , we are talking about a function. There are two meta-operations on step-indexed propositions that we can use to clarify our reasoning.

The first operation is called the *later modality*, and takes a step-indexed proposition and produces a new step-indexed proposition. This operation accesses its argument at the previous step count. If the step is zero, then it holds trivially.

Definition 6.3.1 (Later modality). Define $\triangleright_k \phi$ by case analysis on k .

- $\triangleright_0 \phi$ always holds
- $\triangleright_{(S k)} \phi$ holds iff ϕ_k .

The next operation is called *stepped implication*, and takes *two* step-indexed propositions to produce a new step-indexed proposition.

Definition 6.3.2 (Step-indexed implication). Define $\phi \implies_k \phi'$ as

$$\text{for all } j < k, \phi_j \text{ implies } \phi'_j$$

Definition 6.3.3 (Downward closed). A step-indexed proposition is *downward closed* when $j \leq k$ we have $\phi_k \text{ implies } \phi_j$.

Suppose ϕ and ϕ' are both downward closed. The step-indexed proposition, $\phi_k \text{ implies } \phi'_k$ is not downward closed. Assume $j \leq k$, and that $\phi_k \text{ implies } \phi'_k$. We want to show that $\phi_j \text{ implies } \phi'_j$. So assume ϕ_j . We want to show ϕ'_j . If we could show ϕ'_k , we would be done because ϕ' is downward closed. Furthermore, by assumption, it also suffices to show ϕ_k . In other words, we need $\phi_j \text{ implies } \phi_k$. But here we are stuck! We cannot use the fact that ϕ is downward closed here because that gives us the opposition relation between ϕ_j and ϕ_k .

Lemma 6.3.2 (Step-indexed implication is downward closed). For any ϕ and ϕ' , $\phi \implies \phi'$ is downward closed.

Proof. Assume that $j \leq k$ and $\phi \implies_k \phi'$. We want to show that $\phi \implies_j \phi'$. Unfolding definitions, assume we have some $i < j$ and we want to show that $\phi_i \text{ implies } \phi'_i$. But, we are done, because by transitivity, we have $i < k$, so we can use our original assumption. \square

6.3.2 Logical relation

Now let's define our the logical relation by *well-founded recursion* on the step-count k .

Definition 6.3.4 (Step-indexed Logical Relation).

$$\begin{aligned}
 \mathcal{C}[\![e \in \tau]\!]_k &= e \text{ irreducible implies that there exists } v \text{ such that} \\
 &\quad e = \mathbf{ret} v \text{ and } \mathcal{V}[\![v \in \tau]\!]_k \\
 &\quad \text{and } e \rightsquigarrow e' \text{ implies } \triangleright_k \mathcal{C}[\![e' \in \tau]\!]
 \end{aligned}$$

$$\begin{aligned}
 \mathcal{V}[\![v \in \mathbf{Void}]\!]_k &= \text{never} \\
 \mathcal{V}[\![v \in \mathbf{Nat}]\!]_k &= v \in \mathbb{N} \\
 \mathcal{V}[\![\lambda x. e \in \tau_1 \rightarrow \tau_2]\!]_k &= \forall v_1, \mathcal{V}[\![v_1 \in \tau_1]\!] \implies_k \mathcal{C}[\![e[v_1/x] \in \tau_2]\!] \\
 \mathcal{V}[\![\mathbf{rec} x. v \in \tau_1 \rightarrow \tau_2]\!]_k &= \triangleright_k \mathcal{V}[\![v[\mathbf{rec} x. v/x] \in \tau_1 \rightarrow \tau_2]\!] \\
 \mathcal{V}[\![(v_1, v_2) \in \tau_1 * \tau_2]\!]_k &= \triangleright_k \mathcal{V}[\![v_1 \in \tau_1]\!] \text{ and } \triangleright_k \mathcal{V}[\![v_2 \in \tau_2]\!] \\
 \mathcal{V}[\![\mathbf{rec} x. v \in \tau_1 * \tau_2]\!]_k &= \triangleright_k \mathcal{V}[\![v[\mathbf{rec} x. v/x] \in \tau_1 * \tau_2]\!] \\
 \mathcal{V}[\![\mathbf{inj}_1 v_1 \in \tau_1 + \tau_2]\!]_k &= \triangleright_k \mathcal{V}[\![v_1 \in \tau_1]\!] \\
 \mathcal{V}[\![\mathbf{inj}_2 v_2 \in \tau_1 + \tau_2]\!]_k &= \triangleright_k \mathcal{V}[\![v_2 \in \tau_2]\!] \\
 \mathcal{V}[\![\mathbf{fold} v \in \mu\alpha.\tau]\!]_k &= \triangleright_k \mathcal{V}[\![v \in \tau[\mu\alpha.\tau/\alpha]]\!]
 \end{aligned}$$

This definition is well-defined because recursive occurrences of $\mathcal{C}[\![e \in \tau]\!]$ and $\mathcal{V}[\![v \in \tau]\!]$ are done with steps that are strictly smaller than k . (There is one exception: definition of $\mathcal{C}[\![e \in \tau]\!]_k$ refers to $\mathcal{V}[\![v \in \tau]\!]_k$ at the same index. However, $\mathcal{V}[\![v \in \tau]\!]_k$ makes all of its recursive calls to smaller indices so we never get back to $\mathcal{C}[\![e \in \tau]\!]_k$.)

The way this works is through the use of the two auxiliary operations defined above: if ϕ is a step-indexed set, then $\triangleright_{\phi k}$ accesses ϕ at the predecessor of k . The use of this operation in $\mathcal{C}[\![e \in \tau]\!]_k$ means that after taking a step, the next term must be in some set with a smaller index.

For the case of product types, we include all pairs, where both components are in the previous set, and all recursive values, where the unrolling is in the previous set.

We do a similar thing for functions, including explicitly lambda terms and recursive values. However, the lambda case uses the step-indexed implication defined above. This makes sure that we will be able to prove that $\mathcal{C}[\![e \in \tau]\!]$ and $\mathcal{V}[\![v \in \tau]\!]$ are downward closed.

Lemma 6.3.3 (Downward closed).

1. If $j \leq k$ then $\mathcal{V}[\![v \in \tau]\!]_k$ implies $\mathcal{V}[\![v \in \tau]\!]_j$.
2. If $j \leq k$ then $\mathcal{C}[\![e \in \tau]\!]_k$ implies $\mathcal{C}[\![e \in \tau]\!]_j$.

Proof. We will prove this by strong induction on k . Let $j \leftarrow k$ be arbitrary, and suppose . We want to show that for any v , $\mathcal{V}[\![v \in \tau]\!]_k$ implies $\mathcal{V}[\![v \in \tau]\!]_j$ and for any e , $\mathcal{C}[\![e \in \tau]\!]_k$ implies $\mathcal{C}[\![e \in \tau]\!]_j$. Suppose $\mathcal{V}[\![v \in \tau]\!]_k$. Let's consider the cases for τ .

- τ cannot be α because we only work with closed types.
- If τ is **Void**, we are done as the sets are empty.
- If τ is **Nat**, we are also done as the sets are all natural numbers.
- If τ is $\tau_1 * \tau_2$, then we must have either
 1. v as (v_1, v_2) , and $\triangleright_k \mathcal{V}[\![v_1 \in \tau_1]\!]$ and $\triangleright_k \mathcal{V}[\![v_2 \in \tau_2]\!]$. If $k = 0$ then j is zero, so we are done. If $k > 0$ then we have $\mathcal{V}[\![v_1 \in \tau_1]\!]_{(k-1)}$ and $\mathcal{V}[\![v_2 \in \tau_2]\!]_{(k-1)}$. In this case we can use our induction hypothesis, to show that $\mathcal{V}[\![v_1 \in \tau_1]\!]_j$ and $\mathcal{V}[\![v_2 \in \tau_2]\!]_j$.

2. v as $\text{rec } x.v'$ and $\triangleright_k \mathcal{V}[(v'[\text{rec } x.v'/x]) \in \tau_1 * \tau_2]$. If $k = 0$ then j is zero, so we are done. If $k > 0$, then the result follows via induction, as above.
- If τ is $\tau_1 \rightarrow \tau_2$, then we must have either
 1. v as $\lambda x.e$, and for all v_1 , $\mathcal{V}[v_1 \in \tau_1] \implies_k \mathcal{C}[e[v_1/x] \in \tau_2]$. Stepped-implication is downward closed, so we also know $\mathcal{V}[v_1 \in \tau_1] \implies_j \mathcal{C}[e[v_1/x] \in \tau_2]$
 2. v as $\text{rec } x.v'$ and $\triangleright_k \mathcal{V}[(v'[\text{rec } x.v'/x]) \in \tau_1 \rightarrow \tau_2]$. If $k = 0$ then j is zero, so we are done. If $k > 0$, then the result follows via induction, as above.

□

Examples Let's consider some example terms

$\mathcal{V}[v \in \tau]_0$ includes:

All values

$\mathcal{C}[e \in \tau]_0$ includes:

All expressions, except those that are stuck, e.g. 0 0

$\mathcal{V}[v \in \tau]_1$ includes:

All values, except functions with stuck bodies e.g. $\lambda x.0\ 0$ or $\lambda x.0\ x$

And functions that use their arguments in the wrong way *immediately*.

e.g. $\lambda x.x\ 0$ and $\lambda x.\text{let } z = x \text{ in } i_n(\lambda y.y)\ z$ are not in this set when τ is $\text{Nat} \rightarrow \tau_2$, but these terms are in this set when τ is $(\text{Nat} \rightarrow \tau_1) \rightarrow \tau_2$.

Furthermore $\lambda x.(\lambda y.(x\ 0)\ 1)$ is in this set at any type, but it is not immediately stuck.

$\mathcal{C}[e \in \tau]_1$ includes:

All $e = \text{ret } v$ where v is in $\mathcal{V}[v \in \tau]_1$

plus $e \rightsquigarrow e'$ and e' is not stuck.

$\mathcal{V}[v \in \tau]_2$ includes:

All values, except functions with stuck bodies e.g. $\lambda x.0\ 0$ or $\lambda x.0\ x$

or functions that step to stuck bodies $\lambda x.(\lambda y.0\ 0)\ 0$.

And functions that use their arguments in the wrong way, immediately, or *after one step*.

e.g. $\lambda x.x\ 0$ and $\lambda x.(\lambda y.(x\ 0)\ 1)$ are not in this set when τ is $\text{Nat} \rightarrow \tau_2$.

$\mathcal{C}[e \in \tau]_2$ includes:

All $e = \text{ret } v$ where $\mathcal{V}[v \in \tau]_2$

plus $e \rightsquigarrow \text{ret } v$ where $\mathcal{V}[v \in \tau]_1$

plus $e \rightsquigarrow e' \rightsquigarrow e''$ where e'' is not stuck.

Computation sets include safe terms Our goal is to show that all terms are in the computation set are safe, so let's prove that now.

Lemma 6.3.4 (Semantic safety). If for any k , we have $\mathcal{C}[e \in \tau]_k$, then e is safe.

Proof. We want to show that if $e \rightsquigarrow^* e'$ and e' irreducible then e' is some $\text{ret } v$. We will prove this by induction on the evaluation $e \rightsquigarrow^* e'$.

- If $e = e'$, then e is irreducible, so by the definition of $\mathcal{C}[e \in \tau]_0$, we know that it must be some value.
- If $e \rightsquigarrow e_1$ and $e_1 \rightsquigarrow^* e'$, then we can prove this result by induction, assuming

that we can show that for all k , $\mathcal{C}[\![e_1 \in \tau]\!]_k$. So let k be arbitrary. By assumption, we have $\mathcal{C}[\![e \in \tau]\!]_{(S^k)}$. Because $e \rightsquigarrow e_1$, this means that $\triangleright_{(S^k)} \mathcal{C}[\![e_1 \in \tau]\!]$, which means that $\mathcal{C}[\![e_1 \in \tau]\!]_k$.

□

6.3.3 Semantic typing and semantic typing lemmas

Continuing on, we find that our semantic substitution and semantic typing propositions are also now step-indexed. Furthermore, to make sure that the typing relations are downward closed, we use step-indexed implication.

Definition 6.3.5 (Semantic substitution and typing). 1. Define $[\![\sigma \in \Gamma]\!]_k$ when for all $x \in \text{dom } \Gamma$, we have $\mathcal{V}[\![\sigma x \in \Gamma x]\!]_k$.

2. Define $\Gamma \vDash k e : \tau$ when $[\![\sigma \in \Gamma]\!] \implies_k \mathcal{C}[\![e[\sigma] \in \tau]\!]$.

3. Define $\Gamma \vDash_k v \in \tau$ when $[\![\sigma \in \Gamma]\!] \implies_k \mathcal{V}[\![v[\sigma] \in \tau]\!]$.

Lemma 6.3.5 (Semantic var rule). forall k , we have $\Gamma \vDash_k x \in \Gamma x$.

Proof. We assume some $k, j < k$, and $[\![\sigma \in \Gamma]\!]_j$. We want to show $\mathcal{V}[\![x[\sigma] \in (\Gamma x)]!]_j$. But this is true by unfolding definitions. □

Lemma 6.3.6 (Semantic zero rule). forall k , $\Gamma \vDash_k 0 \in \mathbf{Nat}$.

We assume some $k, j < k$, and $[\![\sigma \in \Gamma]\!]_j$. We want to show $\mathcal{V}[\![0[\sigma] \in \mathbf{Nat}]\!]_j$. But this is again true by unfolding definitions.

Lemma 6.3.7 (Semantic succ rule). forall k , if $\Gamma \vDash_k v \in \mathbf{Nat}$ then $\Gamma \vDash_k S v \in \mathbf{Nat}$.

We assume some $k, j < k$, and $[\![\sigma \in \Gamma]\!]_j$. We want to show $\mathcal{V}[\!(S v)[\sigma] \in \mathbf{Nat}]\!]_j$. Our induction hypothesis gives us $\mathcal{V}[\![v[\sigma] \in \mathbf{Nat}]\!]_j$, which means that $v[\sigma]$ must be a natural number. So its successor must be a natural number.

Lemma 6.3.8 (Semantic abs rule). forall k , if $\Gamma, x : \tau \vDash k e : \tau_1 \rightarrow \tau_2$, then $\Gamma \vDash k \lambda x. e : \tau_1 \rightarrow \tau_2$.

Proof. We assume some $k, j < k$, and $[\![\sigma \in \Gamma]\!]_j$. We want to prove that $\mathcal{V}[\!(\lambda x. e)[\sigma] \in \tau_1 \rightarrow \tau_2]\!]_j$, which can be restated as for any $i < j$, $\mathcal{V}[\![v \in \tau_1]\!]_i$ implies $\mathcal{C}[\![e[v/x, \sigma] \in \tau_2]\!]_i$. Using our assumption, we know that because $i < k$, if $[\![v/x, \sigma \in \Gamma, x : \tau_1]\!]_i$ then $\mathcal{C}[\![e[v/x, \sigma] \in \tau_2]\!]_i$. So it suffices to show that $\mathcal{V}[\![v \in \tau_1]\!]_i$ and $[\![\sigma \in \Gamma]\!]_i$. However, we assumed the former, and the latter holds because $[\![\sigma \in \Gamma]\!]$ is downward closed. □

Before we continue further, what if our definition for function types had instead been in terms of application, instead of substitution. i.e. we combined both cases above into the single case:

$$\mathcal{V}[\![v \in \tau_1 \rightarrow \tau_2]\!]_k = \forall v_1, \mathcal{V}[\![v_1 \in \tau_1]\!] \implies_k \mathcal{C}[\![v_1 v_2 \in \tau_2]\!]$$

We assume some $k, j < k$, and $[\![\sigma \in \Gamma]\!]_j$. We want to prove that $\mathcal{V}[\!(\lambda x. e)[\sigma] \in \tau_1 \rightarrow \tau_2]\!]_j$, which can be restated as for any $i < j$, $\mathcal{V}[\![v_1 \in \tau_1]\!]_i$ implies $\mathcal{C}[\!(\lambda x. e)[\sigma] v_1 \in \tau_2]\!]_i$. Using our assumption, we know that for any $i' < k$, if $[\![v/x, \sigma \in \Gamma, x : \tau_1]\!]_{i'}$ then $\mathcal{C}[\![e[v/x, \sigma] \in \tau_2]\!]_{i'}$. Our logical relation is backwards closed, but it takes a step. We need $\mathcal{C}[\![e[v/x, \sigma] \in \tau_2]\!]_{(S_i)}$ to show $\mathcal{C}[\!(\lambda x. e)[\sigma] v_1 \in \tau_2]\!]$. By instantiating

i' with S_i , it suffices to show that $\mathcal{V}[v \in \tau_1]_{(S_i)}$ and $[\sigma \in \Gamma]_{(S_i)}$. However, here we are stuck. We know these results about i but not i' .

Now let's do the case where we were stuck before.

Lemma 6.3.9 (Semantic rec rule). For all k , if $\Gamma, x : \tau \models_k v \in \tau$ then $\Gamma \models_k \text{rec } x.v \in \tau$.

Proof. Let's again consider the case only when τ is a function type, i.e. $\tau = \tau_1 \rightarrow \tau_2$. We will prove this by strong induction on k . We can assume that the lemma holds for all indices less than k , and we want to show it for k .

Now, assume some $j < k$, and $[\sigma \in \Gamma]_j$. We want to prove that $\mathcal{V}[(\text{rec } x.v)[\sigma]] \in \tau_1 \rightarrow \tau_2]_j$, which can be restated as $\triangleright_j \mathcal{V}[v[\text{rec } x.v/x, \sigma] \in \tau_1 \rightarrow \tau_2]$. We can also assume that $j = S_i$, as this is trivial otherwise. So we want to show $\mathcal{V}[v[\text{rec } x.v/x, \sigma] \in \tau_1 \rightarrow \tau_2]_i$. This result holds by instantiating our assumption with i and the substitution $\text{rec } x.v/x, \sigma$. We also need to show that $[\text{rec } x.v/x, \sigma \in \Gamma, x : \tau_1 \rightarrow \tau_2]_i$. This holds because $[\sigma \in \Gamma]$ is downward closed. So we really only need to show $\Gamma \models_i \text{rec } x.v \in \tau$. But as $i < k$ we can use our induction hypothesis, and reduce this to $\Gamma, x : \tau \models_i v \in \tau$. However, semantic typing is also downward closed, and we assumed this judgment at k , so we can also have it at i . \square

Compared to where we got stuck before, we needed the semantic typing of the recursive value, but only at a smaller index, which was available via induction.

Lemma 6.3.10 (Semantic app rule). For all k , if $\Gamma \models k v_1 : \tau_1 \rightarrow \tau_2$ and $\Gamma \models_k v_2 \in \tau_1$, then $\Gamma \models k v_1 v_2 : \tau_2$.

To prove this lemma, we can assume some $j < k$, and $[\sigma \in \Gamma]_j$. That gives us a goal of $\mathcal{C}[v_1 v_2 \in \tau_2]_k$. We can prove this result with a sub-lemma, that applies after instantiating our premises with the assumptions.

Lemma 6.3.11 (Computation set application). For all k , if $\mathcal{V}[v_1 \in \tau_1 \rightarrow \tau_2]_k$ and $\mathcal{V}[v_2 \in \tau_1]_k$, then $\mathcal{C}[v_1 v_2 \in \tau_2]_k$.

Proof. We will prove this by strong induction on k .

We want to prove that $\mathcal{C}[v_1 v_2 \in \tau_2]_k$. We have two cases for $\mathcal{V}[v_1 \in \tau_1 \rightarrow \tau_2]_k$:

1. If v_1 is $\text{rec } x.v$, then we know that $\triangleright_k \mathcal{V}[v[\text{rec } x.v/x] \in \tau_1 \rightarrow \tau_2]$. We want to show two things: If e irreducible then there exists some v such that $e = \text{ret } v$ and $\mathcal{V}[v \in \tau]_k$ and if $e \rightsquigarrow e'$ then $\triangleright_k \mathcal{C}[e' \in \tau_2]$. The first case does not apply because $(\text{rec } x.v) v_2$ is always reducible, as it steps to $v[\text{rec } x.v/x] v_2$. So we focus on the second, where we need to show that $\triangleright_k \mathcal{C}[v[\text{rec } x.v/x] v_2 \in \tau_2]$. Here we can assume that $k = Sj$ as the zero case is trivial. So we need to show that $\mathcal{C}[v[\text{rec } x.v/x] v_2 \in \tau_2]_j$. As j is less than k , we can use our induction hypothesis, with our initial assumption about v_1 and the downward closure of $\mathcal{C}[v_2 \in \tau_2]$.
2. If v_1 is $\lambda x.e$, then we know that $\forall v_1, \mathcal{V}[v_2 \in \tau_1] \implies_k \mathcal{C}[e[v_2/x] \in \tau_2]$. Again we want to show two things: If e irreducible then there exists some v such that $e = \text{ret } v$ and $\mathcal{V}[v \in \tau]_k$ and if $e \rightsquigarrow e'$ then $\triangleright_k \mathcal{C}[e' \in \tau_2]$. And again we have $(\lambda x.e) v_2 \rightsquigarrow e[v_2/x]$ so the first case doesn't apply. So we need to show $\triangleright_k \mathcal{C}[e[v_2/x] \in \tau_2]$. Again we assume that $k = Sj$, and reduce this to $\mathcal{C}[e[v_2/x] \in \tau_2]_j$. By our assumption about v_1 , as j is strictly smaller than k , we only need to show $\mathcal{V}[v_2 \in \tau_1]_j$. This holds by downward closure.

\square

Finally, let's prove a rule specific to fine-grained CBV. In particular, we have isolated all sequencing to let terms, so here is where we will need to consider how our computation set is closed under evaluation.

Lemma 6.3.12 (Semantic let rule). For all k , if $\Gamma \vdash k e_1 : \tau_1$ and $\Gamma, x : \tau_1 \vdash k e_2 : \tau_2$, then $\Gamma \vdash k \text{ let } x = e_1 \text{ in } e_2 : \tau_2$.

Like the application rule, we show this result with the help of a sub-lemma. However, the details of this proof have been left as an exercise.

Lemma 6.3.13. For all k , if $\mathcal{C}[\![e_1 \in \tau_1]\!]_k$ and $\forall v, \mathcal{V}[\![v \in \tau_1]\!] \implies_k \mathcal{C}[\![\text{let } x = \text{ret } v \text{ in } e_2 \in \tau_2]\!]_k$ then $\mathcal{C}[\!(\text{let } x = e_1 \text{ in } e_2)\!] \in \tau_2$.

Putting the above lemmas together, we can show that all well-typed terms and values are semantically sound.

Lemma 6.3.14 (Semantic Soundness).

- If $\Gamma \vdash e \in \tau$ then for all k , $\Gamma \vdash k e : \tau$

- If $\Gamma \vdash v \in \tau$ then for all k , $\Gamma \vdash_k v \in \tau$

6.4 Further Reading

Step-indexed logical relations were invented by Appel and McAllester [AM01] and developed by Ahmed [Ahm04].

This methodology forms the foundation of the Iris logic. Timany et al [TKDB24] describes how to prove this and more sophisticated results using Iris, while hiding the step-counts altogether using a special purpose logic.

The approach taken in this section was inspired by Xavier Leroy's lectures on step-indexed logical relations, from the lecture series titled "Programming = proving? The Curry-Howard correspondence today". Slides available from <https://xavierleroy.org/CdF/2018-2019/8.pdf>.

Additional reading about the foundations of inductive and coinductive definitions can be found in Ron Garcia's lecture notes available from <https://www.cs.ubc.ca/~rxg/cpsc509-spring-2022/06-coinduction.pdf>

7

Type and Effect systems

The REC language introduced the idea of a computational effect: *nontermination*.

However, this effect is invisible to the type system. As a result, we had to take a pessimistic view of well-typed REC terms: any of them could diverge at any time.

In this chapter, we will refine our type system so that its static analysis can tell us more about a program than just the form of its resulting value (if any). In particular, we will annotate our typing judgment with an effect modality ε that describes the potential *effects* of computation (if any). When tracking non-termination, the effect annotation can either be \perp (indicating that we know the code terminates) or DIV , indicating that it may diverge. However, the structure of the type-and-effect system that we present is more general than that, and could be extended to track other forms of effects.

7.1 Core system

For concreteness, we continue to work with fine-grained CBV REC language, reusing its syntax of terms and values and its small-step operational semantics. The key updates are all in the type system: we slightly modify the syntax of types and we present a completely new pair of typing judgements for values and computations.

This subsection starts with the design of the core system, ignoring (for now) features that introduce potential nontermination. In section 7.2 we extend this discussion to include recursive values and recursive types. These rules are an example of a *type-and-effect* system and they are, in fact, general about the specific effects that are tracked by the type system. We discuss what that means in section 7.5.

Definition 7.1.1 (Effect-annotated types).

$$\begin{array}{lcl} \text{effect flags} & \varepsilon & ::= \perp \mid \text{DIV} \\ \text{types} & \tau & ::= \text{Void} \mid \text{Nat} \mid \tau_1 \xrightarrow{\varepsilon} \tau_2 \mid \tau_1 \times \tau_2 \mid \tau_1 + \tau_2 \end{array}$$

In this type-and-effect system, the type syntax includes effect annotations. These annotations ε occur on the types of values that can be recursive (i.e. function types

and product types).

Non-recursive products **can** use \perp for their annotation, indicating projection has no effect. However some non-recursive functions must use DIV if they could diverge when applied, for example, if their body contains an infinite loop. We say that the flags on this type describe the *latent* effect of the function: these effects happen not when the function is created but could occur when the function is used.

Definition 7.1.2 (Type system).

$$\begin{array}{c}
 \boxed{\Gamma \vdash v \in \tau} \quad \text{(in context } \Gamma, \text{ value } v \text{ has type } \tau\text{)} \\
 \begin{array}{cccc}
 \text{TV-ZERO} & \text{TV-SUCC} & \text{TV-VAR} & \text{TV-ABS-EFF} \\
 \frac{}{\Gamma \vdash 0 \in \mathbf{Nat}} & \frac{\Gamma \vdash v \in \mathbf{Nat}}{\Gamma \vdash \mathbf{S} v \in \mathbf{Nat}} & \frac{x : \tau \in \Gamma}{\Gamma \vdash x \in \tau} & \frac{\Gamma, x : \tau_1 \vdash e \stackrel{\varepsilon}{\in} \tau_2}{\Gamma \vdash \lambda x. e \in \tau_1 \xrightarrow{\varepsilon} \tau_2}
 \end{array} \\
 \begin{array}{ccc}
 \text{TV-PAIR-EFF} & & \text{TV-INJ1} \\
 \frac{\Gamma \vdash v_1 \in \tau_1 \quad \Gamma \vdash v_2 \in \tau_2}{\Gamma \vdash (v_1, v_2) \in \tau_1 \times \tau_2} & & \frac{\Gamma \vdash v_1 \in \tau_1}{\Gamma \vdash \mathbf{inj}_1 v_1 \in \tau_1 + \tau_2} \\
 & & \text{TV-INJ2} \\
 & & \frac{\Gamma \vdash v_2 \in \tau_2}{\Gamma \vdash \mathbf{inj}_2 v_2 \in \tau_1 + \tau_2}
 \end{array} \\
 \boxed{\Gamma \vdash e \stackrel{\varepsilon}{\in} \tau} \quad \text{(in context } \Gamma, \text{ term } e \text{ has type } \tau \text{ with effect } \varepsilon\text{)} \\
 \begin{array}{ccc}
 \text{TEE-LET} & & \text{TEE-APP} \\
 \frac{\text{TEE-RET} \quad \Gamma \vdash v \in \tau}{\Gamma \vdash \mathbf{ret} v \stackrel{\perp}{\in} \tau} & \frac{\Gamma \vdash e_1 \stackrel{\varepsilon_1}{\in} \tau_1 \quad \Gamma, x : \tau_1 \vdash e_2 \stackrel{\varepsilon_2}{\in} \tau}{\Gamma \vdash \mathbf{let} x = e_1 \mathbf{in} e_2 \stackrel{\varepsilon_1 \oplus \varepsilon_2}{\in} \tau} & \frac{\Gamma \vdash v_1 \in \tau_1 \xrightarrow{\varepsilon} \tau_2 \quad \Gamma \vdash v_2 \in \tau_1}{\Gamma \vdash v_1 v_2 \stackrel{\varepsilon}{\in} \tau_2}
 \end{array} \\
 \text{TEE-IFZ} \\
 \frac{\Gamma \vdash v \in \mathbf{Nat} \quad \Gamma \vdash e_1 \stackrel{\varepsilon}{\in} \tau \quad \Gamma, x_2 : \mathbf{Nat} \vdash e_2 \stackrel{\varepsilon}{\in} \tau}{\Gamma \vdash \mathbf{case} v \mathbf{of} \{0 \Rightarrow e_1; \mathbf{S} x_2 \Rightarrow e_2\} \stackrel{\varepsilon}{\in} \tau} \quad \begin{array}{ccc}
 \text{TEE-PRJ1} & & \text{TEE-PRJ2} \\
 \frac{\Gamma \vdash v \in \tau_1 \times \tau_2}{\Gamma \vdash \mathbf{prj}_1 v \stackrel{\varepsilon}{\in} \tau_1} & & \frac{\Gamma \vdash v \in \tau_1 \times \tau_2}{\Gamma \vdash \mathbf{prj}_2 v \stackrel{\varepsilon}{\in} \tau_2}
 \end{array} \\
 \text{TEE-CASE} \\
 \frac{\Gamma \vdash v \in \tau_1 + \tau_2 \quad \Gamma, x_1 : \tau_1 \vdash e_1 \stackrel{\varepsilon}{\in} \tau \quad \Gamma, x_2 : \tau_2 \vdash e_2 \stackrel{\varepsilon}{\in} \tau}{\Gamma \vdash \mathbf{case} v \mathbf{of} \{\mathbf{inj}_1 x_1 \Rightarrow e_1; \mathbf{inj}_2 x_2 \Rightarrow e_2\} \stackrel{\varepsilon}{\in} \tau} \quad \begin{array}{c}
 \text{TEE-SUB-EFF} \\
 \frac{\Gamma \vdash e \stackrel{\varepsilon_1}{\in} \tau \quad \varepsilon_1 <: \varepsilon_2}{\Gamma \vdash e \stackrel{\varepsilon_2}{\in} \tau}
 \end{array}
 \end{array}$$

Values do not have effects, so we do not need to change their typing rules, except for the introduction form for functions and products.

In the effect annotated computation rules, the judgment form $\Gamma \vdash e \stackrel{\varepsilon}{\in} \tau$ uses ε to indicate whether the computation e is pure (i.e. terminates) or may diverge. Returned values must be pure, so rule **TEE-RET** uses \perp for this annotation. However, a let binding sequences the evaluation of two computations e_1 and e_2 . Therefore, the effect of the whole computation can be computed from the effect of these two

subterms, using the operation $\varepsilon_1 \oplus \varepsilon_2$. For nontermination, this computation returns DIV if either effect is DIV and \perp otherwise.

In the application rule, the effect of the computation is completely determined by the latent effect of the function. If the function could diverge, for example if it is $\lambda x.\text{loop}$ or $\lambda x.\text{case } x \text{ of } \{0 \Rightarrow \text{loop}; S y \Rightarrow y\}$, then its type will include the DIV effect. On the other hand, if the function terminates on *all* arguments, then the type uses the \perp effect. If the function is higher-order, i.e. if it takes another function as an argument, then the effect on the type of that argument determines the effect on the whole function. For example, we can give the function $\lambda f.f(f)$ the type $(\text{Nat} \xrightarrow{\perp} \text{Nat}) \xrightarrow{\perp} \text{Nat}$ or the type $(\text{Nat} \xrightarrow{\text{DIV}} \text{Nat}) \xrightarrow{\text{DIV}} \text{Nat}$.

The two projections $\text{prj}_1 v$ and $\text{prj}_2 v$ use the effects from the product types. Because this language includes recursive values, it is possible for projection to cause divergence. For example, the computation

$$\text{prj}_1(\text{rec } x.(x, x)) \rightsquigarrow \text{prj}_1(\text{rec } x.(x, x), \text{rec } x.(x, x)) \rightsquigarrow \dots$$

does not terminate.

The two pattern-matching elimination forms require that the effects for both branches be the same (just as the types of the both branches must be the same). During type-checking time, if the scrutinee is a variable, we don't know which branch will execute. So the effects of the computation could either be the effects of the first branch or the effects of the second.

What if these two branches don't have the same effect, such as $\text{case } x \text{ of } \{0 \Rightarrow \text{loop}; S y \Rightarrow y\}$ from above? In this case, we include a *subsumption rule*, rule TEE-SUB-EFF which allows the type system to weaken its analysis. Every effect system has a (pre)-order on effects. For nontermination, we have $\perp <: \text{DIV}$, meaning that if we know that a system always terminates, then it is sound to weaken that analysis to say that it might diverge.

With this subeffecting rule, this type system becomes our first example of a system that is not syntax directed. In every other system that we have encountered so far, we have had a 1-1 correspondence between typing rules and syntax forms. This gives us strong inversion properties: If we have a derivation, we know exactly what order the rules must be applied for the system to type check.¹

7.2 Recursive values and types

The DIV effect is explicitly introduced to the type system through the typing rules for recursive values and types. A recursively-defined function or product **must** use DIV in its type to indicate that it may diverge. This restriction is governed by the $\tau \text{ ok}$ judgement.

Definition 7.2.1. $\boxed{\tau \text{ ok}}$

(Values of type τ can be recursive)

OK-FUN-EFF

OK-PROD-EFF

$$\frac{}{(\tau_1 \xrightarrow{\text{DIV}} \tau_2) \text{ ok}}$$

$$\frac{}{(\tau_1 \times \tau_2) \text{ ok}}$$

¹Remember that even though a type system is syntax-directed, it does not need to have unique types. All of our type systems allow programs to be given multiple distinct types for some functions.

Definition 7.2.2 (Type system).

$$\begin{array}{c}
 \boxed{\Gamma \vdash v \in \tau} \quad (\text{in context } \Gamma, \text{ value } v \text{ has type } \tau) \\
 \begin{array}{c}
 \text{TV-REC} \\
 \frac{\tau \text{ ok} \quad \Gamma, x:\tau \vdash v \in \tau}{\Gamma \vdash \mathbf{rec} \ x.v \in \tau}
 \end{array}
 \quad
 \begin{array}{c}
 \text{TV-FOLD} \\
 \frac{\Gamma \vdash v \in \tau[\mu\alpha.\tau/\alpha]}{\Gamma \vdash \mathbf{fold} \ v \in \mu\alpha.\tau}
 \end{array} \\
 \\[1em]
 \boxed{\Gamma \vdash e \stackrel{\varepsilon}{\in} \tau} \quad (\text{in context } \Gamma, \text{ term } e \text{ has type } \tau \text{ with effect } \varepsilon) \\
 \begin{array}{c}
 \text{TEE-UNFOLD} \\
 \frac{\Gamma \vdash v \in \mu\alpha.\tau}{\Gamma \vdash \mathbf{unfold} \ v \stackrel{\text{DIV}}{\in} \tau[\mu\alpha.\tau/\alpha]}
 \end{array}
 \end{array}$$

Recall that the introduction rule for recursive values has a precondition that states what types of values can be recursive. To make sure that the infinite loops that these values can create is accurately accounted for, we require function and product types to be marked with DIV when they are assigned to recursive definitions. This annotation is used when these values are eliminated: the application and project rules need to know whether the value could loop when it is in the active position.

Similarly, unfolding a value with a recursive type also triggers the DIV effect. (We pessimistically assume that all values with recursive types could cause non-termination, so there is no need to mark this in the recursive type itself.)

7.3 Syntactic metatheory: Preservation and progress

Because this typing rules for computations are not syntax directed, we cannot easily use inversion. If we have a derivation $\Gamma \vdash e \stackrel{\varepsilon}{\in} \tau$, the rule **SUB-EFF** rule could always have been used, so it is difficult to say more about specific derivations. For example, the following lemma was trivially true in the REC type system. In this system, we must prove it by induction.

Lemma 7.3.1 (Ret inversion). If $\Gamma \vdash \mathbf{ret} \ v \stackrel{\varepsilon}{\in} \tau$ then $\Gamma \vdash v \in \tau$.

Proof. We prove this by induction on the typing derivation. There are only two possible cases: rule **TV-RET** and rule **TV-SUB**. In the first case, the result is immediate. In the second case, we must appeal to the induction hypothesis. \square

In our type safety proof, the place where we commonly use inversion in the progress lemma: we need to know the form of a value given its type. Sometime we explicitly state these inversions as canonical forms lemmas. In this setting, only the computation typing rule includes a non-syntax directed rule. In our value typing rule, there is exactly one rule for each syntactic form of value.

7.4 Effect soundness

To know whether the effect annotations in our type system are meaningful, we need to prove more than type safety — we need to argue that if a term has type $\vdash e \in \tau$ then it must terminate. As before, we can prove this result using a logical relation.

As we are only concerned with terminating programs, we don't need to use a step indexed logical relation. Instead, we can take advantage of the simpler definition by structural recursion on the type structure. In the case of recursive types, we use the empty set because our type system does not allow values with recursive types to be used in computations marked \perp .

The key difference in this definition is that the computation set is also indexed by the effect.

$$\begin{aligned} \mathcal{C}[\tau]^\perp &= \{ e \mid e \rightsquigarrow^* v \text{ and } v \in \mathcal{V}[\tau] \} \\ \mathcal{C}[\tau]^{\text{DIV}} &= \text{all terms} \end{aligned}$$

$$\begin{aligned} \mathcal{V}[\text{Void}] &= \emptyset \\ \mathcal{V}[\text{Nat}] &= \mathbb{N} \\ \mathcal{V}[\tau_1 \xrightarrow{\varepsilon} \tau_2] &= \{ v \mid \forall v_1, v_1 \in \mathcal{V}[\tau_1] \text{ implies } v v_1 \in \mathcal{C}[\tau_2]^\varepsilon \} \\ \mathcal{V}[\tau_1 \times \tau_2] &= \{ v \mid \mathbf{prj}_1 v \in \mathcal{C}[\tau_1]^\varepsilon \text{ and } \mathbf{prj}_2 v \in \mathcal{C}[\tau_2]^\varepsilon \} \\ \mathcal{V}[\mu\alpha.\tau] &= \{ \mathbf{fold} v \mid v \in \text{all values} \} \end{aligned}$$

Note that the computation set is compatible with the effect ordering. This property is important to reason about our subeffecting rule. As this rule is not syntax directed, we need to make sure that the set for the larger effect contains all of the same terms.

Lemma 7.4.1 (Subeffecting). If $\varepsilon_1 <: \varepsilon_2$ then $\mathcal{C}[\tau]^{\varepsilon_1} \subseteq \mathcal{C}[\tau]^{\varepsilon_2}$

As usual, we can prove the fundamental property of this logical relation by validating each of the semantic typing rules. This lemma says that the relation contains all well typed values and terms. (For brevity, we don't list those lemmas here.)

Lemma 7.4.2 (Fundamental lemma). If $\Gamma \vdash e \in \tau$ then $\Gamma \models e \in \tau$. If $\Gamma \vdash v \in \tau$ then $\Gamma \models v : \tau$.

Using the fundamental lemma above, and the definition of $\mathcal{C}[\tau]^\perp$, we can show that the termination effect implies termination.

Lemma 7.4.3 (Effect soundness). If $\vdash e \in \tau$, then there exists some v , such that $e \rightsquigarrow^* v$.

Note that our definition of $\mathcal{C}[\tau]^{\text{DIV}}$ does not tell us *anything* about terms with the DIV effect. We could strengthen this result, by combining it with the step-indexed semantic soundness property of the previous chapter. But, observe that we are able to prove termination for the \perp fragment without doing so.

7.5 Generalizing effects

This type system is specialized to tracking non-termination. However, if we view these rules more abstractly if we consider effects to be an abstract structure.

Looking at the rules, we need the following operations from our abstract structure:

ε	some type of effect annotation
\perp	annotation for “pure” code
$\varepsilon_1 \oplus \varepsilon_2$	combination of effects
$\varepsilon_1 <: \varepsilon_2$	ordering of effects

It will turn out, that to show that our type-and-effect system has the rules that we want, we will also want to assume some properties about this structure. In particular, we will want to make sure that this structure is a *pre-ordered monoid*.

Our structure is a monoid if it satisfies the following three properties:

1. Left identity: $\perp \oplus \varepsilon = \varepsilon$
2. Right identity: $\varepsilon \oplus \perp = \varepsilon$
3. Associativity: $(\varepsilon_1 \oplus \varepsilon_2) \oplus \varepsilon_3 = \varepsilon_1 \oplus (\varepsilon_2 \oplus \varepsilon_3)$

It is a preorder if $<:$ is reflexive and transitive. In other words, if we have $\varepsilon <: \varepsilon$ and $\varepsilon_1 <: \varepsilon_2$ and $\varepsilon_2 <: \varepsilon_3$ implies $\varepsilon_1 <: \varepsilon_3$.

Finally, it is a preordered monoid if the ordering is compatible with the combining operation. In other words, we have $\varepsilon_1 <: \varepsilon'_1$ and $\varepsilon_2 <: \varepsilon'_2$ then $\varepsilon_1 \oplus \varepsilon_2 <: \varepsilon'_1 \oplus \varepsilon'_2$.

It turns out that the specific structure that we use here, with the definition of ε as \perp or DIV, the ordering $\perp <: \text{DIV}$ and the combining operation described above, satisfies these properties. This structure satisfies other properties as well, such as commutativity of $\varepsilon_1 \oplus \varepsilon_2$, but we will find that we won’t need to use that property in our proofs.

To modify this type system to track other forms of effects, we need to modify the effect structure that we use accordingly.

7.6 Expressivity

The type system shown in this chapter is designed to show the essence of how type-and-effect systems can work. As a result, it is rather simplistic and not expressive enough for practical programmer.

Strengthening the termination analysis: recursive definitions The type system that we have defined does not include a very large fragment that can be checked with effect \perp . The reason is the innate pessimism of general recursion. There is no way to conclude that a recursive definition actually terminates.

One way to make the language more expressive is to include more capabilities for terminating recursive definitions. For example, we could add the primitive natural number recursion operation from Chapter 2. Type systems that enforce termination (such as the type theories of Rocq, Agda or Lean, include many such recursion principles for inductive datatypes.

Subtype polymorphism Many languages with effect systems also support subtype polymorphism. What this means in this context is the addition of two new rules that allow the type of a value or term to be replaced by any of its supertypes.

$$\begin{array}{c}
 \text{TEE-SUB-EFF-TY} \\
 \Gamma \vdash e \in \tau_1 \quad \varepsilon_1 <: \varepsilon_2 \quad \tau_1 <: \tau_2 \\
 \frac{\text{TV-SUB}}{\Gamma \vdash v \in \tau_2} \qquad \frac{\Gamma \vdash e \in \tau_1 \quad \varepsilon_1 <: \varepsilon_2 \quad \tau_1 <: \tau_2}{\Gamma \vdash e \in \tau_2}
 \end{array}$$

In essence, the subsumption rules weaken what we know: for example, we may forget that a function will always terminate when applied.

The subtyping relation is induced by the effect annotations function and product types. For example, a function annotated by \perp must always terminate, whereas a function annotated by DIV may terminate or may diverge. Similarly, projection from a tuple annotated by \perp will always terminate, but from a tuple annotated by DIV could diverge.

Definition 7.6.1 (Subtyping).

$$\boxed{\tau_1 <: \tau_2} \qquad (\tau_1 \text{ is a subtype of } \tau_2)$$

$$\begin{array}{c}
 \text{S-ARR} \qquad \qquad \qquad \text{S-PROD} \qquad \qquad \qquad \text{S-SUM} \\
 \frac{\tau_2 <: \tau_1 \quad \varepsilon_1 <: \varepsilon_2 \quad \tau'_1 <: \tau'_2}{\tau_1 \xrightarrow{\varepsilon_1} \tau'_1 <: \tau_2 \xrightarrow{\varepsilon_2} \tau'_2} \quad \frac{\tau_1 <: \tau_2 \quad \varepsilon_1 <: \varepsilon_2 \quad \tau'_1 <: \tau'_2}{\tau_1 \times \tau'_1 <: \tau_2 \times \tau'_2} \quad \frac{\tau_1 <: \tau_2 \quad \tau'_1 <: \tau'_2}{\tau_1 + \tau'_1 <: \tau_2 + \tau'_2}
 \end{array}$$

On top of this relationship, the rules for subtyping above extend the primitive subtyping relationship compatibly through the structure of types. Following standard practice, subtyping is *covariant* (i.e. in the same direction) in the components of pairs and sums and the result type of functions. It is *contravariant* in the argument of functions.

For example, we can show the following relationship between these two types

$$(\tau_1 \xrightarrow{\text{DIV}} \tau_2) \perp \rightarrow (\tau_1 \perp \rightarrow \tau_2) <: (\tau_1 \perp \rightarrow \tau_2) \perp \rightarrow (\tau_1 \xrightarrow{\text{DIV}} \tau_2)$$

(Both of these types can be assigned to the higher-order function $\lambda x.\text{ret}(\lambda y.y)$).

Effect polymorphism Although terms of this type system do not include any type or effect annotations, type *inference* for this type system should be decidable.²

However, although an inference algorithm exists, it is not a *compositional* algorithm. We cannot break a program into parts and type check each part individually; we need to know how definitions are used to figure out what type their types should be.

Most efficient type inference systems are based on the idea of *principal types*. This means that for any term in a given context, we can assign a unique “best” type for that term, i.e., one that can be used in any typing derivation featuring that term.

Principal typing fails for this language. To see why, note that we can give the term $\lambda x.\text{ret}(\lambda y.x y)$ *many* different types of the form: $(\tau_1 \xrightarrow{\varepsilon_1} \tau_2) \xrightarrow{\varepsilon_1} (\tau_1 \xrightarrow{\varepsilon_3} \tau_2)$.

There are eight possible combinations of these annotations, and six of them are valid types for the term. (Options (5) and (7) do not type check. Only configura-

²This is a conjecture, we don't have a proof yet.

tions where $\varepsilon_1 <: \varepsilon_3$ are allowed.)

	ε_1	ε_2	ε_3
(1)	\perp	\perp	\perp
(2)	\perp	\perp	DIV
(3)	\perp	DIV	\perp
(4)	\perp	DIV	DIV
<i>Doesn't type check</i>	(5)	DIV	\perp
	(6)	DIV	DIV
<i>Doesn't type check</i>	(7)	DIV	\perp
	(8)	DIV	DIV

If we look at the subtyping relationships between these types, we can arrange them into the following hierarchy, where subtypes are to the left and super types are to the right.

$$\begin{array}{ccc} & (2) & \\ (6) & <: & \\ & (8) & (4) \\ (1) & <: & \\ & (3) & \end{array}$$

Type (4), i.e., $(\tau_1 \xrightarrow{\perp} \tau_2) \xrightarrow{\text{DIV}} (\tau_1 \xrightarrow{\text{DIV}} \tau_2)$ is the maximum type—all other valid types are a subtype of this type. However, there is no *minimum* type. Both types (1) and (6) are subtypes of the rest of the types, but they are not comparable to each other.

The reason is that both of these types are instances of a more general pattern:

$$(\tau_1 \xrightarrow{\perp} \tau_2) \xrightarrow{\perp} (\tau_1 \xrightarrow{\perp} \tau_2) \quad \text{and} \quad (\tau_1 \xrightarrow{\text{DIV}} \tau_2) \xrightarrow{\perp} (\tau_1 \xrightarrow{\text{DIV}} \tau_2)$$

The effect of the output function (ε_3) depends on the effect of the input function (ε_1). We can capture this relationship using *effect polymorphism*, which would allow us to assign the following type to the function.

$$\forall \epsilon. (\tau_1 \xrightarrow{\epsilon} \tau_2) \xrightarrow{\perp} (\tau_1 \xrightarrow{\epsilon} \tau_2)$$

Note that effect polymorphism is not the only reason that this type system lacks principal types. To really develop a compositional system, we also need to include *type polymorphism*, as is found in ML or Haskell. Therefore, the best type of the term can be defined as:

$$\forall \alpha_1. \forall \alpha_2. \forall \epsilon. (\alpha_1 \xrightarrow{\epsilon} \alpha_2) \xrightarrow{\perp} (\alpha_1 \xrightarrow{\epsilon} \alpha_2)$$

At the same time, the more that we add to the language, the more difficult type inference becomes. When we add effect and type polymorphism, we need to extend any hypothetical type inference algorithm to include those features. The former, especially with this minimal (two point) effect structure is not difficult to accommodate. However, for the latter, we must be careful and stick to restrictions such as *prenex polymorphism* [Mil78] that are known to be well-behaved, if we would like to infer all types.

Value-dependent or path-dependent effects The type-and-effect analysis approximates runtime behavior when it comes to branching. The rule for case analysis

rule **TEE-IFZ** requires that after testing a natural number, both branches produce the same type of value and have the same effect. This rule is designed this way because static type checking does not evaluate programs, but instead must explore all paths.

However, there may be some relationship between the value of natural number and the effects of each branch. A dependent type system, can express how the effect of a conditional may be determined by this value.

7.7 Variation: Syntax-directed effects

The subsumption rule makes the type-and-effect system not syntax directed. However, we can reformulate the system so that this rule is not necessary. Instead, subeffecting is worked into the other rules.

Definition 7.7.1 (Syntax-directed Type system).

$\Gamma \vdash_{SD} e \stackrel{\varepsilon}{\in} \tau$	<i>(in context Γ, term e has type τ with effect ε)</i>
SD-TEE-LET	
$\frac{\begin{array}{c} \Gamma \vdash_{SD} e_1 \stackrel{\varepsilon_1}{\in} \tau_1 \\ \Gamma, x:\tau_1 \vdash_{SD} e_2 \stackrel{\varepsilon_2}{\in} \tau \\ \varepsilon_1 \oplus \varepsilon_2 <: \varepsilon \end{array}}{\Gamma \vdash_{SD} \text{let } x = e_1 \text{ in } e_2 \stackrel{\varepsilon}{\in} \tau}$	$\frac{\begin{array}{c} \Gamma \vdash_{SD} v_1 \in \tau_1 \xrightarrow{\varepsilon_1} \tau_2 \\ \Gamma \vdash_{SD} v_2 \in \tau_1 \\ \varepsilon_1 <: \varepsilon \end{array}}{\Gamma \vdash_{SD} v_1 v_2 \stackrel{\varepsilon}{\in} \tau_2}$
	$\frac{\begin{array}{c} \Gamma \vdash_{SD} v \in \tau_1 \times \tau_2 \\ \varepsilon_1 <: \varepsilon \end{array}}{\Gamma \vdash_{SD} \text{prj}_1 v \stackrel{\varepsilon}{\in} \tau_1}$
SD-TEE-PRJ2	
	$\frac{\begin{array}{c} \Gamma \vdash_{SD} v \in \tau_1 \times \tau_2 \\ \varepsilon_1 <: \varepsilon \end{array}}{\Gamma \vdash_{SD} \text{prj}_2 v \stackrel{\varepsilon}{\in} \tau_2}$

By looking at the rules, we can see that sub-effecting is an *admissible* rule for the system. Even though we don't have a specific rule, we can always turn a derivation into another with a larger effect.

Lemma 7.7.1 (Subeffecting for syntax-directed system). If $\Gamma \vdash_{SD} e \stackrel{\varepsilon_1}{\in} \tau$ and $\varepsilon_1 <: \varepsilon_2$ then $\Gamma \vdash_{SD} e \stackrel{\varepsilon_2}{\in} \tau$.

Using the subeffecting property, we can show that this type system is equivalent to our previous version. We state the property as follows.

Lemma 7.7.2 (Equivalence between syntax-directed and core system).

1. $\Gamma \vdash e \stackrel{\varepsilon}{\in} \tau$ iff $\Gamma \vdash_{SD} e \stackrel{\varepsilon}{\in} \tau$.
2. $\Gamma \vdash v \in \tau$ iff $\Gamma \vdash_{SD} v \in \tau$.

When thinking about syntax-directed judgements, it is useful to assign *modes* to the components of the judgement. These modes correspond to whether that component is an input or an output to a function that determines whether the judgement is derivable. In this case, we have designed the type system as if the effect is an input to the judgement: i.e. it should be larger the actual effect of the term.

7.8 Historical notes and further reading

Type-and-effect systems [[GL86](#), [LG88](#)] were originally developed to track memory usage.

8

A Nontermination monad

In this chapter we contrast the effect system from the previous chapter with a *monadic* treatment of effects. The key idea is that the type system separates the terminating part of the language from programs that have the potential to diverge. To do so, we use a monadic type that encapsulates potentially nonterminating code and isolates it from the rest of the language. If a term has the type $\Box\tau$ then it could possibly diverge. Terms of all other types are guaranteed to terminate. (As in the previous chapter, we call the terminating part of the language *pure* because it lacks the nontermination effect.)

For comparison with the previous languages, we continue to work with a variant of a fine-grained call-by-value language. However, for simplicity we have removed a few constructs to the language (sums, products, and recursive types).

8.1 Language definition

This monadic language starts with a pure core, a fine-grained CBV language with natural numbers and nonrecursive functions. On top of this core, we add the monadic type $\Box\tau$ and its operations: return and bind, written `box v` and `x ← e1; e2`, and recursive function values `fun f x.e`. The two value forms `box v` and `fun f x.e`, introduce the monadic type. The bind operation sequences potentially nonterminating computations.

Definition 8.1.1 (MON: Syntax).

$\tau ::= \text{Nat} \mid \tau_1 \rightarrow \tau_2 \mid \square\tau$	types
$v ::= x$	Variables
zero	zero
S v	successor
$\lambda x.e$	Functions
fun $f x.e$	Recursive functions
box v	monadic pure
$e ::= v_1 v_2$	application
case v of $\{0 \Rightarrow e_1; S x \Rightarrow e_2\}$	test for zero
ret v	value
let $x = e_1$ in e_2	pure sequencing
$x \leftarrow e_1; e_2$	monadic bind

We show the new rules in the definitions below. The core rules (for nats and nonrecursive functions) are the same as in the REC language, so we omit them from the definition.

First, we include three new small-step rules in the operational semantics. When recursive function values are applied, they perform two substitutions: one for the argument and one for the recursive definition. The bind operation reduces its first subterm until it gets to a returned, boxed value. It then substitutes this value into the body and continues evaluation.

Definition 8.1.2 (Small-step rules).	$e \rightsquigarrow e'$	(term e steps to e')
S-APP-FUN		S-BIND
$(\text{fun } f x.e) v \rightsquigarrow e[\text{fun } f x.e/f, v/x]$		$x \leftarrow \text{ret}(\text{box } v); e_2 \rightsquigarrow e_2[v/x]$
	$\frac{\text{S-BIND-CONG}}{e_1 \rightsquigarrow e'_1}$	
	$x \leftarrow e_1; e_2 \rightsquigarrow x \leftarrow e'_1; e_2$	

We also include three new typing rules, and modify the existing rule for let terms.

The **box** v

Definition 8.1.3 (Type system).

$\boxed{\Gamma \vdash v \in \tau}$		(in context Γ , value v has type τ)
	$\frac{\text{TV-PURE}}{\Gamma \vdash v \in \tau}$	
	$\frac{\text{TV-RFUN}}{\Gamma, f:\tau_1 \rightarrow \square\tau_2, x:\tau_1 \vdash e \in \square\tau_2}$	
$\boxed{\Gamma \vdash e \in \tau}$		(in context Γ , term e has type τ)
	$\frac{\text{TE-BIND}}{\Gamma \vdash e_1 \in \square\tau_1}$	
	$\frac{\Gamma, x:\tau_1 \vdash e_2 \in \square\tau_2}{\Gamma \vdash x \leftarrow e_1; e_2 \in \square\tau_2}$	
	$\frac{\text{TE-LET-NB}}{\Gamma \vdash e_1 \in \tau_1}$	
	$\frac{\tau_1 \text{ is not a box type}}{\Gamma, x:\tau_1 \vdash e_2 \in \tau_2}$	
	$\frac{}{\Gamma \vdash \text{let } x = e_1 \text{ in } e_2 \in \tau_2}$	

The **box** v term injects a (terminating) value into the monadic part of the language. Even though we know this value terminates, it is sound to forget this information. Recursive functions are type checked like nonrecursive functions, with two changes: they have access to recursion through the first bound variable, and they must return a monadic result of $\square\tau$. Finally, the typing rule for bind operation is like let, but requires that both subterms have monadic types. We also modify the typing rule for let to require that its first subterm *not* have a monadic type, for reasons that we describe below.

8.1.1 Examples

Here are some examples of terms in this calculus:

1. A function that loops forever when applied. This function is itself a value, so its type doesn't have a box at the top-level, only on the return type of the function.

$$\vdash \mathbf{fun} f x. f x \in \mathbf{Nat} \rightarrow \square\mathbf{Nat}$$

2. A nonterminating term. (This term always steps to itself). It has a monadic type.

$$\vdash (\mathbf{fun} f x. f x) 3 \in \square\mathbf{Nat}$$

3. An addition function, of type $\mathbf{Nat} \rightarrow \mathbf{Nat} \rightarrow \square\mathbf{Nat}$. From its type, any potential divergence comes after providing two arguments (but it won't diverge).

$$\lambda x. \mathbf{ret} (\mathbf{fun} f y. \mathbf{case} y \text{ of } \{0 \Rightarrow \mathbf{ret} (\mathbf{box} x); \mathbf{S} z \Rightarrow x_1 \leftarrow f y; \mathbf{ret} (\mathbf{box} (\mathbf{succ} x_1))\})$$

8.2 Restricted typing rule for let

Consider the difference between these two expressions:

$$\vdash y \leftarrow (\mathbf{fun} f x. f x) 2; \mathbf{ret} (\mathbf{box} 4) \in \square\mathbf{Nat}$$

$$\vdash \mathbf{let} y = ((\mathbf{fun} f x. f x) 2) \mathbf{in} \mathbf{ret} 4 \in \mathbf{Nat}$$

Under the operational semantics, both expressions evaluate the same way, i.e., they both go into an infinite loop. Yet, there is a problem with the second example: its type is \mathbf{Nat} , but it doesn't terminate!

The issue with the second example is that in $\mathbf{let} x = e_1 \mathbf{in} e_2$, the subterm e_2 is not required to use x . But if x is not used (say in some bind), then any trace in the type system of its nontermination effect is lost. This is in contrast with bind, which requires the type of the entire expression to be the monadic type, no matter whether the body uses x or not.

Therefore, we have two forms of sequencing in this language. The let operation sequences a pure term, in any type of computation. The bind operation sequences a monadic term, but its result must be monadic. This is also why we need to have

two separate terms with congruence rules; we cannot do all computation with let only.

Because this is a CBV language, the restriction in the let rule permeates through the derived forms that make this look like a “normal” language. For example, our derived typing rule for full application includes an additional condition that the *argument* of the application is not a box type.

$$\begin{array}{c} \text{TE-APP-NB} \\ \Gamma \vdash e_1 \in \tau_1 \rightarrow \tau_2 \\ \Gamma \vdash e_2 \in \tau_1 \\ \tau_1 \text{ is not a box type} \\ \hline \Gamma \vdash e_1 e_2 \in \tau_2 \end{array}$$

The restriction in the let rule is enough to be sure that our type system accurately characterizes the effect of computation. We can use a logical relations proof to show that if $\vdash e \in \tau$ and τ is not a box type then $e \rightsquigarrow \text{ret } v$ for some value.

8.3 Translation to Dependent Type Theory

The point of a monadic type system is that it *isolates* the effectful part of the computation from the pure computation. What this means is that we can give this language a clean denotational interpretation in a dependent type theory, as long as we can find an interpretation in that type theory for the monadic type.

One option for interpreting potentially divergent terms in dependent type theory is to use a *clock* with a timeout: i.e. a count of how many recursive iterations are available.

The type `Timed A` is a monad: it passes around the current clock and causes a recursive definition to fail if there is no more time left on the clock. The return of this monad ignores the current clock and returns its value directly. The bind for this monad supplies both subterms with the current clock value, but only evaluates the second if the first succeeds. Finally, the loop operation defines recursive functions. It checks the clock and if there is time remaining, executes the function body, providing a recursive definition that uses the smaller clock.

```

Definition Timed (A : Type) : Type := nat -> option A.

Definition pure {A} (v:A) : Timed A := fun (x:nat) => Some v.

Definition bind {A B} (t : Timed A) (u : A -> Timed B) : Timed B :=
  fun x => match t x with
    | Some y => u y x
    | None => None
  end.

Fixpoint loop {A B}
  (f: (A -> Timed B) -> (A -> Timed B)) : A -> Timed B :=
  fun a x =>
  match x with
  | O => None
  | S n => f (fun b y => loop f b n) a (S n)
  end.

```

With these operations, we can translate well-typed values and terms to dependent type theory, i.e. constructive logic.

Definition 8.3.1 (Translation to dependent type theory).

$$\begin{aligned}
 |\text{Nat}| &= \mathbb{N} \\
 |\tau_1 \rightarrow \tau_2| &= |\tau_1| \rightarrow |\tau_2| \\
 |\square\tau| &= \text{Timed } |\tau| \\
 \\
 |x| &= x \\
 |\text{zero}| &= 0 \\
 |\text{succ } v| &= S|v| \\
 |\lambda x. e| &= \lambda x. |e| \\
 |\text{fun } f \ x. e| &= \text{loop } (\lambda f. \lambda x. |e|) \\
 |\text{box } v| &= \text{pure } |v| \\
 \\
 |v_1 \ v_2| &= |v_1| \ |v_2| \\
 |\text{let } x = e_1 \text{ in } e_2| &= (\lambda x. |e_2|) \ |e_1| \\
 |x \leftarrow e_1; e_2| &= \text{bind } |e_1| \ (\lambda x. |e_2|)
 \end{aligned}$$

Lemma 8.3.1 (Type preservation).

If a term or value typechecks in MON, then its translation type checks in dependent type theory.

1. If $\Gamma \vdash e \in \tau$ then $|\Gamma| \vdash |e| \in |\tau|$
2. If $\Gamma \vdash v \in \tau$ then $|\Gamma| \vdash |v| \in |\tau|$

8.4 The translation to EFF

We can also translate our monadic language to our effect-tracking language.

The key idea of the translation is that we are going to map all terms in MON to pure terms in EFF. But what about nontermination? We will do that by suspending monadic computations in “thunks” (i.e. functions that take a trivial argument) and recording their potentially effect in the latent effect of the function type. For example, we have:

$$\square\text{Nat} = \text{Nat} \xrightarrow{\text{DIV}} \text{Nat}$$

(Note that the argument Nat in the translation will always be 0. We are using the function to defer computation only, not pass fuel around.)

This translation depends on the ability to defer effects in the target language.

Definition 8.4.1 (Deferred binary function).

$$v_\eta = \lambda x. \text{ret } (\lambda y. v \ x \ y)$$

Lemma 8.4.1. If $\Gamma \vdash v \in \tau_1 \xrightarrow{\varepsilon_1} (\tau_2 \xrightarrow{\varepsilon_2} \tau_3)$ then $\Gamma \vdash v_\eta \in \tau_1 \xrightarrow{\perp} \tau_2 \xrightarrow{(\varepsilon_1 \oplus \varepsilon_2)} \tau_3$

Note that deferring a function can change its behavior, as evidenced by the change in the effects. Consider the function:

$$v = \text{rec } f. f : (\text{Nat} \xrightarrow{\perp} \text{Nat}) \xrightarrow{\text{DIV}} (\text{Nat} \xrightarrow{\perp} \text{Nat})$$

This function diverges if you apply it to any argument.

Now, if it is deferred, we get a function that terminates on a single application.

$$v_\eta = \lambda x.\text{ret}(\lambda y.(\text{rec } f.f) x y) : (\text{Nat} \xrightarrow{\perp} \text{Nat}) \xrightarrow{\perp} (\text{Nat} \xrightarrow{\text{DIV}} \text{Nat})$$

Here is the translation in full:

$ \text{Nat} $	$= \text{Nat}$
$ \tau_1 \rightarrow \tau_2 $	$= \tau_1 \xrightarrow{\perp} \tau_2 $
$ \square\tau $	$= \text{Nat} \xrightarrow{\text{DIV}} \tau $
$ x $	$= x$
$ \lambda x.e $	$= \lambda x. e $
$ k $	$= k$
$ \text{box } v $	$= \lambda x.\text{ret} v $
$ \text{fun } f.x.e $	$= (\text{rec } f.\lambda x. e [f_\eta/f])_\eta$
$ v_1 v_1 $	$= v_1 v_2 $
$ \text{ret } v $	$= \text{ret} v $
$ \text{case } v \text{ of } \{0 \Rightarrow e_0; \mathbf{S} x \Rightarrow e_1\} $	$= \text{case } v \text{ of } \{0 \Rightarrow e_0 ; \mathbf{S} x \Rightarrow e_1 \}$
$ \text{let } x = e_1 \text{ in } e_2 $	$= \text{let } x = e_1 \text{ in } e_2 $
$ x \leftarrow e_1; e_2 $	$= \text{ret}(\lambda y.\text{let } x = e_1 \text{ in } i_n e_1 0)$

Lemma 8.4.2 (Type preservation).

If a term or value typechecks in MON, then its translation type checks in EFF and has the pure effect.

1. If $\Gamma \vdash e \in \tau$ then $|\Gamma| \vdash |e| \in |\tau|$

2. If $\Gamma \vdash v \in \tau$ then $|\Gamma| \vdash |v| \in |\tau|$

You may wonder whether this transformation changes the operational semantics of the language, given that we are suspending all potentially divergent computations, and differing recursive definitions.

This is a valid worry.

For example, in the source language, we have:

$$\text{loop} = (\text{fun } f.x.f x) 3 : \square\text{Nat}$$

But after the translation, this turns into the suspension:

$$\text{ret}(\lambda z.(\lambda x_1.\lambda x_2.(\text{rec } f.\lambda x.(\lambda x_1.\lambda x_2.f x_1 x_2) x) x_1 x_2) 3) : \text{Nat} \xrightarrow{\text{DIV}} \text{Nat}$$

What this means is that we will have to be careful about how we state the soundness theorem for our language when it comes to the operational semantics.

8.5 What does a CBN Monadic language look like?

What if we had started with call-by-name (CBN) base language instead of our fine-grained CBV language?

Let's first contrast the structure of a CBN base language with our fine-grained CBV language. Note that CBN languages do not make a strict syntactic restriction between values and terms: values are merely a sublanguage of terms.

$$\begin{array}{lcl} e & ::= & x \mid \lambda x. e \mid e_1 e_2 \\ & | & \text{zero} \mid \text{succ } e \mid \text{case } e \text{ of } \{0 \Rightarrow e_0; \mathbf{S} y \Rightarrow e_1\} \\ & | & \text{box } e \mid x \leftarrow e_1; e_2 \mid \text{fun } f x. e \\ v & ::= & \lambda x. e \mid \text{zero} \mid \text{succ } e \mid \text{box } e \end{array}$$

What makes CBN languages CBN is that they β -reduce function applications before reducing their arguments. This is possible because substitution in CBN replaces variables by *terms* instead of values.

$$\frac{\text{S-APP-CBN}}{(\lambda x. e_1) e_2 \rightsquigarrow e_1[e_2/x]}$$

Because of this change, we can also change the semantics of other operations as well. We can make our natural numbers nonstrict: in other words, we can treat **succ** e as a value.

$$\frac{\text{S-CASE-SUCC-CBN}}{\text{case } (\text{succ } e) \text{ of } \{0 \Rightarrow e_0; \mathbf{S} y \Rightarrow e_1\} \rightsquigarrow e_1[e/x]}$$

And we can make our recursive definitions more eager. We do not need to consider a recursive function to be a value. It can unfold immediately.

$$\frac{\text{S-FUN-CBN}}{(\text{fun } f x. e) e \rightsquigarrow \lambda x. (e[\text{fun } f x. e/f])}$$

In the fine-grained CBV language we combined all congruence rules into one through the use of let expressions. Here we cannot do that. Even though we can add a strict let expression to this language, we cannot require eliminators to take values only because variables are not values. Therefore we are back to adding congruence rules for all of the strict subterms in the language.

$$\frac{\begin{array}{c} \text{S-APP-CONG-CBN} \\ e_1 \rightsquigarrow e'_1 \\ \hline e_1 e_2 \rightsquigarrow e'_1 e_2 \end{array}}{} \quad \frac{\begin{array}{c} \text{S-CASE-CONG-CBN} \\ e \rightsquigarrow e' \\ \hline \text{case } e \text{ of } \{0 \Rightarrow e_0; \mathbf{S} y \Rightarrow e_1\} \rightsquigarrow \text{case } e' \text{ of } \{0 \Rightarrow e_0; \mathbf{S} y \Rightarrow e_1\} \end{array}}{}$$

Finally, we also make box-terms nonstrict, delaying the evaluation of the argument inside the box.

$$\frac{\text{S-BIND-CBN}}{x \leftarrow \text{box } e_1; e_2 \rightsquigarrow e_2[e_1/x]}$$

The upshot of this design is that we don't need to include any τ is not a box type restrictions into the type system. In a CBN language, the term $(\lambda x. 3)$ *loop* terminates, so functions are free to discard the effects of their arguments.

9

Explicit control stacks

In this chapter, we look at what happens if we make the *control stack* explicit in the semantics of a programming language. The control stack records where we are during the execution of a program. In both the small-step and the big-step semantics, this record is stored directly (but implicitly) in the term as it evaluates.

We have two primary motivations for making the control stack explicit:

First, if we want to give a semantics to some programming language features that interact with the control stack, such as exceptions or continuations, we need to have access to this structure.

Second, if we want to define what it means for two programs to be “the same” then we want to think about all of the ways that a program can interact with its environment. A control stack is a way to “use” a closed program and gives us a definition of equivalence based on both programs always producing the same result on any stack that we might run them on.

9.1 Stack-based semantics

Let’s first get comfortable with stack based semantics. We’ve already seen two different ways of defining the operational semantics of a programming language: *small-step* and *big-step* semantics. In this section, we will think about variants of those semantics that use an explicit control stack.

We continue to work with REC, a fine-grained CBV language, in this chapter. To recap, this language explicitly separates values and terms. Furthermore, most constructs in the term language do only one step at a time. The only exception is `let $x = e_1$ in e_2` , which sequences one computation after another.

$\tau ::=$	Unit Void Nat $\tau_1 \rightarrow \tau_2$ $\tau_1 * \tau_2$ $\tau_1 + \tau_2$	types
$v ::=$	x unit 0 S v $\lambda x.e$ (v_1, v_2) inj ₁ v inj ₂ v rec $x.v$ fold v	variables unit zero successor functions pairs sums recursive value recursive type
$e ::=$	$v_1 v_2$ case v of {0 $\Rightarrow e_1$; S $x \Rightarrow e_2$ } prj ₁ v prj ₂ v case v of { inj ₁ $x \Rightarrow e_1$; inj ₂ $x \Rightarrow e_2$ } unfold v ret v let $x = e_1$ in e_2	application test for zero projection case unfold recursive value value sequencing

9.1.1 Primitive vs. control reductions

Before we introduce explicit stacks, will will first refactor our existing small step and big-step reduction relations to first isolate *primitive* reductions as a separate judgement. That structure will allow us to focus only on the reductions that affect the control stack. Notice that all of these rules are *axioms*, they do not have any premises.

Definition 9.1.1 (Primitive reductions). $\boxed{e \twoheadrightarrow e'}$ (*e primitive reduces to e'*)

$$\begin{array}{c}
\text{PRIM-BETA} \quad \frac{}{\boxed{(\lambda x.e) v \twoheadrightarrow e[v/x]}} \quad \text{PRIM-CASE-ZERO} \quad \frac{}{\boxed{\text{case } 0 \text{ of } \{0 \Rightarrow e_1; \mathbf{S} x \Rightarrow e_2\} \twoheadrightarrow e_1}} \\
\text{PRIM-CASE-SUCC} \quad \frac{}{\boxed{\text{case } (\mathbf{succ } k) \text{ of } \{0 \Rightarrow e_1; \mathbf{S} x \Rightarrow e_2\} \twoheadrightarrow e_2[k/x]}} \quad \text{PRIM-PRJ1} \quad \frac{}{\boxed{\mathbf{prj}_1(v_1, v_2) \twoheadrightarrow \mathbf{ret } v_1}} \\
\text{PRIM-PRJ2} \quad \frac{}{\boxed{\mathbf{prj}_2(v_1, v_2) \twoheadrightarrow \mathbf{ret } v_2}} \quad \text{PRIM-CASE-INJ1} \quad \frac{}{\boxed{\text{case } (\mathbf{inj}_1 v) \text{ of } \{\mathbf{inj}_1 x \Rightarrow e_1; \mathbf{inj}_2 x \Rightarrow e_2\} \twoheadrightarrow e_1[v/x]}} \\
\text{PRIM-CASE-INJ2} \quad \frac{}{\boxed{\text{case } (\mathbf{inj}_2 v) \text{ of } \{\mathbf{inj}_1 x \Rightarrow e_1; \mathbf{inj}_2 x \Rightarrow e_2\} \twoheadrightarrow e_2[v/x]}} \quad \text{PRIM-UNFOLD} \quad \frac{}{\boxed{\mathbf{unfold } (\mathbf{fold } v) \twoheadrightarrow \mathbf{ret } v}} \\
\text{PRIM-APP-REC} \quad \frac{}{\boxed{(\mathbf{rec } x.v_1) v_2 \twoheadrightarrow v_1[\mathbf{rec } x.v_1/x] v_2}} \quad \text{PRIM-PRJ1-REC} \quad \frac{}{\boxed{\mathbf{prj}_1(\mathbf{rec } x.v) \twoheadrightarrow \mathbf{prj}_1 v[\mathbf{rec } x.v/x]}} \\
\text{PRIM-PRJ2-REC} \quad \frac{}{\boxed{\mathbf{prj}_2(\mathbf{rec } x.v) \twoheadrightarrow \mathbf{prj}_2 v[\mathbf{rec } x.v/x]}}
\end{array}$$

Our revised version of the small-step semantics for REC includes all primitive reductions, plus the two rules that work with let expressions. This semantics is trivially equivalent to the version presented in Chapter 5. (You should convince yourself of this fact.) As before, the only rule of the semantics with a premise is rule S-LET-CONG.

Definition 9.1.2 (Small-step semantics). $\boxed{e \rightsquigarrow e'} \quad (e \text{ small-steps to } e')$

$$\frac{\begin{array}{c} \text{S-PRIM} \\ e \rightarrow e' \end{array}}{e \rightsquigarrow e'} \quad \frac{\text{S-LET-RET}}{\text{let } x = \text{ret } v \text{ in } e \rightsquigarrow e[v/x]} \quad \frac{\begin{array}{c} \text{S-LET-CONG} \\ e_1 \rightsquigarrow e'_1 \end{array}}{\text{let } x = e_1 \text{ in } e_2 \rightsquigarrow \text{let } x = e'_1 \text{ in } e_2}$$

Now let's design a big-step semantics for the REC language. We can use primitive reductions to do so using only three rules.

Definition 9.1.3 (Big-step semantics). $\boxed{e \Rightarrow v} \quad (e \text{ big-steps to } v)$

$$\frac{\begin{array}{c} \text{BS-PRIM} \\ e_1 \rightarrow e_2 \quad e_2 \Rightarrow v \end{array}}{e_1 \Rightarrow v} \quad \frac{\text{BS-VAL}}{v \Rightarrow v} \quad \frac{\begin{array}{c} \text{BS-LET} \\ e_1 \Rightarrow v_1 \\ e_2[v_1/x] \Rightarrow v_2 \end{array}}{\text{let } x = e_1 \text{ in } e_2 \Rightarrow v_2}$$

Both of these semantics are deterministic. Furthermore, we can show that they are also equivalent to each other, using the same technique that we used in Chapter 3.

Theorem 9.1.1 (Equivalence of semantics). For closed expressions e , we have $e \rightsquigarrow^* \text{ret } v$ if and only if $e \Rightarrow v$.

The proofs of both directions of these theorems is by straightforward induction on the derivations, with the assistance of the following two easy lemmas.

As in Chapter 3, we needed a step-expansion lemma to show the forward direction of the lemma.

Lemma 9.1.1 (step-expansion). If $e_1 \rightsquigarrow e_2$ and $e_2 \Rightarrow v$ then $e_1 \Rightarrow v$.

As in Chapter 3, to prove the reverse direction of this lemma, we need to show that we can multistep let expressions.

Lemma 9.1.2 (ms-let-cong). If $e_1 \rightsquigarrow^* e'_1$ then $\text{let } x = e_1 \text{ in } e_2 \rightsquigarrow^* \text{let } x = e'_1 \text{ in } e_2$

9.2 Stack-based semantics

Now let's make the stack explicit in our operational semantics. What is the point of the control stack?

In the stack-based semantics, each step describe the reduction of an abstract *machine*. A machine $\langle s, e \rangle$ is a pair of some control *stack* and some term e executing on that stack.

$$\begin{array}{lll} \text{frame} & f & ::= \quad \text{let } x = _ \text{ in } e \\ \text{stack} & s & ::= \quad f : s \mid [] \\ \text{machine} & m & ::= \quad \langle s, e \rangle \end{array}$$

A control stack is a *list of frames* which record “what we need to do next” after evaluating the current term e . As we accumulate more computational work, we *push* a new frame on the stack. Once we have evaluated the current term to a value, we *pop* the current frame from the stack (if any) to get the next step in our evaluation. If the stack is empty, we have reached the end of the computation.

In the fine-grained REC language, we only have one form of term that sequences computation: $\text{let } x = e_1 \text{ in } e_2$. That means that we only need one kind of frame, written $\text{let } x = _ \text{ in } e_2$. This frame gets pushed on the stack when we are evaluating the right-hand-side of a let definition. Once this result has been calculated, the frame remembers what we should do with the value of x in e_2 the body of the term.

9.2.1 Stack-based small-step semantics

Using this framework, we can say how (abstract) machines evaluate, step-by-step. There are only three rules: either the term does a primitive reduction, the term is a let so we push a frame on the stack, or the term is a returned value, and there is a let-frame waiting on the stack, so we pop it to access the next step.

Definition 9.2.1 (Stack based small-steps).

$$\begin{array}{c}
 \boxed{m \mapsto m'} \quad (m \text{ small-steps to } m') \\
 \begin{array}{ccc}
 \text{SSM-PRIM} & & \text{SSM-PUSH} \\
 \dfrac{e \rightarrow e'}{\langle s, e \rangle \mapsto \langle s, e' \rangle} & & \dfrac{}{\langle s, \text{let } x = e_1 \text{ in } e_2 \rangle \mapsto \langle \text{let } x = _ \text{ in } e_2 : s, e_1 \rangle} \\
 \text{SSM-POP} \\
 \dfrac{}{\langle \text{let } x = _ \text{ in } e_2 : s, \text{ret } v \rangle \mapsto \langle s, e_2[v/x] \rangle}
 \end{array}
 \end{array}$$

We also think about the reflexive, transitive closure of this relation, written $m \mapsto^* m'$. A complete evaluation is one where m' is some *terminal* machine state.

Definition 9.2.2 (Terminal machine state). A machine has finished evaluation when the stack is empty and the term is a returned value. In other words, when it is of the form: $\langle \[], \text{ret } v \rangle$

The key property of this stack-based semantics is that it *linearizes* computation. None of the rules are inductive. Instead, we have replaced the congruence rule for let with pushing and popping on the stack. That means that a computation may need more steps of evaluation to produce the final value. However, each individual step is simpler as we don't need to look deep within the term to find a reduction; it is always apparent from the top-level configuration of the machine.

9.2.2 Example: small-step semantics without and with stacks

Consider the small step evaluation of this term:

```

let x = (let y = ((λx.ret x) 0) in 1 + y) in let z = (let y = (2 + x) in ret y) in (3 + z)
~~ function application (primitive)
let x = (let y = (ret 0) in 1 + y) in let z = (let y = (2 + x) in ret y) in (3 + z)
~~ inline y
let x = 1 + 0 in let z = (let y = (2 + x) in ret y) in (3 + z)
~~ addition (primitive)
let x = ret 1 in let z = (let y = (2 + x) in ret y) in (3 + z)
~~ inline x
let z = (let y = (2 + 1) in ret y) in (3 + z)
~~ addition (primitive)
let z = (let y = (ret 3) in ret y) in (3 + z)
~~ inline y
let z = (ret 3) in (3 + z)
~~ inline z
(3 + 3)
~~ addition (primitive)
ret 6

```

During this evaluation, to figure out how to step this term, we needed to find the innermost reduction $(\lambda x.\text{ret } x) 0$ hiding inside the nested let expression. Then, we had to search of the next additions to do, still hiding inside nested lets.

In a stack based semantics, we will consider the evaluation of a stack s paired with a term e . When the term is a let expression, we will push the congruent part of the let expression on the stack and then focus on the right hand side.

This takes a few extra steps on our example to do the pushing.

```

⟨[], let x = (let y = ((λx.ret x) 0) in 1 + y) in let z = (let y = (2 + x) in ret y) in (3 + z)⟩
~~ push let x
⟨let x = _ in let z = (let y = (2 + x) in ret y) in (3 + z) : [], let y = ((λx.ret x) 0) in 1 + y⟩
~~ push let y
⟨let y = _ in 1 + y : let x = _ in let z = (let y = (2 + x) in ret y) in (3 + z) : [], (λx.ret x) 0⟩
~~ function application (primitive)
⟨let y = _ in 1 + y : let x = _ in let z = (let y = (2 + x) in ret y) in (3 + z) : [], ret 0⟩
~~ pop (subst for y)
⟨let x = _ in let z = (let y = (2 + x) in ret y) in (3 + z) : [], 1 + 0⟩
~~ addition (primitive)
⟨let x = _ in let z = (let y = (2 + x) in ret y) in (3 + z) : [], 1⟩
~~ pop (subst for x)
⟨[], let z = (let y = (2 + 1) in ret y) in (3 + z)⟩
~~ push let z
⟨let z = _ in (3 + z) : [], (let y = (2 + 1) in ret y)⟩
~~ push let y
⟨let y = _ in ret y : let z = _ in (3 + z) : [], 2 + 1⟩
~~ addition (primitive)
⟨let y = _ in ret y : let z = _ in (3 + z) : [], 3⟩
~~ pop (subst for y)
⟨let z = _ in (3 + z) : [], ret 3⟩
~~ pop (subst for z)
⟨[], 3 + 3⟩
~~ addition (primitive)
⟨[], ret 6⟩

```

9.2.3 Stack-based big-step semantics

Now let's design a big-step version of the semantics that uses an explicit stack. Note that in this version of the big-step semantics, each rule has at most one premise. Where derivations in the original bigstep semantics could be wide (due to the let rule) in this version they are always very tall. This version of the semantics also linearizes the computation.

$$\begin{array}{c}
 \boxed{\langle s, e \rangle \Downarrow v} \qquad \qquad \qquad (e \text{ small-steps to } e') \\
 \text{BSM-FINAL} \qquad \qquad \qquad \text{BSM-PRIM} \qquad \qquad \qquad \text{BSM-PUSH} \\
 \frac{}{\langle[], \text{ret } v \rangle \Downarrow v} \qquad \qquad \frac{e \rightarrow e' \quad \langle s, e' \rangle \Downarrow v}{\langle s, e \rangle \Downarrow v} \qquad \qquad \frac{\langle \text{let } x = _ \text{ in } e_2 : s, e_1 \rangle \Downarrow v}{\langle s, \text{let } x = e_1 \text{ in } e_2 \rangle \Downarrow v} \\
 \text{BSM-POP} \\
 \frac{\langle s, e_2[v_1/x] \rangle \Downarrow v}{\langle \text{let } x = _ \text{ in } e_2 : s, \text{ret } v_1 \rangle \Downarrow v}
 \end{array}$$

9.2.4 Example: big-step semantics with stacks

Here's the same example as above, expressed with our big-step semantics. This derivation is the “vertical” version of the small-step derivation. Note that every

step of the machine is the same—the only difference is that we have a single tall derivation instead of a long sequence of steps.

$$\begin{array}{c}
 \frac{}{\langle[], \mathbf{ret} 6\rangle \Downarrow 6} \\
 \frac{}{\langle[], 3 + 3\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } y = _ \mathbf{in } \mathbf{ret} y : \mathbf{let } z = _ \mathbf{in } (3 + z) : [], 3\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } y = _ \mathbf{in } \mathbf{ret} y : \mathbf{let } z = _ \mathbf{in } (3 + z) : [], 2 + 1\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } z = _ \mathbf{in } (3 + z) : [], (\mathbf{let } y = (2 + 1) \mathbf{in } \mathbf{ret} y)\rangle \Downarrow 6} \\
 \frac{}{\langle[], \mathbf{let } z = (\mathbf{let } y = (2 + 1) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z)\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } x = _ \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z) : [], 1\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } x = _ \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z) : [], 1 + 0\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } y = _ \mathbf{in } 1 + y : \mathbf{let } x = _ \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z) : [], \mathbf{ret} 0\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } y = _ \mathbf{in } 1 + y : \mathbf{let } x = _ \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z) : [], (\lambda x. \mathbf{ret} x) 0\rangle \Downarrow 6} \\
 \frac{}{\langle\mathbf{let } x = _ \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z) : [], \mathbf{let } y = ((\lambda x. \mathbf{ret} x) 0) \mathbf{in } 1 + y\rangle \Downarrow 6} \\
 \frac{}{\langle[], \mathbf{let } x = (\mathbf{let } y = ((\lambda x. \mathbf{ret} x) 0) \mathbf{in } 1 + y) \mathbf{in } \mathbf{let } z = (\mathbf{let } y = (2 + x) \mathbf{in } \mathbf{ret} y) \mathbf{in } (3 + z)\rangle \Downarrow 6}
 \end{array}$$

9.3 Equivalence between these semantics

Theorem 9.3.1 (Stack machine correctness). The following are equivalent:

1. $e \rightsquigarrow^* \mathbf{ret} v$
2. $\langle[], e\rangle \mapsto^* \langle[], \mathbf{ret} v\rangle$
3. $e \Rightarrow v$
4. $\langle[], e\rangle \Downarrow v$

From above, already have the equivalence between the original small and big step semantics, i.e. (1) iff (3). In this section, we will complete the relationships between these judgements.

9.3.1 Stack completeness

The original big-step and small-step semantics can be simulated by the stack-based small-step semantics. It turns out that it is simplest to first show the connection between the big-step semantics and the stack based semantics, i.e. that (3) implies (2). This also gives us (1) implies (2).

Lemma 9.3.1 (Stack completeness (big-step)). If $e \Rightarrow v$ then $\langle s, e\rangle \mapsto^* \langle s, \mathbf{ret} v\rangle$, for any s .

This proof is by induction on the big-step semantics.

Lemma 9.3.2 (Stack completeness (small-step)). If $e \rightsquigarrow^* \mathbf{ret} v$ then $\langle[], e\rangle \mapsto^* \langle[], \mathbf{ret} v\rangle$

This follows from the big-step completeness and the equivalence between big and small step semantics.

9.3.2 Soundness

It is more difficult to prove the soundness lemma: the fact that all executions on the machine semantics can be simulated by the original semantics.

Definition 9.3.1 (Unravel). Define $s\{e\}$ by recursion on the stack:

- $\llbracket e \rrbracket = e$
- $(\text{let } x = _ \text{ in } e_2 : s)\{e\} = s\{\text{let } x = e \text{ in } e_2\}$

Lemma 9.3.3 (Stack congruence). If $e \rightsquigarrow e'$ then $s\{e\} \rightsquigarrow s\{e'\}$

Lemma 9.3.4 (Unravel step). If $\langle s, e \rangle \mapsto \langle s', e' \rangle$ then $s\{e\} \rightsquigarrow^* s'\{e'\}$.

Lemma 9.3.5 (Unravel multistep). If $\langle s, e \rangle \mapsto^* \langle s', \text{ret } v \rangle$ then $s\{e\} \rightsquigarrow^* s'\{\text{ret } v\}$.

The soundness lemma is a corollary of the multistep lemma above.

Corollary 9.3.1 (Soundness). If $\langle \llbracket \cdot \rrbracket, e \rangle \mapsto^* \langle \llbracket \cdot \rrbracket, \text{ret } v \rangle$ then $e \rightsquigarrow^* \text{ret } v$.

9.3.3 Small-step and big-step stack equivalence

Because the small-step and big-step stack semantics are so similar it is not difficult to show the equivalence of (2) and (4). Both of the directions follow by straightforward induction on the derivations.

Lemma 9.3.6. If $\langle s, e \rangle \mapsto^* \langle \llbracket \cdot \rrbracket, v \rangle$ then $\langle s, e \rangle \Downarrow v$.

Lemma 9.3.7. If $\langle s, e \rangle \Downarrow v$ then $\langle s, e \rangle \mapsto^* \langle \llbracket \cdot \rrbracket, v \rangle$.

9.4 References and Further reading

This chapter is adapted from Chapter 28 of PFPL, translating to a fine-grained CBV language and adding a big-step version of the machine semantics.

10

Control Effects

Many languages include features for throwing and catching exceptions. Exceptions provide a way for programs to signal errors or break out of loops early.

Fewer languages provide first-class access to the control stack through control operators. Scheme (and its modern-day variant Racket), includes `call-with-current-continuation`, or `call-cc`. (Other languages that have included this feature are Standard ML of New Jersey, Ruby, and Smalltalk.) This feature acts a little like exceptions — by passing a captured continuation, users can go back to earlier parts of the stack — but it is much more powerful. Captured continuations are first class. They can be saved in data structures and stored in function closures. That means that invoking a continuation may take your program to a stack that is very different than the current one.

First-class continuations can be used to implement other high-level control structures that are difficult implement cleanly otherwise, such as exceptions, coroutines or generators, or non-deterministic backtracking. Furthermore, even if languages do not include support for first-class continuations, they can be implemented through the use of “continuation-passing-style” or a concurrency monad.

Effect handlers are newer still and have only recently been added to the OCaml language. This feature, based on an older idea of *delimited continuations* combines features of first-class continuations and exception handling to provide more compositionality and more efficient implementation. This language feature is being actively explored in research languages such as Koka and Eff, and was briefly included in the Scala language.

In this chapter we will consider how to extend the semantics of our stack based language with new forms of control operations.

10.1 Exceptions

Most languages include an extensible variant type for exceptions. For convenience, this type `exn` or `Exception` can be extended with new options corresponding various forms of errors (`Failure` or `NullPointerException`) and signals (`Not_found` or `NoSuchElementException`). For simplicity in our for-

malism, we will define a new type **Exn** as our exception type, but this type will be isomorphic to natural numbers so that we have as many exceptions as we need and can distinguish between different exception values.

Our extension with exceptions includes two new terms in the language: a way to *raise* or trigger an exception, and a way to install an exception handler with **try**.

$$\begin{array}{ll} \text{term} & e ::= \text{try } e_1 \text{ with exn } x \Rightarrow e_2 \mid \text{raise } v \\ \text{frame} & f ::= \text{try_with exn } x \Rightarrow e_2 \end{array}$$

There are three new small-step rules that work with exceptions. The rule **SSM-TRY** pushes an exception handler frame onto the stack and continues executing the body of the “try”. The rule **SSM-DISCARD** pops this exception handler off the stack when the body returns a value. Finally, the rule **SSM-RAISE** uses an auxiliary function (shown below) to scan the stack for a current exception handler when an exception is raised.

$$m \mapsto m' \quad (\text{Stack-based small-step rules})$$

SSM-TRY

$$\overline{\langle s, \text{try } e_1 \text{ with exn } x \Rightarrow e_2 \rangle \mapsto \langle \text{try_with exn } x \Rightarrow e_2 : s, e_1 \rangle}$$

SSM-DISCARD

$$\overline{\langle \text{try_with exn } x \Rightarrow e_2 : s, \text{ret } v \rangle \mapsto \langle s, \text{ret } v \rangle}$$

SSM-RAISE

$$\overline{\langle \text{frame} : s, \text{raise } v \rangle \mapsto \text{find_exn}(\text{frame} : s) v}$$

$$\begin{array}{ll} \text{find_exn}(\text{try_with exn } x \Rightarrow e_2 : s) v &= \langle s, e_2[v/x] \rangle \\ \text{find_exn}(\text{frame} : s) v &= \text{find_exn } s v \\ \text{find_exn} [] v &= \langle [], \text{raise } v \rangle \end{array}$$

When **find_exn** $s v$ successfully finds an exception handler frame on the stack (see first line of the definition) then the machine switches to the stack immediately below the handler and the code in the handler. The value v is the thrown exception value (a natural number in this case) and is substituted for x in the handler code. If the **find_exn** $s v$ reaches the end of the stack without finding an exception handler, the machine stops with an empty stack and an uncaught exception (signified by the **raise** v term). This term cannot step because the rule **SSM-RULE** applies only when the stack is not empty.

10.1.1 Type Safety in the presence of exceptions

The typing rule for the two new terms are shown below:

$$\boxed{\Gamma \vdash v \in \tau}$$

(Typing rules for values)

TV-EXN

$$\overline{\Gamma \vdash \text{exn } k \in \mathbf{Exn}}$$

$$\begin{array}{c}
 \boxed{\Gamma \vdash e \in \tau} \qquad \qquad \qquad (\textit{Typing rules for terms}) \\
 \begin{array}{c}
 \text{TE-TRY} \\
 \text{TE-RAISE} \\
 \dfrac{\Gamma \vdash v \in \mathbf{Exn}}{\Gamma \vdash \mathbf{raise} \ v \in \tau} \qquad \qquad \dfrac{\begin{array}{c} \Gamma \vdash e_1 \in \tau \\ \Gamma, x : \mathbf{Exn} \vdash e_2 \in \tau \end{array}}{\Gamma \vdash \mathbf{try} \ e_1 \ \mathbf{with} \ \mathbf{exn} \ x \Rightarrow e_2 \in \tau}
 \end{array}
 \end{array}$$

Given any natural number, the value `exn k` has the type `Exn`. Then the term `raise v` expects its argument to be an exception value. Note that this term could have any type, including `Void`. The `raise v` term never produces a value—it just jumps to some other part of the program. So it is safe to pretend that this term produces any type of value that we want.

The term `try e1 with exn x ⇒ e2` requires both the body of the try block `e1` and the exception handler `e2` to have the same type. This expression will either terminate normally, or it will trigger the handler. We can't tell which will happen statically, so the type system requires them to be the same, just as in pattern matching expressions. The variable `x` is a name for the exception value that was thrown.

We can prove type safety for our stack-based language, but we need to extend our type system to be able to talk about stacks, frames, and machines. All three of these judgments are for runtime constructs, so we only care about type checking them with the empty context.

Definition 10.1.1 (Frame, Stack and Machine typing).

$$\begin{array}{c}
 \boxed{\vdash \text{frame} \in \tau_1 \rightsquigarrow \tau_2} \qquad \qquad \qquad (\textit{Frame frame takes values of type } \tau_1 \textit{ to stacks of type } \tau_2) \\
 \begin{array}{c}
 \text{TF-LET} \\
 \dfrac{\begin{array}{c}, x : \tau_1 \vdash e_2 \in \tau_2 \\ \vdash \text{let } x = _ \text{ in } e_2 \in \tau_1 \rightsquigarrow \tau_2 \end{array}}{\vdash \text{try } _ \text{ with exn } x \Rightarrow e_2 \in \tau \rightsquigarrow \tau} \qquad \qquad \text{TF-TRY} \\
 \dfrac{\begin{array}{c}, x : \mathbf{Exn} \vdash e_2 \in \tau \\ , x : \tau \end{array}}{\vdash \text{try } _ \text{ with exn } x \Rightarrow e_2 \in \tau \rightsquigarrow \tau}
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 \boxed{\vdash s \in \tau_1} \qquad \qquad \qquad (\textit{Stack s accepts terms of type } \tau_1) \\
 \begin{array}{c}
 \text{TSS-NIL} \\
 \dfrac{}{\vdash [] \in \tau} \qquad \qquad \qquad \text{TSS-CONS} \\
 \dfrac{\begin{array}{c} \vdash \text{frame} \in \tau_1 \rightsquigarrow \tau_2 \\ \vdash s \in \tau_2 \end{array}}{\vdash (\text{frame} : s) \in \tau_1}
 \end{array}
 \end{array}$$

$$\begin{array}{c}
 \boxed{\vdash m \text{ ok}} \qquad \qquad \qquad (\textit{Machine m is well-typed}) \\
 \begin{array}{c}
 \text{M-OK-S} \\
 \dfrac{\begin{array}{c} \vdash s \in \tau_1 \qquad \vdash e \in \tau_1 \\ \vdash \langle s, e \rangle \text{ ok} \end{array}}{\vdash \langle s, e \rangle \text{ ok}}
 \end{array}
 \end{array}$$

We expect our machines to stop executing when the stack is empty and the term has returned a value, or when the stack is empty and the term has raised an exception (which is uncaught).

Definition 10.1.2 (Terminal machines). A machine is in a terminal state if it is of the form $\langle [], \text{ret } v \rangle$ or $\langle [], \text{raise } v \rangle$.

We can use the definition of terminal machine to talk about what it means for computation to not get stuck.

Theorem 10.1.1 (Type Safety). If $\vdash m \text{ ok}$ and $m \mapsto^* m'$ then either m' is a terminal machine or there exists some m'' such that $m' \mapsto m''$.

The proof of this property is via the preservation and progress theorems. There are a few observations to make about this proof.

First, note that the rule **SSM-RAISE** case of the preservation theorem requires showing that looking for an exception handler produces an well-typed machine. It doesn't matter what type of stack we have when we raise the exception; if its well-formed, then any exception handler that we find on the stack will produce a well-formed machine.

Lemma 10.1.1 (Preservation for `find_exn`). If $\vdash v \in \mathbf{Exn}$ and $\vdash s \in \tau_1$, then $\vdash \mathbf{find_exn}\ s\ v \text{ ok}$

Second, notice more generally that we don't (need to) track the machine type in the $\vdash m \text{ ok}$ judgment. This means that our preservation theorem doesn't necessarily "preserve" the type of the machine, it just preserves the fact that the machine is well-formed as it steps.

Lemma 10.1.2 (Preservation). If $\vdash m \text{ ok}$ and $m \mapsto m'$ then $\vdash m \text{ ok}$.

Third, note our type safety theorem does not rule out uncaught exceptions. The machine configuration $\langle[], \text{raise } v\rangle$ results from throwing an exception when there is no handler on the stack, and is considered a well-formed terminal machine.

It would be difficult for our type system to do better. If we want to be sure that all exceptions are caught by appropriate handlers, i.e. that our language is *exception safe* in addition to *type safe*, we might use some sort of type-and-effect system that tracks which exceptions are thrown and removes those effects when a try-block is in play. The Java Language includes this capability with its mechanism for checked exceptions. However, this feature has turned out to not be very useful to developers.

Without exception safety, type safety is a weaker theorem. Instead of saying that a well-typed program either diverges or returns a value, we say now that a well-typed program either diverges, returns a value, or throws an exception. More things could happen. What this means is that if we have a language that is not type-safe because it gets stuck, we can trivially make it type-safe by changing the semantics to throw an exception in all places where it would get stuck. (This is a bit better than changing the semantics to go into an infinite loop in all places where it would get stuck, an option that was already available to us.) Our type safety theorem is meaningful, but we have to realize that its utility depends on the language semantics (as always). We need to make sure that there are enough reduction rules to actually compute the answers that we need. A trivial language that always and immediately raises an exception would also be type safe.

10.2 Continuations

Control operations, such as `call/cc`, reify the control stack as a data structure. This provides a more expressive way to "jump" to new computation contexts.

Definition 10.2.1 (Control operations).

$$\begin{array}{lcl} \text{types} & \tau & ::= \quad \mathbf{Cont} \, \tau \\ \text{value} & v & ::= \quad \mathbf{cont} \, s \\ \text{term} & e & ::= \quad \mathbf{letcc} \, x \, \mathbf{in} \, e \mid \mathbf{throw} \, v_1 \, v_2 \end{array}$$

Let's add a **letcc** x in e operation to the REC language. This term captures the current stack as a new continuation value **cont** s , bound to the variable x . In the scope of x , the **throw** $v_1 \, v_2$ operation changes whatever the current execution stack is to this saved stack, which also receives the value v_2 .

Definition 10.2.2 (Small-step stack based semantics for control).

$m \mapsto m'$	<i>(Stack-based small-step rules)</i>
SSM-LETCC	SSM-THROW
$\langle s, \mathbf{letcc} \, x \, \mathbf{in} \, e \rangle \mapsto \langle s, e[\mathbf{cont} \, s/x] \rangle$	$\langle s, \mathbf{throw} \, (\mathbf{cont} \, s') \, v \rangle \mapsto \langle s', \mathbf{ret} \, v \rangle$

Because we have added a new form of value to the language, we also add a new type, called **Cont** τ . The **cont** s expression introduces a new continuation value, formed from a stack of type τ . In a **letcc** x in e term, x has a continuation type that must be the same as the type of the whole term; this is the point of the computation that we will come back to when the continuation is invoked. The **throw** term takes any continuation and any value

$\Gamma \vdash v \in \tau$	<i>(Value typing)</i>
	TV-CONT
	$\vdash s \in \tau$
	$\frac{}{\Gamma \vdash \mathbf{cont} \, s \in \mathbf{Cont} \, \tau}$

$\Gamma \vdash e \in \tau$	<i>(Term typing)</i>
TE-LETCC	TE-THROW
$\Gamma, x : \mathbf{Cont} \, \tau \vdash e \in \tau$	$\Gamma \vdash v_1 \in \mathbf{Cont} \, \tau'$
$\frac{}{\Gamma \vdash \mathbf{letcc} \, x \, \mathbf{in} \, e \in \tau}$	$\Gamma \vdash v_2 \in \tau'$
	$\frac{}{\Gamma \vdash \mathbf{throw} \, v_1 \, v_2 \in \tau}$

10.2.1 Continuations and the Curry-Howard Isomorphism

The Curry-Howard Isomorphism relates constructive logic with typed functional programming. It observes that there is a correspondence between logical propositions and types, and with the proofs of those propositions and programs that have those types. This idea is the basis of modern proof assistants based on dependent type theory, such as Rocq, Lean or Agda. To show that a proposition is true, you just need to find a program that has that type.

Definition 10.2.3 (Correspondence between propositions and types).

$A \wedge B$	$\tau_1 * \tau_2$
$A \vee B$	$\tau_1 + \tau_2$
$A \Rightarrow B$	$\tau_1 \rightarrow \tau_2$
$\neg A$	$A \rightarrow \text{Void}$
<i>True</i>	Unit
<i>False</i>	Void

This chart shows the correspondence between common logical combinators (conjunction, disjunction, implication, negation) and types. And, it makes sense that constructive proofs of these combinators correspond to programs. For example, a constructive proof of a conjunction $A \wedge B$ is a pair of proofs, one which shows A and another that shows B . We can see this in the definition of conjunction in Rocq standard library:

```
Inductive and (A B : Prop) : Prop :=
  conj : A -> B -> A /\ B.
```

Similarly, a proof of a disjunction $A \vee B$ is either a proof of A , or a proof of B . This corresponds to a sum type.

```
Inductive or (A B : Prop) : Prop :=
  or_introl : A -> A \vee B
  | or_intror : B -> A \vee B.
```

The proof of an implication $A \Rightarrow B$ is a function that takes a proof of A to a proof of B . This corresponds to a function. In constructive logic, negation $\neg A$ is the same as saying that A is impossible. In other words, it is a function that takes a proof of A to a proof of *False*—i.e. a type that has no inhabitants (so there can never be a proof of *False*).

For example, we can prove various properties by creating programs in Gallina, the language of the Rocq proof assistant. (Usually we don't write these programs directly; tactics let us construct them more conveniently.) For example, we can show (half) of one of the de Morgan laws as follows:

```
Definition demorgan1 : not P /\ not Q -> not (P \vee Q) :=
  fun npnq => match npnq with
    | conj np nq =>
      fun p_or_q =>
        match p_or_q with
          | or_introl p => np p
          | or_intror q => nq q
        end
      end.
```

The inverse of this law is also expressible using a proof term in Rocq.

```
Definition demorgan2 : not (P \vee Q) -> not P /\ not Q :=
  fun h => conj (fun p => h (or_introl p)) (fun q => h (or_intror q)).
```

However, there is a difficulty with the other de Morgan identity. One direction is also easy.

```
Definition demorgan3 : not P \vee not Q -> not (P /\ Q) :=
  fun h => match h with
```

```

| or_introl np => fun pq => match pq with
|   conj p q => np p
end
| or_intror nq => fun pq => match pq with
|   conj p q => pq q
end
end.

```

However, the remaining direction is not constructively provable.

```
Lemma demorgan4 : not (P /\ Q) -> not P \vee not Q.
```

There is no way to prove this proposition in Rocq without using any axioms.

What does this have to do with control operators? Suppose Rocq had explicit control over continuations. We can pretend that is the case by assuming that there is some continuation type `cont`, and the two operations `callcc` and `throw`. (To make it easier to add the control operator to Rocq, we use `callcc` instead of `letcc`. The former uses a function argument instead of a let expression to bind the current continuation.)

Furthermore, we can give reasonable definitions to the continuation type and the throw operation. All we *really* need to axiomatize is `callcc`.

```

Definition cont : Prop -> Prop := fun A => False.
Definition throw : forall A, cont A -> A -> False := fun A k v => k v.
Axiom callcc : forall A, ((cont A) -> A) -> A.

```

If you look closely at the type of `callcc`, you may recognize it as Pierce's law, usually written $((P \Rightarrow Q) \Rightarrow P) \Rightarrow P$. This axiom is equivalent to the law of the excluded middle, which states that any proposition is either true or false, $P \vee \neg P$. Adding either Pierce's law or the law of the excluded middle to constructive logic turns it into classical logic. That means that these axioms are consistent with constructive logic, but cannot be proven constructively. If you would like to reason using classical logic in Rocq, you need to assume one of these axioms.

These axioms are equivalent. For example, you can derive the law of the excluded middle using `callcc`.

```

Definition excluded_middle {A} : A \vee cont A :=
  callcc (fun k => or_intror (fun a =>
    throw k (or_introl a))).

```

This proof captures the current continuation k and then immediately asserts that A is not true. But if it was the case that A was actually true (i.e. someone called this function with some a that is a proof of A), then we can throw this proof to the captured continuation to assert that A was really true all along. It's a bit like time travel—we can assert whatever we want without evidence, but if we ever find out that we were wrong, we can go back in time and replace our answer with the other version.

Furthermore, with `callcc`, we can show the problematic de Morgan law from above:

```

Definition demorgan4 : not (P /\ Q) -> not P \vee not Q :=
  fun npq => callcc (fun k =>
    or_introl (fun p => throw k (or_intror (fun q => npq (conj p q))))).

```

What does this proof mean? To show that $\text{not}(P \wedge Q)$ implies $\text{not}P \vee \text{not}Q$, we take

a proof that both P and Q cannot be true simultaneously and then capture the current continuation k . At this point, we assert that it is P that is not true, by showing a proof that P leads to *False*. However, if we ever receive a proof of P , we use the captured continuation to go back to where we started and assert that no, it was q that was not true. If, in this case, we also receive a proof of Q , we can use our assumption with the evidence for P and Q that we have already received.

Another proposition that is equivalent to Pierce's law and the law of the excluded middle is double negation elimination.

$$\text{not}(\text{not}P) \Rightarrow P$$

This proposition holds classically, but not constructively. In other words, P in classical logic P is the same as its double negation, but they are different propositions in constructive logic.

In fact, even though a property P may only be provable using classical logic, we can prove its double negation, i.e. $\text{not}(\text{not}P)$, using constructive logic. To show this fact, we need a technique that allows us to transform a proof term that uses `callcc` into one that does not. And that technique is called *Continuation-Passing-Style conversion* or CPS conversion.

10.2.2 CPS conversion

CPS conversion is a translation from a language that includes explicit continuation operations, like `letcc` or `callcc` into one that does not include these constructs. It is useful for logicians, but it is more useful for programmers, providing access to continuation-based programming even when it is not supported by the underlying programming language.

The key idea of CPS is that we will translate the program so that every function will take a second argument: its continuation. Furthermore, once the function has computed its result it will call this continuation with that result instead of returning it. Functions never return when written in continuation-passing-style, they only jump to the next instruction. That means that in order to call a function, we need to give it not just its argument, but what it should do next—its continuation—sort of like putting a return address on the stack when we call a function.

Because the continuation is always available for every term, it is easy to implement `letcc x in e`; we only need to bind x to that continuation.

Our goal is to define a translations for terms $\mathcal{E}(e)_w$ and for types $\mathcal{T}(\tau)$, such that the following theorem holds.

Theorem 10.2.1 (CPS translation is type preserving).

If $\vdash e \in \tau$ then $\vdash \lambda x.\mathcal{E}(e)_x \in (\mathcal{T}(\tau) \rightarrow \mathbf{Void}) \rightarrow \mathbf{Void}$

The term translation $\mathcal{E}(e)_w$ takes two arguments: a term to translate e and a continuation w . This continuation should have type $\mathcal{T}(\tau) \rightarrow \mathbf{Void}$, so that the type of the top-level translation is the double negation (plus inner type translation) of the original type of the term.

The type translation mostly leaves the structure of the types alone. However, it makes two changes. For function types, instead of returning a result of type $\mathcal{T}(\tau_2)$, the function is transformed to take a continuation argument instead. The type of the continuation argument is $\mathcal{C}(\tau_2)$. The function doesn't return a type anymore (so its result type is `Void`), instead it passes its result to this continuation. We also use $\mathcal{C}(\tau)$ as the translation of the continuation type.

Definition 10.2.4 (Type translation).

$$\begin{aligned}
 \mathcal{T}(\text{Nat}) &= \text{Nat} \\
 \mathcal{T}(\text{Void}) &= \text{Void} \\
 \mathcal{T}(\tau_1 * \tau_2) &= \mathcal{T}(\tau_1) * \mathcal{T}(\tau_2) \\
 \mathcal{T}(\tau_1 + \tau_2) &= \mathcal{T}(\tau_1) + \mathcal{T}(\tau_2) \\
 \mathcal{T}(\alpha) &= \alpha \\
 \mathcal{T}(\mu\alpha.\tau) &= \mu\alpha.\mathcal{T}(\tau) \\
 \mathcal{T}(\tau_1 \rightarrow \tau_2) &= \mathcal{T}(\tau_1) \rightarrow \mathcal{C}(\tau_2) \rightarrow \text{Void} \\
 \mathcal{T}(\text{Cont } \tau) &= \mathcal{C}(\tau) \\
 \\
 \mathcal{C}(\tau) &= \mathcal{T}(\tau) \rightarrow \text{Void}
 \end{aligned}$$

A continuation $\mathcal{C}(\tau)$ is a function that wants an argument of type $\mathcal{T}(\tau)$. But what should its result type be? If we want the theorem above to hold, then it has to be **Void**. So continuations themselves can never return any answers.

Wait a minute. If continuations never return, and we translate e to $\lambda x.\mathcal{E}(e)_x$, i.e. a term waiting for its continuation, what continuation do we pass in for x to run the translated term? How do we get the result? Continuations cannot return results, so they need to communicate them in another way. Therefore, we add the **exit** v construct to the language. This term aborts the computation immediately. It doesn't return v , it *exits* with v , where an exited machine is a new terminal state.

$$\frac{\text{SSM-EXIT}}{\langle s, \text{exit } v \rangle \mapsto \langle [], \text{exit } v \rangle} \qquad \frac{\text{TE-EXIT} \quad \Gamma \vdash v \in \tau}{\Gamma \vdash \text{exit } v \in \tau'}$$

It may seem like a subtle difference between returning a value and exiting with a value, but the payoff is in the typing rule. A returned value may be used by another computation. So its type needs to be the same as the type of the value. But an exited value is *done*. No one can ever use it. So the typing rule for exit can have any type, including **Void**. That means that to run the translated term, we can pass in $\lambda y.\text{exit } y$ as the initial continuation.

The translation for terms $\mathcal{E}(e)_w$ is defined mutually with a translation for values $\mathcal{V}(v)$ and stacks $\mathcal{S}(s)$. We use the metavariable w for values used as continuations by the term translation. (The more common practice is to use a k , but already use that letter to stand for natural numbers.)

Definition 10.2.5 (Term, Value and Stack translation).

$$\begin{aligned}
\mathcal{V}(x) &= x \\
\mathcal{V}(\text{zero}) &= \text{zero} \\
\mathcal{V}(\text{succ } v) &= \text{succ } \mathcal{V}(v) \\
\mathcal{V}(\text{unit}) &= \text{unit} \\
\mathcal{V}(\lambda x.e) &= \lambda x.\text{ret } \lambda y.\mathcal{E}(e)_y \\
\mathcal{V}((v_1, v_2)) &= (\mathcal{V}(v_1), \mathcal{V}(v_2)) \\
\mathcal{V}(\text{inj}_1 v) &= \text{inj}_1 \mathcal{V}(v) \\
\mathcal{V}(\text{inj}_2 v) &= \text{inj}_2 \mathcal{V}(v) \\
\mathcal{V}(\text{fold } v) &= \text{fold } \mathcal{V}(v) \\
\mathcal{V}(\text{rec } x.v) &= \text{rec } x.\mathcal{V}(v) \\
\mathcal{V}(\text{cont } s) &= \mathcal{S}(s) \\
\\
\mathcal{E}(\text{ret } v)_w &= w \mathcal{V}(v) \\
\mathcal{E}(\text{let } x = e_1 \text{ in } e_2)_w &= \mathcal{E}(e_1)_{(\lambda x.\mathcal{E}(e_2)_w)} \\
\mathcal{E}(v_1 v_2)_w &= \text{let } x = \mathcal{V}(v_1) \mathcal{V}(v_2) \text{ in } x \text{ w} \\
\mathcal{E}(\text{letcc } x \text{ in } e)_w &= \text{let } x = \text{ret } w \text{ in } \mathcal{E}(e)_w \\
\mathcal{E}(\text{throw } v_1 v_2)_w &= \mathcal{V}(v_1) \mathcal{V}(v_2) \\
\mathcal{E}(\text{exit } v)_w &= \text{exit } \mathcal{V}(v) \\
\mathcal{E}(\text{prj}_1 v)_w &= \text{let } x = \text{prj}_1 \mathcal{V}(v) \text{ in } w \text{ x} \\
\mathcal{E}(\text{prj}_2 v)_w &= \text{let } x = \text{prj}_2 \mathcal{V}(v) \text{ in } w \text{ x} \\
\mathcal{E}(\text{unfold } v)_w &= \text{let } x = \text{unfold } \mathcal{V}(v) \text{ in } w \text{ x} \\
\mathcal{E}(\text{case } v \text{ of } \{0 \Rightarrow e_1; \mathbf{S} y \Rightarrow e_2\})_w &= \text{case } \mathcal{V}(v) \text{ of } \{0 \Rightarrow \mathcal{E}(e_1)_w; \mathbf{S} y \Rightarrow \mathcal{E}(e_2)_w\} \\
\mathcal{E}(\text{case } v \text{ of } \{\text{inj}_1 x \Rightarrow e_1; \text{inj}_2 y \Rightarrow e_2\})_w &= \text{case } \mathcal{V}(v) \text{ of } \{\text{inj}_1 x \Rightarrow \mathcal{E}(e_1)_w; \text{inj}_2 y \Rightarrow \mathcal{E}(e_2)_w\} \\
\\
\mathcal{S}([]) &= \lambda x.\text{exit } x \\
\mathcal{S}(\text{let } x = _ \text{ in } e_2 : s) &= \lambda x.\mathcal{E}(e_2)_{\mathcal{S}(s)}
\end{aligned}$$

The value translation produces a new value through structural recursion. For function values, the translation introduces a second abstraction (for the continuation y that will be passed to the function). This continuation is used to translate the body of the function. For continuation values, $\text{cont } s$, the translation uses the *stack translation* to convert an explicit stack into a continuation function. If the stack is empty, then the continuation function should exit with its value. Otherwise, if there is a frame on the stack, the continuation function uses the term translation to process that value, with the continuation computed from the rest of the stack.

The term translation is more interesting. Instead of returning a value, the translation passes it to the continuation w . Let expressions use the continuation argument for sequencing. In an application, the translated function expects a continuation argument in addition to the usual one. After the first (partial) application, the result is applied to the current continuation. For *letcc*, the current continuation value is named by a let expression before continuing the translation. For *throw*, the continuation value is applied to the argument, the current continuation is discarded. Similarly *exit*, discards the current continuation and just halts the program. The remainder of the cases merely propagate the continuation through the translation.

Finally, we need to define the translation of substitutions. We translate substitution by translating each of the types or values in the range of the substitution.

Definition 10.2.6 (Substitution translation).

$$\mathcal{T}(\sigma) \alpha = \mathcal{T}(\sigma \alpha)$$

$$\mathcal{V}(\sigma) x = \mathcal{V}(\sigma x)$$

Properties of the type translation The type translation commutes with renaming and substitution.

Lemma 10.2.1 (Type translation renaming).

$$\mathcal{T}(\tau)\langle\xi\rangle = \mathcal{T}(\tau\langle\xi\rangle)$$

Lemma 10.2.2 (Type translation substitution).

$$\mathcal{T}(\tau)[\mathcal{T}(\sigma)] = \mathcal{T}(\tau[\sigma])$$

This final property establishes the result that we need for recursive types:

Corollary 10.2.1 (Recursive type translation).

$$\mathcal{T}(\tau)[\mu\alpha.\mathcal{T}(\tau)/\alpha] = \mathcal{T}(\tau[\mu\alpha.\tau/\alpha])$$

Properties of the term and value translations To prove our main lemma, we need to generalize it to talk about open terms. However, the scope of the result of the translation is not the same scope as that of the input. When the scope is extended when going under an abstraction, the translation also adds a continuation argument. We can establish the relationship between the input and output contexts using the following inductive definition. Bindings introduced by abstractions use rule **TC-ABS**, whereas other bindings use rule **TC-UP**.

Definition 10.2.7 (Context translation).

	<i>(Typing context relation)</i>
TC-NIL	TC-ABS
	$\frac{\Gamma \sim \Delta}{(\Gamma, x:\tau_1) \sim (\Delta, x:\mathcal{T}(\tau_1), y:\mathcal{C}(\tau_2))}$
	TC-UP
	$\frac{\Gamma \sim \Delta}{(\Gamma, x:\tau_1) \sim (\Delta, x:\mathcal{T}(\tau_1))}$

We can now state our main lemma:

Lemma 10.2.3 (CPS translation is type preserving). Suppose $\Gamma \sim \Delta$.

1. If $\Gamma \vdash v \in \tau$ then $\Delta \vdash \mathcal{V}(v) \in \mathcal{T}(\tau)$
2. If $\Gamma \vdash e \in \tau$ and $\Delta \vdash w \in \mathcal{C}(\tau)$ then $\Delta \vdash \mathcal{E}(e)_w \in \mathbf{Void}$
3. If $\vdash s \in \tau_1 \rightsquigarrow \tau$ then $\vdash \mathcal{S}(s) \in \mathcal{C}(\tau)$

The proof of this lemma is by mutual induction on the typing derivation. By choosing a single variable as the initial continuation, we can observe how this translation uses double negation on the input type. Alternatively, we can use an initial continuation to get a complete program.

Corollary 10.2.2 (CPS translation is type preserving). If $\vdash e \in \tau$ then $\vdash \lambda x.\mathcal{E}(e)_x \in (\mathcal{T}(\tau) \rightarrow \mathbf{Void}) \rightarrow \mathbf{Void}$ and $\vdash \mathcal{E}(e)_{(\lambda y.\mathbf{exit}\ y)} \in \mathbf{Void}$.

To show that the semantics is preserved by the simulation, we first need to show two substitution lemmas. The first shows that substitution commutes with the translation. Note that the “continuation” w is part of the output scope, isn’t translated like the term, value or stack. Therefore, we still need to translate the value as we push it through the continuation.

Lemma 10.2.4 (Substitution). 1. $(\mathcal{E}(e)_w)[\mathcal{V}(v)/x] = \mathcal{E}(e[v/x])_{(w[\mathcal{V}(v)/x])}$

2. $(\mathcal{V}(v_1))[\mathcal{V}(v)/x] = \mathcal{V}(v_1[v/x])$

3. $(\mathcal{S}(s))[\mathcal{V}(v)/x] = \mathcal{S}(s[v/x])$

The second lemma, called *Kubstitution* by Pottier [Pot17], lets us substitute for variables that are free in the continuation only, and not in the expression.

Lemma 10.2.5 (Kubstitution). 1. If $x \notin \text{fv}(e)$ then $(\mathcal{E}(e)_w)[v/x] = \mathcal{E}(e[v/x])_{(w[\mathcal{V}(v)/x])}$

We can use these substitution lemmas to show that a step with the stack machine semantics can be simulated by zero-or-more steps with the usual small-step semantics (augmented with a small-step rule for `exit`).

Lemma 10.2.6 (Semantic correctness). If $\langle s_1, e_1 \rangle \mapsto \langle s_2, e_2 \rangle$ then $\mathcal{E}(e_1)_{\mathcal{S}(s_1)} \rightsquigarrow^* \mathcal{E}(e_2)_{\mathcal{S}(s_2)}$

10.3 Effect handlers

The control operator `letcc` demonstrates how to *reify* the stack as a first class value. However, in some ways this feature is both too powerful and too inflexible.

On one hand, `letcc` can be inefficient to implement because saved continuations can be used multiple times. This means that the programming language’s runtime needs to allocate memory for the saved stack and keep this memory around until the stack reference goes out of scope and can be garbage collected. A more efficient version is *one shot* continuations: saved stacks that can be jumped to at most once.

At the same time, these continuations are not *composeable*. The issue is that `letcc` captures the *entire* stack. But in some situations it may be useful to only capture *part* of the stack and compose that part with the current stack in a new context.

OCaml 5’s new feature, called *effect handlers*, are based on *one-shot delimited continuations*. In this section, we will work through their support for *delimited* continuations. However, for simplicity, we won’t model the fact that they are single use here.

10.3.1 Grammar and stack-based semantics

Below, we introduce the grammar for this extension. Just like we used `Nat` as the type of exceptions (which means that we could tell what exception was raised by looking at that number), we also represent effects using natural numbers. When a user *performs* an effect with `perform v`, v must be the number corresponding to an effect. Then an effect handler is a frame on the stack that executes some code if that particular effect has been performed.

The term $\text{try } e_1 \text{ with } x \Rightarrow e_2 ; \text{eff}_n y \Rightarrow e'_2$ installs an effect handler during the execution of e_1 . If e_1 returns a value, then the first branch e_2 is taken, where x is

bound to that value. However, if e_1 performs the effect n , then the second branch e'_2 is selected, and y is bound to the continuation of the effect. This is the part of the stack that occurs between where the effect was performed up-to and including the effect handler itself.

Definition 10.3.1 (Grammar for effect handlers).

$$\begin{array}{lll} \text{types} & \tau & ::= \quad \mathbf{DeCont} \, \tau_1 \, \tau_2 \mid \mathbf{Eff} \, \tau \\ \text{value} & v & ::= \quad \mathbf{cont} \, s \mid \mathbf{effect} \, n \\ \text{term} & e & ::= \quad \mathbf{perform} \, v \mid \mathbf{try} \, e_1 \, \mathbf{with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 \mid \mathbf{continue} \, v_1 \, v_2 \\ \text{frame} & frame & ::= \quad \mathbf{try_with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 \end{array}$$

Definition 10.3.2 (Effect handler operational semantics).

$m \mapsto m'$	<i>(Stack-based small-step rules)</i>
SSM-HANDLE	
$\overline{\langle s, \mathbf{try} \, e_1 \, \mathbf{with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 \rangle \mapsto \langle \mathbf{try_with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 : s, e_1 \rangle}$	
SSM-HANDLE-RET	
$\overline{\langle \mathbf{try_with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 : s, \mathbf{ret} \, v \rangle \mapsto \langle s, e_2[v/x] \rangle}$	
SSM-PERFORM	
$\overline{\langle frame : s, \mathbf{perform} \, (\mathbf{effect} \, n) \rangle \mapsto \mathbf{find_eff} \, (frame : s) \, n \, []}$	
SSM-CONTINUE	
$\overline{\langle s_1, \mathbf{continue} \, (\mathbf{cont} \, s_2) \, v \rangle \mapsto \langle s_2 ++ s_1, \mathbf{ret} \, v \rangle}$	

The rule **SSM-HANDLE** merely installs the effect handler by pushing the relevant information on the stack. This handler is popped in rule **SSM-HANDLE-RET**, where the returned value is passed to the first branch of the handler.

As before, the value $\mathbf{cont} \, s$ embeds a list of frames as value—what is different this time is that when we perform an effect, we need to search the current stack to find the handler for that effect, along the way saving all of the stack frames that we encounter while searching for the effect.

The meta-operation $\mathbf{find_eff} \, s \, n \, s'$ searches stack s , looking for a handler for effect n , and accumulating frames in argument s' . The rule **SSM-PERFORM** calls this operation with an initial accumulator of the empty stack.

$$\begin{aligned} \mathbf{find_eff} \, (frame : s) \, n \, s' &= \langle s, e'_2[\mathbf{cont} \, (\mathbf{rev} \, (frame : s')) / y] \rangle \\ &\text{when } frame = \mathbf{try_with} \, x \Rightarrow e_2 ; \mathbf{eff}_n \, y \Rightarrow e'_2 \\ &= \mathbf{find_eff} \, s \, n \, (frame : s') \\ &\text{otherwise} \\ \mathbf{find_eff} \, [] \, n \, s' &= \langle [], \mathbf{perform} \, (\mathbf{effect} \, n) \rangle \end{aligned}$$

When a matching frame is found, then the next machine stack is the one where the handler was installed s . The next machine term is the branch for that effect (e'_2). Furthermore, the variable y is replaced with a continuation composed of the accumulated stack frames along with the handler. (By keeping the current handler,

this semantics implements so called *deep* effect handlers. Alternative semantics discard this frame.) Because the accumulator is constructed while recurring down the stack, we must reverse the frames when constructing the continuation.

If the frame is not a handler for the current effect, then the operation calls itself recursively, adding that frame to the accumulator. Finally, if there is no handler for the effect, then the machine steps to a configuration with a unhandled effect.

This saved continuation, which is available in the effect handler can be invoked with the `continue` $v_1 v_2$ term. The operation of `continue` is a little like that of `throw`. However, instead of discarding the current stack, this rule appends the saved stack to the current stack with the notation $s_1 ++ s_2$. Stacks are represented as lists, so this is merely a list append operation.

10.3.2 Typing rules and type safety

To type check programs that involve effect handlers, we need a more informative type for our saved continuations. Our original typing judgment for stacks only tracked the type of the term that was compatible with the frame at the top of the stack. It did not say anything about the type of value that would eventually be returned by the machine. This is a feature of exceptions and our previous control operators: we can prove type safety for this type system without statically tracking this information. This makes sense because the only thing that is done with the final returned value is to return it. No more computation can happen because the stack is empty.

In this setting, though, we need to know the bottom of the partial stack that is reified by the effect handler because it will be composed with another stack. We need to know that this composition will work out. If this type information is not necessary (such as in checking a machine, or in the typing rule for a non-delimited continuation) we can merely ignore the second type.

Definition 10.3.3 (Delimited stack typing and values).

$\vdash m \text{ ok}$	<i>(Machine is well typed)</i>
$\frac{\begin{array}{c} \text{M-OK} \\ \vdash s \in \tau_1 \rightsquigarrow \tau_2 \\ \vdash e \in \tau_1 \end{array}}{\vdash \langle s, e \rangle \text{ ok}}$	
$\vdash s \in \tau_1 \rightsquigarrow \tau_2$	<i>(Stack typing)</i>
$\frac{\begin{array}{c} \text{TS-NIL} \\ \text{TS-CONS} \quad \vdash \text{frame} \in \tau_1 \rightsquigarrow \tau_2 \\ \vdash s \in \tau_2 \rightsquigarrow \tau_3 \end{array}}{\vdash (\text{frame} : s) \in \tau_1 \rightsquigarrow \tau_3}$	
$\Gamma \vdash v \in \tau$	<i>(Value typing)</i>
$\frac{\begin{array}{c} \text{TV-DECONT} \\ \vdash s \in \tau_1 \rightsquigarrow \tau_2 \end{array}}{\Gamma \vdash \text{cont } s \in \text{DeCont } \tau_1 \tau_2}$ $\frac{\text{TV-EFF}}{\Gamma \vdash \text{effect } n \in \text{Eff } (\Phi(n))}$	

Therefore, our revised stack typing rule includes two types. These two types are recorded in **DeCont** $\tau_1 \tau_2$, the type of a delimited continuation.

We also need to be careful about the typing of effect values. We assume that our type system is parameterized by some map Φ , which maps effect numbers to their types.

Definition 10.3.4 (Typing rules for effect handlers).

$$\begin{array}{c}
 \boxed{\Gamma \vdash e \in \tau} \qquad \qquad \qquad (\text{Term typing}) \\
 \\
 \begin{array}{c}
 \text{TE-PERFORM} \\
 \dfrac{\Gamma \vdash v \in \mathbf{Eff} \tau}{\Gamma \vdash \mathbf{perform} v \in \tau} \qquad \qquad \qquad \text{TE-CONTINUE} \\
 \dfrac{\Gamma \vdash v_1 \in \mathbf{DeCont} \tau_1 \tau \quad \Gamma \vdash v_2 \in \tau_1}{\Gamma \vdash \mathbf{continue} v_1 v_2 \in \tau}
 \end{array} \\
 \\
 \text{TE-HANDLE} \\
 \dfrac{\begin{array}{c} \Gamma \vdash e_1 \in \tau_1 \\ \Gamma, x:\tau_1 \vdash e_2 \in \tau \\ \Gamma, y:\mathbf{DeCont}(\Phi(n)) \tau \vdash e'_2 \in \tau \end{array}}{\Gamma \vdash \mathbf{try} e_1 \text{ with } x \Rightarrow e_2 ; \mathbf{eff}_n y \Rightarrow e'_2 \in \tau} \\
 \\
 \boxed{\vdash f \in \tau_1 \rightsquigarrow \tau_2} \qquad \qquad \qquad (\text{Frame typing}) \\
 \\
 \text{TF-HANDLE} \\
 \dfrac{\begin{array}{c} , x:\tau_1 \vdash e_2 \in \tau \\ , y:\mathbf{DeCont}(\Phi(n)) \tau \vdash e'_2 \in \tau \end{array}}{\vdash \mathbf{try_} \text{ with } x \Rightarrow e_2 ; \mathbf{eff}_n y \Rightarrow e'_2 \in \tau_1 \rightsquigarrow \tau}
 \end{array}$$

When we perform an effect, the type of the effect must match the context where we perform that effect. The reason is that when we capture the continuation, its type includes the context of the perform: the τ_1 of **DeCont** $\tau_1 \tau_2$ says that the sort of value expected by the captured stack is τ_1 — i.e. the sort of value that we can put in this spot. When we **continue** with this continuation, we must give it a value of the appropriate type. We also must only **continue** in a context that matches the bottom of the saved stack, as this stack will be composed with the one at this point in the execution.

In rule **TE-HANDLE**, the two branches must have the same type τ . In the first case, the term acts like a let, and continues with e_2 with x as the result of e_1 . In the second case, the code e'_2 executes in response to e_1 performing an effect of type $\mathbf{Eff}(\Phi(n))$. That means that the saved delimited continuation will accept a value of type $\Phi(n)$, and produce a value of type τ . Therefore y , which is bound to this saved continuation, must have type $\mathbf{DeCont}(\Phi(n)) \tau$.

The typing rule for the frame (i.e. effect handler on the stack) mirrors that of the term. It records that it should be paired with an expression of type e_1 and produces a value of type τ .

Most of the preservation and progress proof from the previous section carries over. The most difficult part of setting this proof up is getting the definitions right. If we say, tried to use type **Cont** τ instead of **DeCont** $\tau_1 \tau_2$ for the type of the saved stacks, we would not be able to show our results.

The key part of the preservation lemma is showing that we perform an effect, then the next machine state is well-formed.

Lemma 10.3.1 (Effect finding). If $\vdash s \in \tau_1 \rightsquigarrow \tau$ and $\vdash \mathbf{rev} s_1 \in \Phi(n) \rightsquigarrow \tau_1$ (for any s_1), then $\vdash \mathbf{find_eff} s n s_1 \text{ ok}$

The proof of this lemma is by induction on the stack typing for s . Our accumulation stack, s_1 , needs to accept a term compatible with effect n and produce a value that is compatible with the top of the stack s .

When we perform an effect, the top of the stack must have type $\Phi(n)$, according to rule **T-PERFORM**. The accumulator starts empty, so its bottom type must also be $\Phi(n)$. The lemma ensures that starting from these initial conditions, no matter where we find the effect handler we will have a good machine state.

10.4 Further reading

The treatment of exceptions and control operations was inspired by PFPL chapters 29 and 30. The material on effect handlers was inspired by Xavier Leroy's course "Control Structures" <https://xavierleroy.org/CdF/2023-2024/>. The draft book for this course is available at <https://xavierleroy.org/control-structures/>. More information about CPS conversion and its verification using Autosubst is available in Pottier's paper [Pot17].

11

(Untyped) program equivalence

What does it mean for two closed programs e_1 and e_2 to be *equivalent*?

- They have the same syntax.
- They have the same syntax, up to renaming of bound variables.
- One evaluates to the other. In other words, $e_1 \rightsquigarrow^* e_2$ or $e_2 \rightsquigarrow^* e_1$.
- They evaluate to the same value. In other words, there exists some v , such that $e_1 \rightsquigarrow^* \text{ret } v$ and $e_2 \rightsquigarrow^* \text{ret } v$.
- They evaluate to some common term (not necessarily a value). In other words, there exists some e , such that $e_1 \rightsquigarrow^* e$ and $e_2 \rightsquigarrow^* e$.
- They either both evaluate to a value or both diverge. In other words, $\exists v_1, e_1 \rightsquigarrow^* \text{ret } v_1$ if and only if $\exists v_2, e_2 \rightsquigarrow^* \text{ret } v_2$.
- No closed function that returns a natural number can distinguish them. In other words *for all* $e, e_1 \rightsquigarrow^* \text{ret } k$ iff $e_2 \rightsquigarrow^* \text{ret } k$.

And what if the programs are open? We cannot evaluate them, so how can we figure out whether they are equivalent? And why does it matter whether programs are equivalent? How can we make use of this property?

11.1 Preliminaries

In this chapter, we will examine a few different definitions of equivalence for a call-by-value version of the *untyped* lambda calculus, augmented with numbers, recursive functions and unit. As in the previous chapters, we will work with a fine-grained language, and use a stack-based operational semantics for evaluation.

Definition 11.1.1 (Syntax).

$$\begin{array}{ll}
 \text{values} & v ::= x \mid \text{unit} \mid \text{fun } x \ y \ . \ e \mid 0 \mid \mathbf{S} \ v \\
 \text{terms} & e ::= \text{ret } v \mid \text{let } x = e_1 \text{ in } e_2 \mid v_1 \ v_2 \mid \text{case } v \text{ of } \{0 \Rightarrow e_0; \mathbf{S} \ y \Rightarrow e_1\} \\
 \text{frame} & \text{frame} ::= \text{let } x = _ \text{ in } e_2 \\
 \text{stack} & s ::= [] \mid \text{frame} : s
 \end{array}$$

The small-step operational semantics is similar to the one found in Chapter 9, and makes use of a stack of frames. We define the stack semantics in terms of an underlying primitive reduction relation (written \rightarrow but not shown here).

Definition 11.1.2 (Small-step semantics).

$$\boxed{m \mapsto m'} \quad (m \text{ small-steps to } m')$$

$$\frac{\text{SSM-PRIM}}{e \rightarrow e'} \quad \frac{\text{SSM-PUSH}}{\langle s, e \rangle \mapsto \langle s, e' \rangle} \quad \frac{\text{SSM-POP}}{\langle \text{let } x = _ \text{ in } e_2 \rangle \mapsto \langle \text{let } x = _ \text{ in } e_2 : s, e_1 \rangle}$$

$$\frac{}{\langle \text{let } x = _ \text{ in } e_2 : s, \text{ret } v \rangle \mapsto \langle s, e_2[v/x] \rangle}$$

We also work with the reflexive, transitive closure of this small-step relation, written $m \mapsto^* m'$, and a similar variant that tracks the number of steps taken during evaluation.

Definition 11.1.3 (Stepped small-step semantics).

$$\boxed{m \mapsto^k m'} \quad (m \text{ small-steps to } m' \text{ in exactly } k \text{ steps})$$

$$\frac{\text{SN-REFL}}{m \mapsto^0 m} \quad \frac{\text{SN-TRANS}}{\frac{m_0 \mapsto m_1}{\frac{m_1 \mapsto^k m_2}{m_0 \mapsto^{\mathbf{S} k} m_2}}}$$

We also define predicates that talk about when a machine evaluates to a terminal state: i.e. an empty stack paired with a returned value.

Definition 11.1.4 (Halts).

A machine *halts*, written $\text{halts } m$, when $m \mapsto^* \langle [], \text{ret } v \rangle$. A machine *halts in k*, written $\text{halts}^k m$, when $m \mapsto^k \langle [], \text{ret } v \rangle$.

These predicates are closed under forward and reverse evaluation.

Lemma 11.1.1 (Halts forward and reverse). Suppose $m \mapsto m'$.

1. $\text{halts } m$ implies $\text{halts } m'$.
2. $\text{halts } m'$ implies $\text{halts } m$.
3. $\text{halts}^k m'$ implies $\text{halts}^{(\mathbf{S} k)} m'$

4. $\text{halts}^k m$ implies $\text{halts}^j m'$ for some j , such that $S j = k$.

We can also say when one machine *approximates* (or *approximates in k*) another machine.

Definition 11.1.5 (Machine approximation).

We have $m_1 \sqsubseteq m_2$ when $\text{halts } m_1$ implies $\text{halts } m_2$ and $m_1 \sqsubseteq^k m_2$ when $\text{halts}^k m_1$ implies $\text{halts } m_2$.

The approximation relation $m_1 \sqsubseteq m_2$ is a pre-order as it is both reflexive and transitive. The stepped approximation $m_1 \sqsubseteq^k m_2$ is a reflexive, step-indexed relation. (It is not transitive because m_2 may halt in using any number of steps.) When we want to generalize over steps, we write it as $m_1 \sqsubseteq^\cdot m_2$.

Finally, we can overload notation to say when a term approximates another by putting it on the empty stack.

Definition 11.1.6 (Term approximation).

We have $e_1 \sqsubseteq e_2$ when $\langle \emptyset, e_1 \rangle \sqsubseteq \langle \emptyset, e_2 \rangle$.

11.1.1 Scopes and scoped relations

We are not going to talk about typing for this language, but we will be explicit about *scopes*. A scope, X , is a set variables that are allowed to appear free in a value or term. In the accompanying Rocq formalization, we use a de Bruijn representation of variables, so scopes can be tracked using a single number (an upper bound on the indices that can appear in the term). Here we will use sets of names instead.

Definition 11.1.7 (Well-scoped). A term/value is *well scoped*, written $X \vdash e$ (or $X \vdash v$), when its free variables are a subset of X .

Definition 11.1.8 (Closed). A term/value is *closed*, written $\emptyset \vdash e$ (or $\emptyset \vdash v$), when it is well-scoped with an empty set of variables. (We use \emptyset to indicate the empty set.)

Definition 11.1.9 (*Scoped binary relation*). A scoped binary relation is a relation between two terms or values with the same scope. We use the metavariables $\mathcal{R}_\mathcal{E}$ and $\mathcal{R}_\mathcal{V}$ for arbitrary scoped relations on terms and values and assert that terms and values are in these relations using the judgment forms $X \vdash e \mathcal{R}_\mathcal{E} e'$ and $X \vdash v \mathcal{R}_\mathcal{V} v'$.

Definition 11.1.10 (Closing substitution). A substitution σ is a *closing substitution*, written $X \vdash \sigma$, when it maps all variables in X to closed values. In other words, we have $\emptyset \vdash \sigma x$ for all $x \in X$.

Definition 11.1.11 (*Scoped pre-order*). A scoped relation is a *scoped pre-order* when it is reflexive and transitive.

Definition 11.1.12 (*Scoped equivalence*). A scoped relation is a *scoped equivalence* when it is reflexive, transitive and symmetric.

We can trivially turn any preorder into an equivalence by relating terms in both directions. Therefore, we are going to mainly work with preorders going forward. This is both easier (as we don't have to show both directions) and sometimes more informative (as there are terms that will only be related in one direction).

11.2 Scoped preorders and equivalences

11.2.1 CIU pre-order \leq_{CIU}

Let's look at our first preorder. This relation orders expressions by their termination behavior when evaluated by arbitrary stacks, under arbitrary substitutions.

Definition 11.2.1 (CIU Preorder). We say that $X \vdash e \leq_{CIU} e'$ when for all stacks s , and for all closing substitutions $X \vdash \sigma$, $\text{halts}(s, e[\sigma])$ implies $\text{halts}(s, e'[\sigma])$. In other words, when $\langle s, e[\sigma] \rangle \sqsubseteq \langle s, e'[\sigma] \rangle$.

This relation is trivially a preorder because it is defined in terms of the approximation preorder.

Furthermore, if this relation holds in both directions, then we get an equivalence relation, which we call *observational equivalence* or *CIU equivalence*.

Definition 11.2.2 (Observational equivalence). Two terms are *observationally equivalent* when $X \vdash e \leq_{CIU} e'$ and $X \vdash e' \leq_{CIU} e$.

We call this idea of running the program using arbitrary stacks as making *observations* on the terms. This fact that we quantify over all stacks is important. Just because two terms co-terminate does not mean that they are equivalent. One program could terminate with 3 and the other terminate with 5. However, one way to observe the difference between these terms is to run these two programs on a stack that diverges when its argument is greater than 4. If there is no way to tell them apart then we call them equivalent.

Because we are working with an untyped language, some evaluations will get stuck. However, the definition of *halts* requires the abstract machine to evaluate completely to a terminal configuration. As a result, programs that always get stuck will be equal to each other (and to programs that always diverge).

If we were working with a language where all programs terminate, then we would define observations differently. We would instead choose an observation type, like *bool* or *nat*, and require that both terms produce the same value of this type. It is important that syntactic equivalence is sufficient. We can't use functions as our observation type because we sometimes want to identify functions that are not syntactically equal, like $\lambda x.x + 0$ and $\lambda x.x$. However, limiting observations to booleans or natural numbers is not a problem. When we quantify over all stacks, we can relate expressions that return any type of value.

Usage equivalence is sometimes called an *extensional* equivalence because it is based on what the program does (how it interacts with the external environment) instead of how the program is implemented (its internal definitions).

The letters \leq_{CIU} in the name of this relation come from *uses of closed instantiations* [MT91], i.e. we create closed instantiations with arbitrary substitutions, and use them with arbitrary stacks.

11.2.2 Contextual pre-order and equivalence (\leq_C)

A program context is any part of a term or value that can surrounds a single hole, written $\#$. We use the metavariable C to refer to contexts, formed by the following grammar:

Definition 11.2.3 (Program Context).

$$\begin{aligned} C ::= & \# \mid C v_2 \mid v_1 C \mid \text{let } x = C \text{ in } e_2 \mid \text{let } x = e_1 \text{ in } C \\ & \mid \text{ret } C \mid \text{case } v \text{ of } \{0 \Rightarrow C; S y \Rightarrow e_1\} \mid \text{case } v \text{ of } \{0 \Rightarrow e_0; S y \Rightarrow C\} \\ & \mid \text{succ } C \mid \text{fun } x y. C \mid \text{case } C \text{ of } \{0 \Rightarrow e_0; S y \Rightarrow e_1\} \end{aligned}$$

We use the notation $C\{e\}$ and $C\{v\}$ for plugging the hole in a context with a term or value. This notation has the implicit constraint that the hole is in a term position for the first operation and a value position for the second operation. For example, we can only plug a value into the context $\#\text{zero}$, and we can only plug an expression into the context $\text{fun } x y. \#$.

We use contexts to define the *contextual* pre-order on terms and values. Instead of using stacks and closing substitutions, we instead quantify over arbitrary contexts.

Definition 11.2.4 (Contextual pre-order (terms)). Define $X \vdash e_1 \leq_C e_2$ when we have $C\{e_1\} \sqsubseteq C\{e_2\}$ for any C .

Definition 11.2.5 (Contextual pre-order (values)). Define $X \vdash v_1 \leq_C v_2$ when $C\{v_1\} \sqsubseteq C\{v_2\}$ for any C .

As above, we can turn this definition into an equivalence by requiring it in both directions.

This definition is the one most commonly referred to by the name “contextual equivalence”. And it indeed is the most flexible to use. If we know that two terms are equivalent then we can replace any one term with another in any context without changing the termination behavior of the program.

Furthermore, we will also be able to show that programs with contextually equivalent parts are contextually equivalent.

Lemma 11.2.1 (Contextual replacement). If $X \vdash e_1 \leq_C e_2$ then $\emptyset \vdash C\{e_1\} \leq_C C\{e_2\}$.

11.2.3 Contextual Equivalence (Pitts) (\leq_{CTX})

Yet another way to define when two terms are related is *coinductively*. We will define a *contextual pre-order* as the largest pre-order that is adequate and compatible. This terminology overlaps with the definition in the previous section. While a little confusing, we will soon see that they are the same relation.

Adequacy describes a minimal requirement on relations. If they say that two closed terms are equal, then both terms must halt when evaluated. While this seems too weak to be a basis of an equivalence, in conjunction with the other properties, we will see that it is enough.

Definition 11.2.6 (Adequate). A scoped relation \mathcal{R}_E is *adequate* when $\emptyset \vdash e \mathcal{R}_E e'$ means that *halts* $\langle[], e\rangle$ implies *halts* $\langle[], e'\rangle$. In other words, $\langle[], e\rangle \sqsubseteq \langle[], e'\rangle$.

Compatibility means that terms are related when their subterms are related. Because our syntax is mutually defined by terms and values, we actually need compatibility to hold for a pair of relations.

Definition 11.2.7 (Compatible). A pair of scoped relations \mathcal{R}_E and \mathcal{R}_V are *compatible* when the following properties hold:

1. $X \vdash x \mathcal{R}_V x$
2. $X \vdash \mathbf{unit} \mathcal{R}_V \mathbf{unit}$
3. $X \vdash \mathbf{zero} \mathcal{R}_V \mathbf{zero}$
4. $X \vdash v_1 \mathcal{R}_V v_2$ implies $X \vdash \mathbf{succ} v_1 \mathcal{R}_V \mathbf{succ} v_2$
5. $X, x, y \vdash e_1 \mathcal{R}_E e_2$ implies $X \vdash \mathbf{fun} x. y. e_1 \mathcal{R}_V \mathbf{fun} x. y. e_2$
6. $X \vdash v_1 \mathcal{R}_V v_2$ implies $X \vdash \mathbf{ret} v_1 \mathcal{R}_E \mathbf{ret} v_2$
7. $X \vdash e_1 \mathcal{R}_E e_2$ and $X, x \vdash e'_1 \mathcal{R}_E e'_2$ implies $X \vdash \mathbf{let} x = e_1 \mathbf{in} e'_1 \mathcal{R}_E \mathbf{let} x = e_2 \mathbf{in} e'_2$
8. $X \vdash v_1 \mathcal{R}_V v_2$ and $X \vdash v'_1 \mathcal{R}_V v'_2$ implies $X \vdash v_1 v'_1 \mathcal{R}_E v_2 v'_2$.
9. $X \vdash v_1 \mathcal{R}_V v_2$ and $X \vdash e_1 \mathcal{R}_E e_2$ and $X, x \vdash e'_1 \mathcal{R}_E e'_2$ implies $X \vdash \mathbf{case} v_1 \mathbf{of} \{0 \Rightarrow e_1; S x \Rightarrow e'_1\} \mathcal{R}_E \mathbf{case} v'_1 \mathbf{of} \{0 \Rightarrow e'_1; S x \Rightarrow e'_2\}$

If we have more syntactic forms in our language, we need to add more clauses to this definition. The property that we want to hold is that we have one clause for each syntactic form, so that we can guarantee that the relation is reflexive.

Lemma 11.2.2 (Compatible relations are reflexive). If we have Compatible $\mathcal{R}_E \mathcal{R}_V$, then both \mathcal{R}_E and \mathcal{R}_V are reflexive.

Now, here is our key definition:

Definition 11.2.8 (Contextual preorder). Define $X \vdash e \leq_{CTX} e'$ if there is some compatible, adequate pre-order \mathcal{R}_E , such that $X \vdash e \mathcal{R}_E e'$.

As above, we define contextual equivalence when this preorder holds in both directions, i.e. $X \vdash e \leq_{CTX} e'$ and $X \vdash e' \leq_{CTX} e$.

What this means is that to show two terms are equivalent, we just need to find any relation that equates them, as long as that relation has the right properties. That is a useful definition because it gives us a practical way to show that terms are contextually equivalent.

11.2.4 A logical relation \leq_{log}

Finally, we can define a *step-indexed* logical relation. This definition is mutually defined in terms of a term relation $\mathcal{C}[\![e_1 \leq e_2]\!]_k$, stack relation $\mathcal{S}[\![s_1 \leq s_2]\!]_k$, and value relation $\mathcal{V}[\![v_1 \leq v_2]\!]_k$.

Definition 11.2.9 (Step-indexed Logical Relation).

$$\begin{aligned}
 \mathcal{C}[\![e_1 \leq e_2]\!]_k &= \forall s_1, s_2, \mathcal{S}[\![s_1 \leq s_2]\!] \implies_k \langle s_1, e_1 \rangle \sqsubseteq^\cdot \langle s_2, e_2 \rangle \\
 \mathcal{S}[\![s_1 \leq s_2]\!]_k &= \forall v_1, v_2, \mathcal{V}[\![v_1 \leq v_2]\!] \implies_k \langle s_1, \mathbf{ret} v_1 \rangle \sqsubseteq^\cdot \langle s_2, \mathbf{ret} v_2 \rangle \\
 \mathcal{V}[\![\mathbf{unit} \leq \mathbf{unit}]\!]_k &= \text{always} \\
 \mathcal{V}[\![\mathbf{zero} \leq \mathbf{zero}]\!]_k &= \text{always} \\
 \mathcal{V}[\![(\mathbf{succ} v_1) \leq (\mathbf{succ} v_2)]\!]_k &= \triangleright_k \mathcal{V}[\![v_1 \leq v_2]\!] \\
 \mathcal{V}[\![(\mathbf{fun} x. y. e) \leq v_2]\!]_k &= \forall v v', \triangleright_k (\mathcal{V}[\![v \leq v']\!] \implies \mathcal{C}[\![e[v/x, (\mathbf{fun} x. y. e)/y] \leq (v_2 v')]\!])
 \end{aligned}$$

This relation extends to a step-indexed relation on closing substitutions.

Definition 11.2.10 (Related closing substitutions).

$$X \vdash \sigma_1 \leq_{\log}^k \sigma_2 \text{ when } \mathcal{V}[\sigma_1 x] \leq \sigma_2 x]_k$$

Which we can use to create scoped preorders on values, terms and stacks.

Definition 11.2.11 (Related values, terms and stacks).

1. $X \vdash v_1 \leq_{\log} v_2$ when $X \vdash \sigma_1 \leq_{\log}^k \sigma_2$ implies $\mathcal{V}[v_1[\sigma_1]] \leq v_2[\sigma_2]]_k$ for all k .
2. $X \vdash e_1 \leq_{\log} e_2$ when $X \vdash \sigma_1 \leq_{\log}^k \sigma_2$ implies $\mathcal{C}[e_1[\sigma_1]] \leq e_2[\sigma_2]]_k$, for all k
3. $\vdash s_1 \leq_{\log} s_2$ when $\mathcal{S}[s_1] \leq s_2]_k$, for all k .

11.3 Examples

The amazing result is that all of these definitions coincide. All of the pre-orders (and induced equivalence relations) from the previous section relate the same terms.

But why do we have so many definitions? How can we use them?

Pitts' definition of contextual equivalence reassures us that, from our starting definitions of what it means for programs to approximate one another, we have the coarsest definition possible. There isn't some other definition of equivalence out there that would relate more terms. Furthermore, if we find any preorder that is compatible and adequate, that relates two terms, then we know that they are in the contextual preorder.

The traditional definition of Contextual equivalence justifies the refactoring that programmers do, or that compilers do when they optimize code. It says that you can start with any program identify some subterm anywhere and replace it with something equivalent without changing the semantics of the program.

The CIU pre-order is the most convenient definition for reasoning about reduction. For example, we can use this definition to show that inlining or “beta-reduction” is a sound program transformation.

Lemma 11.3.1 (Beta reduction and expansion).

1. $X \vdash (\mathbf{fun} x y. e) v \leq_{\text{CIU}} e[v/x, (\mathbf{fun} x y. e)/y]$
2. $X \vdash e[v/x, (\mathbf{fun} x y. e)/y] \leq_{\text{CIU}} (\mathbf{fun} x y. e) v$

This sort of reasoning is available for any pair of terms, where one steps to another.

1. $X \vdash e[v/x] \leq_{\text{CIU}} \mathbf{let} x = \mathbf{ret} v \mathbf{in} e$
2. $X \vdash \mathbf{let} x = \mathbf{ret} v \mathbf{in} e \leq_{\text{CIU}} e[v/x]$

From these (and other) facts, we can show that the usual definition of untyped β -equivalence implies contextual equivalence. However, β -equivalence is not as large as contextual equivalence: there are terms that are not β -equivalent that are not contextually equivalent.

$$\begin{array}{ccccc}
 & (1) & (6) & (4) & \\
 \leq_C & \Rightarrow & \Rightarrow & \Rightarrow & \\
 & (2) & (5) & (3) & \\
 & \Leftarrow & \Leftarrow & \Leftarrow &
 \end{array}
 \quad
 \begin{array}{c}
 \leq_{CTX} \\
 (1) \quad (6) \quad (4) \\
 \Rightarrow \quad \Rightarrow \quad \Rightarrow \\
 \Leftarrow \quad \Leftarrow \quad \Leftarrow
 \end{array}
 \quad
 \begin{array}{c}
 \leq_{CIU} \\
 (2) \quad (5) \quad (3) \\
 \Leftarrow \quad \Leftarrow \quad \Leftarrow
 \end{array}
 \quad
 \begin{array}{c}
 \leq_{log} \\
 (4)
 \end{array}$$

Figure 11.1: Equivalence steps

For example, we cannot use β -equivalence to show that $\mathbf{fun} \ x \ y. \ x$ is equal to $\mathbf{fun} \ x \ y. \ x + 0$.

Finally, the logical relation is the most convenient for showing the soundness of “ η -reductions”.

Lemma 11.3.2 (Function η -reduction). $X \vdash \mathbf{fun} \ x \ y. \ v \ x \leq_{log} v$

Lemma 11.3.3 (Nat η -reduction). $X \vdash \mathbf{case} \ v \ \mathbf{of} \ \{0 \Rightarrow 0; \mathbf{S} \ x \Rightarrow \mathbf{S} \ x\} \leq_{log} v$

On the other hand, using our definitions, η -expansions are not validated by our definitions. We can use the traditional definition of contextual equivalence to show that

Lemma 11.3.4 (Failure of η -expansion). It is not the case that $X \vdash v \leq_C \mathbf{fun} \ x \ y. \ v \ x$ for all v .

Choose v as zero and the context as $\# \text{zero}$.

11.4 Relating the relations

Finally, let’s prove that these relations are equivalent.

Theorem 11.4.1 (The four relations are equivalent). We have $X \vdash e \leq_C e'$ iff $X \vdash e \leq_{CTX} e'$ iff $X \vdash e \leq_{CIU} e'$ iff $X \vdash e \leq_{log} e'$.

We don’t do this all at once: it takes several steps to make all of these definitions line up.

1. Show that \leq_C implies \leq_{CTX} by showing that it is an adequate and compatible preorder.
2. Show that \leq_{CTX} implies \leq_C , by using the adequacy of any relation in \leq_{CTX} , the fact that contexts preserve compatible relations.
3. Show that \leq_{log} implies \leq_{CIU} , by showing that the logical relation is reflexive.
4. Show that \leq_{CIU} implies \leq_{log} . This also follows from the reflexivity of the logical relation, and the fact that it is closed under CIU-equivalence.
5. Show that \leq_{CIU} implies \leq_{CTX} . This requires showing that it is adequate and transitive (easy). To show that it is compatible, we use the fact that it is equivalent to a compatible relation (\leq_{log}).
6. Show that \leq_{CTX} implies \leq_{CIU} . This happens in two steps, first for closed terms, by induction over the arbitrary stack. Then, we extend the proof to open terms by showing that contextual equivalence is closed under substitution.

11.4.1 \leq_C implies \leq_{CTX}

Showing that \leq_C implies \leq_{CTX} means showing that it is an adequate and compatible pre-order.

The reflexivity and transitivity of \leq_C immediately fall out from the reflexivity and transitivity of the approximation relation \sqsubseteq . Furthermore, showing that contextual equivalence is adequate is similarly straightforward. We merely need to use the context that is a single hole, and the result immediately holds.

However, to show compatibility, we need to be able to compose contexts, written $C_1 \circ C_2$ and reason about how that composition interacts with plugging terms.

Lemma 11.4.1 (Compose plug).

$$C_2\{C_1\{e\}\} = C_1 \circ C_2\{e\}$$

In the compatibility proof, we instantiate premises with the appropriate contexts. For example, so show that contextual equivalence is compatible for application expressions, we prove the following sublemma.

Lemma 11.4.2 (App Compatibility). If $X \vdash v_1 \leq_C v_2$ and $X \vdash v'_1 \leq_C v'_2$ then $X \vdash v_1 v'_1 \leq_C v_2 v'_2$.

Proof. To show $X \vdash v_1 v'_1 \leq_C v_2 v'_2$, we pick an arbitrary context C , and show that $C\{v_1 v'_1\} \sqsubseteq C\{v_2 v'_2\}$. We instantiate our first assumption with the context $v_2 \# \circ C$ to conclude that $C\{v_2 v'_1\} \sqsubseteq C\{v_2 v'_2\}$. Then, we instantiate the second assumption with $\# v_1 \circ C$ to conclude that $C\{v_1 v'_1\} \sqsubseteq C\{v_2 v'_1\}$. Chaining these approximations together transitively gives us the desired result. \square

11.4.2 Show that \leq_{CTX} implies \leq_C

Next, note that if $X \vdash e_1 \leq_{CTX} e_2$, then we have $X \vdash e_1 \mathcal{R}_E e_2$ for some adequate compatible relation \mathcal{R}_E .

Lemma 11.4.3 (\leq_{CTX} implies \leq_C). If $X \vdash e_1 \leq_{CTX} e_2$ then $X \vdash e_1 \leq_C e_2$.

Proof. Let C be an arbitrary context. We want to show that $C\{e_1\} \sqsubseteq C\{e_2\}$. By adequacy, it suffices to show that $\emptyset \vdash C\{e_1\} \mathcal{R}_E C\{e_2\}$. But this follows by lemma 11.4.4 and the compatibility of \leq_{CTX} . \square

Lemma 11.4.4 (Compatible plug). If \mathcal{R}_E and \mathcal{R}_V are compatible, then $X \vdash e_1 \mathcal{R}_E e_2$ implies $X \vdash C\{e_1\} \mathcal{R}_E C\{e_2\}$ and $X \vdash v_1 \mathcal{R}_V v_2$ implies $X \vdash C\{v_1\} \mathcal{R}_V C\{v_2\}$.

Proof. This result follows by induction on the structure of C , using the compatibility properties of \mathcal{R}_E and \mathcal{R}_V along with reflexivity. \square

11.4.3 \leq_{log} implies \leq_{CIU}

First, before we reason about \leq_{log} , because we have defined it as a step-indexed relation, we need to show that it is downward closed. This is not difficult to do.

Next, we show the fundamental property for this logical relation.

The fundamental property of *binary logical relations* is that they are reflexive. In other words, we want to show that for any $X \vdash e$ and we have $X \vdash e \leq_{log} e$. However, we have already observed that reflexivity is a consequence of compatibility. Therefore, we show that the term and value relations are compatible.

Lemma 11.4.5 (Logical equivalence is compatible). The term and value relations are compatible (see definition 11.2.7).

Reflexivity is all we need to show that the logical preorder implies the CIU preorder. The compatibility relation above gives us reflexivity for our relations between terms, values, stacks and closing substitutions.

Lemma 11.4.6 (Reflexivity).

1. If $X \vdash e$ then $X \vdash e \leq_{\log} e$.
2. If $X \vdash v$ then $X \vdash v \leq_{\log} v$.
3. If $\vdash s$ then $\vdash s \leq_{\log} s$.
4. If $X \vdash \sigma$ then $X \vdash \sigma \leq_{\log} \sigma$.

Lemma 11.4.7 (\leq_{\log} implies \leq_{CIU}). If $X \vdash e \leq_{\log} e'$ then $X \vdash e \leq_{\text{CIU}} e'$.

Proof. To show that e is CIU-related to e' , we need to pick an arbitrary stack s and arbitrary substitution σ , and show that $\langle s, e[\sigma] \rangle \sqsubseteq \langle s, e'[\sigma] \rangle$. Now, suppose $\langle s, e[\sigma] \rangle \mapsto^k \langle \[], \text{ret } v \rangle$ in some number of steps k . From our reflexivity property, we know that $\emptyset \vdash \sigma \leq_{\log} \sigma$ and $\vdash s \leq_{\log} s$. Therefore, by the definition of the logical relation, we know that $\mathcal{C}[\langle e[\sigma] \leq e'[\sigma] \rangle]_k$, and thus, $\langle s, e[\sigma] \rangle \sqsubseteq^k \langle s, e'[\sigma] \rangle$. \square

11.4.4 \leq_{CIU} implies \leq_{\log}

To show this direction, we first establish that the logical relation is closed under the CIU relation.

Lemma 11.4.8 (\leq_{\log} is closed under \leq_{CIU}). If $X \vdash e \leq_{\log} e'$ and $X \vdash e' \leq_{\text{CIU}} e''$, then $X \vdash e \leq_{\log} e''$.

Proof. Let k , $X \vdash \sigma \leq_{\log} \sigma'$ be arbitrary. We want to show that $\mathcal{C}[\langle e[\sigma] \leq e''[\sigma'] \rangle]_k$. This means showing that for all $j \leq k$ and for all stacks $S[\langle s \leq s' \rangle_j]$, we have $\langle s, e[\sigma] \rangle \sqsubseteq^j \langle s', e''[\sigma'] \rangle$. By our second assumption, we know that $\langle s', e'[\sigma'] \rangle \sqsubseteq \langle s', e''[\sigma'] \rangle$. This holds transitively, if we can show $\langle s, e[\sigma] \rangle \sqsubseteq^j \langle s', e'[\sigma'] \rangle$. However, that is a consequence of our first assumption as the stacks and substitutions are related. \square

Lemma 11.4.9 (\leq_{CIU} implies \leq_{\log}). If $X \vdash e \leq_{\text{CIU}} e'$ then $X \vdash e \leq_{\log} e'$.

Proof. Reflexivity of the logical relation gives us $X \vdash e \leq_{\log} e$. With Lemma 11.4.8 above, we can conclude $X \vdash e \leq_{\log} e'$. \square

At this point we know that the logical and CIU preorders coincide.

11.4.5 \leq_{CIU} implies \leq_{CTX}

To show that \leq_{CIU} implies \leq_{CTX} , we need to show that it is an adequate, compatible preorder.

Showing that it is adequate and transitive is straightforward from the definition. To show that it is compatible, we use the fact that it is equivalent to a compatible relation (\leq_{\log}).

Lemma 11.4.10 (Compatibility). \leq_{CIU} is compatible (using the logical relation as the value relation).

11.4.6 \leq_{CTX} implies \leq_{CIU}

This happens in two steps. We first show this result for closed terms. Then, we extend the proof to open terms by showing that contextual equivalence is closed under substitution.

Lemma 11.4.11 (**Closed \leq_{CTX} implies \leq_{CIU}**). $\emptyset \vdash e \leq_{\text{CTX}} e'$ implies $\langle s, e \rangle \sqsubseteq \langle s, e' \rangle$ for any s .

Proof. By assumption we know that $\emptyset \vdash e \mathcal{R}_{\mathcal{E}} e'$ using some preorder that is adequate and compatible. We prove $\langle s, e \rangle \sqsubseteq \langle s, e' \rangle$ by induction on the stack s . If the stack is empty, then the result follows via adequacy. If the stack is $\text{let } x = _ \text{ in } e_2 : s$, by compatibility and reflexivity of the arbitrary relation we know that $\emptyset \vdash \text{let } x = e \text{ in } e_2 \mathcal{R}_{\mathcal{E}} \text{let } x = e' \text{ in } e_2$. By induction, we know that $\langle s, \text{let } x = e \text{ in } e_2 \rangle \sqsubseteq \langle s, \text{let } x = e' \text{ in } e_2 \rangle$. By the fact that approximation is closed under forward and reverse reduction, we can conclude $\langle \text{let } x = _ \text{ in } e_2 : s, e \rangle \sqsubseteq \langle \text{let } x = _ \text{ in } e_2 : s, e' \rangle$. \square

Lemma 11.4.12 (**Substitutivity**). If $X \vdash e \leq_{\text{CTX}} e'$ and $X \vdash \sigma$ then $\emptyset \vdash e[\sigma] \leq_{\text{CTX}} e'[\sigma]$.

Proof. We prove this by induction on the length of X (i.e. the number of variables in the domain of σ). We can show the substitutivity result for single substitutions by using the transitivity of \leq_{CTX} to join the following results:

1. $X \vdash e[v/x] \leq_{\text{CTX}} \text{let } x = \text{ret } v \text{ in } e$ (by showing them \leq_{CIU} equivalent)
2. $X \vdash \text{let } x = \text{ret } v \text{ in } e \leq_{\text{CTX}} \text{let } x = \text{ret } v \text{ in } e'$ (by compatibility of any underlying relation in \leq_{CTX})
3. $X \vdash \text{let } x = \text{ret } v \text{ in } e' \leq_{\text{CTX}} e'[v/x]$ (by showing them \leq_{CIU} equivalent).

\square

At this point we have shown that the relations \leq_{CTX} , \leq_{\log} and \leq_{CIU} are equivalent. Because of this equivalence, it is now not difficult to show that \leq_{CTX} is compatible.

11.5 Further Reading

This chapter is inspired by Pitts [Pit10]. Similar material is available in PFPL chapter 47 [Har16], although that is for a typed language, not the untyped lambda calculus. Furthermore, similar mechanized results for CBPV are available in Forster et al. [FSSS19].

12

Typed Program Equivalence

In the last chapter we developed several definitions of contextual equivalence for an untyped lambda calculus and showed them equivalent.

Here, we add types.

Our goal is twofold.

First, with the addition of a simple type system, all programs in our language terminate. That means that we cannot use the approximation relation from the last chapter as the basis of our equivalence. Instead, we need to use an *observable* base type, such as booleans or natural numbers. The good news is that because all well-typed terms produce values, we get to work with equality directly. There is no need to bother with approximation; we will work with equivalence relations.

Second, the addition of types means that our equivalences will be able to relate *more* terms and values. We saw in the last chapter that η -expansion was not valid in the untyped language. Here, because we can use types to rule out stuck terms, and control the contexts, substitutions and stacks that we use for observations, we will be able to relate more programs.

That said, the general format of this development follows the previous chapter.

12.1 Preliminaries

The language that we are working with is similar to that from before: a fine-grained call-by-value language with a stack based semantics. The semantics and type system is similar to that of Chapter 10. For concreteness, we show the syntax of the language below.

Definition 12.1.1 (Syntax).

```
values      v ::= x | unit | λx.e | 0 | S v
terms      e ::= ret v | let x = e1 in e2 | v1 v2 | case v of {0 ⇒ e0; S y ⇒ e1}
frame     frame ::= let x = _ in e2
stack      s ::= [] | frame : s
```

12.1.1 Nat-equivalence

Instead of relating terms by their termination behavior, we equate closed machines when they return the same natural number. (Note: types are not yet part of this definition. Machines that do not finish with natural numbers or that get stuck will all be related to each other.)

Definition 12.1.2 (Nat equivalence). We say $m \equiv_N m'$ when $m \mapsto^* \langle[], \text{ret } k\rangle$ iff $m' \mapsto^* \langle[], \text{ret } k\rangle$.

This is an equivalence relation: it is reflexive, symmetric and transitive. Furthermore, it is closed under evaluation. If we have $m \mapsto^* m'$, then $m \equiv_N m'$.

12.1.2 Typed (equivalence) relations

Next, instead of talking about relations between terms in the same scope, as we did in the previous chapter, we will talk about relations between terms that have the same *type*. This means that we need to define what it means to have a typed relation and define their properties.

Definition 12.1.3 (Typed Term Relation). A typed term relation $\Gamma \vdash e \mathcal{R}_E e' : \tau$ is a relation between terms e and e' , such that $\Gamma \vdash e \in \tau$ and $\Gamma \vdash e' \in \tau$.

Similarly, we also have typed *value* relations, written $\Gamma \vdash v \mathcal{R}_V v' : \tau$, which imply that both values are well-typed. When we want to generalize over typed term and typed value relations, we use the notation $\Gamma \vdash t \in \tau$ for the typing relation and $\Gamma \vdash t \mathcal{R} t' : \tau$ for the typed binary relation.

Definition 12.1.4 (Typed Reflexive). A typed relation is *reflexive* when $\Gamma \vdash t \in \tau$ implies $\Gamma \vdash t \mathcal{R} t : \tau$.

Definition 12.1.5 (Typed Transitive). A typed relation is *transitive* when $\Gamma \vdash t_1 \mathcal{R} t_2 : \tau$ and $\Gamma \vdash t_2 \mathcal{R} t_3 : \tau$ implies $\Gamma \vdash t_1 \mathcal{R} t_3 : \tau$.

We also have typed variants of adequacy (for term relations) and compatibility (for a term and value relation together).

Definition 12.1.6 (Adequate). A typed relation is *adequate* when $\vdash e_1 \mathcal{R}_E e_2 : \mathbf{Nat}$ implies $e_1 \equiv N e_2$.

Definition 12.1.7 (Compatible). A pair of scoped relations \mathcal{R}_E and \mathcal{R}_V are *compatible* when the following properties hold:

1. $\Gamma \vdash x \mathcal{R}_V x : \Gamma x$
2. $\Gamma \vdash \mathbf{unit} \mathcal{R}_V \mathbf{unit} : \mathbf{Unit}$
3. $\Gamma \vdash \mathbf{zero} \mathcal{R}_V \mathbf{zero} : \mathbf{Nat}$
4. $\Gamma \vdash v_1 \mathcal{R}_V v_2 : \mathbf{Nat}$ implies $\Gamma \vdash \mathbf{succ} v_1 \mathcal{R}_V \mathbf{succ} v_2 : \mathbf{Nat}$
5. $\Gamma, x:\tau_1 \vdash e_1 \mathcal{R}_E e_2 : \tau_2$ implies $\Gamma \vdash \lambda x. e_1 \mathcal{R}_V \lambda x. e_2 : \tau_1 \rightarrow \tau_2$
6. $\Gamma \vdash v_1 \mathcal{R}_V v_2 : \tau$ implies $\Gamma \vdash \mathbf{ret} v_1 \mathcal{R}_E \mathbf{ret} v_2 : \tau$
7. $\Gamma \vdash e_1 \mathcal{R}_E e_2 : \tau_1$ and $\Gamma, x:\tau_1 \vdash e'_1 \mathcal{R}_E e'_2 : \tau_2$ implies $\Gamma \vdash \mathbf{let} x = e_1 \mathbf{in} e'_1 \mathcal{R}_E \mathbf{let} x = e_2 \mathbf{in} e'_2 : \tau_2$

8. $\Gamma \vdash v_1 \mathcal{R}_V v_2 : \tau_1 \rightarrow \tau_2$ and $\Gamma \vdash v'_1 \mathcal{R}_V v'_2 : \tau_1$ implies $\Gamma \vdash v_1 v'_1 \mathcal{R}_E v_2 v'_2 : \tau_2$.
9. $\Gamma \vdash v_1 \mathcal{R}_V v_2 : \mathbf{Nat}$ and $\Gamma \vdash e_1 \mathcal{R}_E e_2 : \tau$ and $\Gamma, x:\mathbf{Nat} \vdash e'_1 \mathcal{R}_E e'_2 : \tau$ implies
 $\Gamma \vdash \mathbf{case}\ v_1\ \mathbf{of}\ \{0 \Rightarrow e_1; S\ x \Rightarrow e'_1\} \mathcal{R}_E \mathbf{case}\ v'_1\ \mathbf{of}\ \{0 \Rightarrow e'_1; S\ x \Rightarrow e'_2\} : \tau$

Compatible relations are reflexive on well typed terms and values.

Lemma 12.1.1 (Compatible relations are reflexive). If the typed relations $\Gamma \vdash e \mathcal{R}_E e' : \tau$ and $\Gamma \vdash v \mathcal{R}_V v' : \tau$ are compatible, then

1. If $\Gamma \vdash e \in \tau$ then $\Gamma \vdash e \mathcal{R}_E e : \tau$.
2. If $\Gamma \vdash v \in \tau$ then $\Gamma \vdash v \mathcal{R}_V v \in \tau$.

12.2 Contextual, CIU, and logical equivalence

As before, we have several definitions of equivalence. (For simplicity, we drop the traditional notion of contextual equivalence and concentrate on the other three definitions in this chapter.)

Definition 12.2.1 (Pitt's contextual equivalence). Define $\Gamma \vdash e \equiv_{CTX} e' : \tau$ if there is some transitive, adequate typed term relation \mathcal{R}_E , with compatible typed value relation \mathcal{R}_V , such that $\Gamma \vdash e \mathcal{R}_E e' : \tau$.

Definition 12.2.2 (CIU equivalence). Define $\Gamma \vdash e \equiv_{CIU} e' : \tau$ when $\Gamma \vdash e \in \tau$, and $\Gamma \vdash e' \in \tau$, and for all $\Gamma \vdash \sigma$ and $\vdash s \in \tau \rightsquigarrow \mathbf{Nat}$, we have $\langle s, e[\sigma] \rangle \equiv_N \langle s, e'[\sigma] \rangle$.

Definition 12.2.3 (Logical relation). Define the following mutual logical relations by recursion on type structure.

$$\begin{aligned}
 \mathcal{C}[e_1 \equiv e_2 \in \tau] &= \vdash e_1 \in \tau \text{ and } \vdash e_2 \in \tau \text{ and} \\
 &\quad \forall s_1, s_2, \mathcal{S}[s_1 \equiv s_2 \in \tau] \text{ implies } \langle s_1, e_1 \rangle \equiv_N \langle s_2, e_2 \rangle \\
 \mathcal{S}[s_1 \equiv s_2 \in \tau] &= \vdash s_1 \in \tau \rightsquigarrow \mathbf{Nat} \text{ and } \vdash s_2 \in \tau \rightsquigarrow \mathbf{Nat} \text{ and} \\
 &\quad \forall v_1, v_2, \mathcal{V}[v_1 \leq v_2] \text{ implies } \langle s_1, \mathbf{ret}\ v_1 \rangle \equiv_N \langle s_2, \mathbf{ret}\ v_2 \rangle \\
 \mathcal{V}[v_1 \equiv v_2 \in \tau] &= \vdash v_1 \in \mathbf{Unit} \text{ and } \vdash v_2 \in \mathbf{Unit} \\
 \mathcal{V}[k_1 \equiv k_2 \in \mathbf{Nat}] &= k_1 = k_2 \\
 \mathcal{V}[\lambda x. e_1 \equiv \lambda x. e_2 \in \tau_1 \rightarrow \tau_2] &= x:\tau_1 \vdash e_1 \in \tau_2 \text{ and } x:\tau_1 \vdash e_2 \in \tau_2 \text{ and} \\
 &\quad \forall v v', \mathcal{V}[v \equiv v' \in \tau_1] \text{ implies } \mathcal{C}[e_1[v_1/x] \equiv e_2[v_2/x] \in \tau_2]
 \end{aligned}$$

We can then extend the logical relation to open terms and values in the usual way.

Definition 12.2.4 (Open Logical Relation).

1. $\Gamma \vdash \sigma_1 \equiv_{log} \sigma_2$ when $\mathcal{V}[(\sigma_1 x) \equiv (\sigma_2 x)] \in (\Gamma x)$
2. $\Gamma \vdash e \equiv_{log} e' : \tau$ when $\mathcal{C}[e[\sigma] \equiv e'[\sigma'] \in \tau]$, for arbitrary $\Gamma \vdash \sigma_1 \leq_{log} \sigma_2$.
3. $\Gamma \vdash v \equiv_{log} v' : \tau$ when $\mathcal{V}[v[\sigma] \equiv v'[\sigma'] \in \tau]$, for arbitrary $\Gamma \vdash \sigma_1 \leq_{log} \sigma_2$.

12.3 Examples

We can prove the same equalities as in the previous chapter. For example, we can use this definition to show that inlining or “beta-reduction” is a sound program transformation. This sort of reasoning is available for any pair of terms, where one steps to another.

Lemma 12.3.1 (Beta reduction). If $\Gamma \vdash (\lambda x.e) v \in \tau$ then $\Gamma \vdash (\lambda x.e) v \equiv_{CIU} e[v/x] : \tau$.

Lemma 12.3.2 (Let reduction). If $\Gamma \vdash \text{let } x = \text{ret } v \text{ in } e_2 \in \tau$ then $\Gamma \vdash \text{let } x = \text{ret } v \text{ in } e_2 \equiv_{CIU} e_2[v/x] : \tau$

However, with the addition of types, we can also show *new properties* that only hold for typed relations.

Any value with a function type is equal to an explicit function value.

Lemma 12.3.3 (Unit eta equivalence).

If $\Gamma \vdash v \in \tau_1 \rightarrow \tau_2$ then $\Gamma \vdash v \equiv_{log} \lambda x.v x : \tau_1 \rightarrow \tau_2$.

Any value with type **unit** is equal to the unit value. (And by transitivity and symmetry, all values with type **unit** are equal.)

Lemma 12.3.4 (Unit eta equivalence).

If $\Gamma \vdash v \in \text{Unit}$ then $\Gamma \vdash v \equiv_{log} \text{unit} : \text{Unit}$.

12.4 Equivalence between the relations

We follow the same path as the previous chapter when showing that our definitions are the same. We are skipping the first two steps, but we can recreate the remaining four steps in order.

1. We show that logical equivalence is compatible. This gives us the fact that it is reflexive, which we can use to show that logical equivalence implies CIU equivalence.
2. Next we show that logical equivalence is closed under CIU equivalence. This, plus the reflexivity of logical equivalence, gives us that CIU equivalence implies logical equivalence.
3. Now, because CIU equivalence is the same as logical equivalence, we can show that it is compatible. This, along with easy proofs of adequacy the fact that CIU is an equivalence relation, allows us to show that CIU equivalence implies CTX equivalence.
4. Finally, we next show that CTX equivalence implies CIU equivalence for closed terms. After showing that CTX equivalence is closed under substitution, we can lift the result to open terms terms too.

12.5 What next?

Types strengthen our definition of program equivalence. In CIU equivalence, we only need to think about arbitrary stacks and substitutions of the correct type. That means that we limit in the sorts of observations we need to make to argue that terms are equivalent. We don't need to consider the case where we put a function on a stack that will immediately test for zero. That gives us the important η -laws about values. Notably, these η -laws hold even if we add nontermination (via recursive definitions or recursive types) to the language, because they are properties of values.

However, our definitions of program equivalence in this chapter is weaker than we might expect because it doesn't tell us that programs terminate. We have presented it this way for easy comparison to the previous chapter, but this is not the only way to set up these definitions. For example, PFPL Chapter 46 [Har16] defines (traditional) contextual equivalence such that terms are related only when they both terminate with the same natural number, after being plugged into any arbitrary, appropriately typed program context. (PFPL calls this relation observational equivalence.) This chapter then shows that this definition is equivalent to a logical relation that relates terms of type `nat` that evaluate to the same natural number, and terms of function type that take related arguments to related results. One can use this logical relation to show that all programs terminate.

Bibliography

- [Ahm04] Amal Ahmed. *Semantics of Types for Mutable State*. PhD thesis, Princeton University, 2004.
- [AM01] Andrew W. Appel and David McAllester. An indexed model of recursive types for foundational proof-carrying code. *ACM Trans. Program. Lang. Syst.*, 23(5):657–683, September 2001.
- [Bar84] Hendrik Pieter Barendregt. *The Lambda Calculus: Its Syntax and Semantics*. Sole distributors for the U.S.A. and Canada, Elsevier Science Pub. Co., New York, N.Y., 1984.
- [Cha13] Arthur Charguéraud. Pretty-big-step semantics. In Matthias Felleisen and Philippa Gardner, editors, *Programming Languages and Systems*, pages 41–60, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [de 72] N.G de Bruijn. Lambda calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the church-rosser theorem. *Indagationes Mathematicae (Proceedings)*, 75(5):381–392, 1972.
- [FSSS19] Yannick Forster, Steven Schäfer, Simon Spies, and Kathrin Stark. Call-by-push-value in coq: operational, equational, and denotational theory. In *Proceedings of the 8th ACM SIGPLAN International Conference on Certified Programs and Proofs*, CPP 2019, page 118–131, New York, NY, USA, 2019. Association for Computing Machinery.
- [Gö58] Kurt Gödel. Über eine bisher noch nicht benützte erweiterung des finiten standpunktes. *Dialectica*, 12(3):280, 1958.
- [GL86] David K. Gifford and John M. Lucasen. Integrating functional and imperative programming. In *Proceedings of the 1986 ACM Conference on LISP and Functional Programming*, LFP ’86, page 28–38, New York, NY, USA, 1986. Association for Computing Machinery.
- [Har16] Robert Harper. *Practical Foundations for Programming Languages*. Cambridge University Press, USA, 2nd edition, 2016.
- [Kah87] G. Kahn. Natural semantics. In Franz J. Brandenburg, Guy Vidal-Naquet, and Martin Wirsing, editors, *STACS 87*, pages 22–39, Berlin, Heidelberg, 1987. Springer Berlin Heidelberg.

- [Ler06] Xavier Leroy. Coinductive big-step operational semantics. In Peter Sestoft, editor, *Programming Languages and Systems*, pages 54–68, Berlin, Heidelberg, 2006. Springer Berlin Heidelberg.
- [LG88] J. M. Lucassen and D. K. Gifford. Polymorphic effect systems. In *Proceedings of the 15th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL ’88, page 47–57, New York, NY, USA, 1988. Association for Computing Machinery.
- [LPT03] Paul Blain Levy, John Power, and Hayo Thielecke. Modelling environments in call-by-value programming languages. *Information and Computation*, 185(2):182–210, 2003.
- [Men87] Paul Francis Mendler. Inductive definitions in type theory, 1987.
- [Mil78] Robin Milner. A theory of type polymorphism in programming. *Journal of Computer and System Sciences*, 17(3):348–375, 1978.
- [MT91] Ian Mason and Carolyn Talcott. Equivalence in functional languages with effects. *Journal of Functional Programming*, 1(3):287–327, 1991.
- [PdAC⁺25] Benjamin C. Pierce, Arthur Azevedo de Amorim, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cătălin Hrițcu, Vilhelm Sjöberg, Andrew Tolmach, and Brent Yorgey. *Programming Language Foundations*, volume 2 of *Software Foundations*. Electronic textbook, 2025. Version 6.7, <http://softwarefoundations.cis.upenn.edu>.
- [Pie02] Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, 2002.
- [Pit10] Andrew M. Pitts. Step-Indexed Biorthogonality: a Tutorial Example. In Amal Ahmed, Nick Benton, Lars Birkedal, and Martin Hofmann, editors, *Modelling, Controlling and Reasoning About State*, volume 10351 of *Dagstuhl Seminar Proceedings (DagSemProc)*, pages 1–10, Dagstuhl, Germany, 2010. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [Pot17] François Pottier. Revisiting the cps transformation. Available from <https://cambium.inria.fr/~fpottier/publis/fpottier-cps.pdf>, July 2017.
- [SNO⁺07] Peter Sewell, Francesco Zappa Nardelli, Scott Owens, Gilles Peskine, Thomas Ridge, Susmit Sarkar, and Rok Strniša. Ott: effective tool support for the working semanticist. In *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming*, ICFP ’07, page 1–12, New York, NY, USA, 2007. Association for Computing Machinery.
- [TKDB24] Amin Timany, Robbert Krebbers, Derek Dreyer, and Lars Birkedal. A logical approach to type soundness. *J. ACM*, 71(6), November 2024.
- [WF94] A.K. Wright and M. Felleisen. A syntactic approach to type soundness. *Information and Computation*, 115(1):38–94, 1994.